

# TECHNOLOGY



## Container Orchestration using Kubernetes

# TECHNOLOGY

## Core Concepts

# A Day in the Life of a DevOps Engineer

You have recently started working as a DevOps engineer and are struggling to comprehend the functionality, different components, and features of Kubernetes.

The goal is to understand how the cluster is configured and how pods are created. Your team is looking for someone who can deploy applications, maintain their status, and provide all cluster details.

To achieve these and some additional features, you will learn a few concepts in this lesson that will help you find a solution for the given scenario.





# Learning Objectives

By the end of this lesson, you will be able to:

- Distinguish between Docker and Containerd for effective container runtime selection
- Explore Containerd thoroughly and effectively apply its critical commands for container management
- Comprehend the workings of container orchestration in Kubernetes for efficient container management
- Describe the functions of Etcd, Controller, Scheduler, and Kubelet for seamless container orchestration within Kubernetes



# Learning Objectives

By the end of this lesson, you will be able to:


- Differentiate kube-proxy, pods, ReplicaSet, and deployment to manage Kubernetes clusters
- Classify common deployments use cases to facilitate organized management
- Summarize the concepts of containers and policies in Kubernetes to articulate their significance in practical use



## Overview of Kubernetes

# Kubernetes Design Principles

Kubernetes is a collection of building components. The components work together to offer mechanisms for deploying, maintaining, and scaling applications.

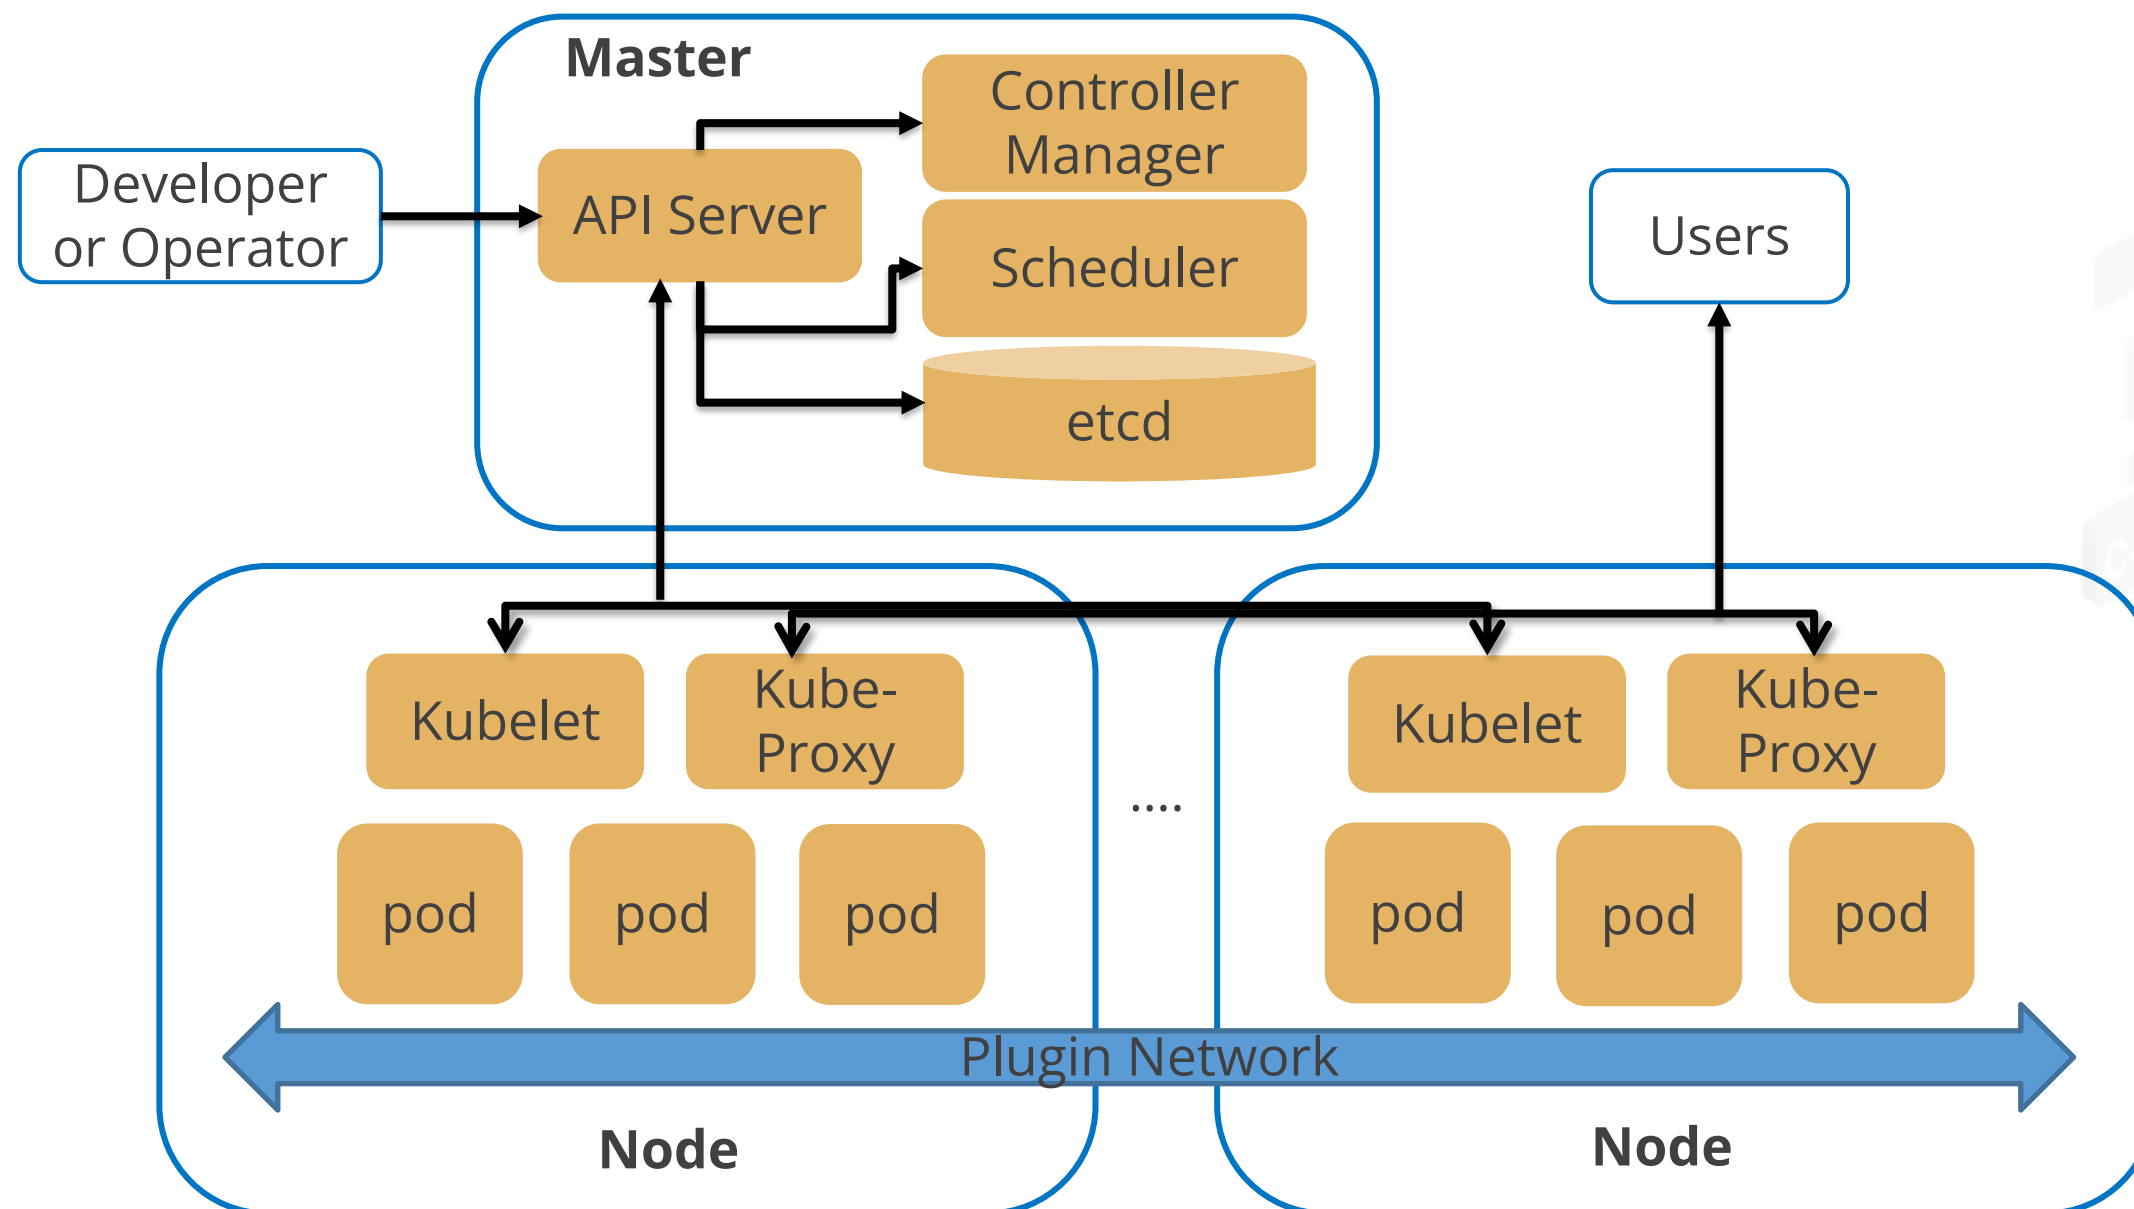


Kubernetes is coupled and expanded to handle a variety of workloads and scheduling scenarios.

The platform gains control over computing and storage resources by defining resources.

# Kubernetes Architecture

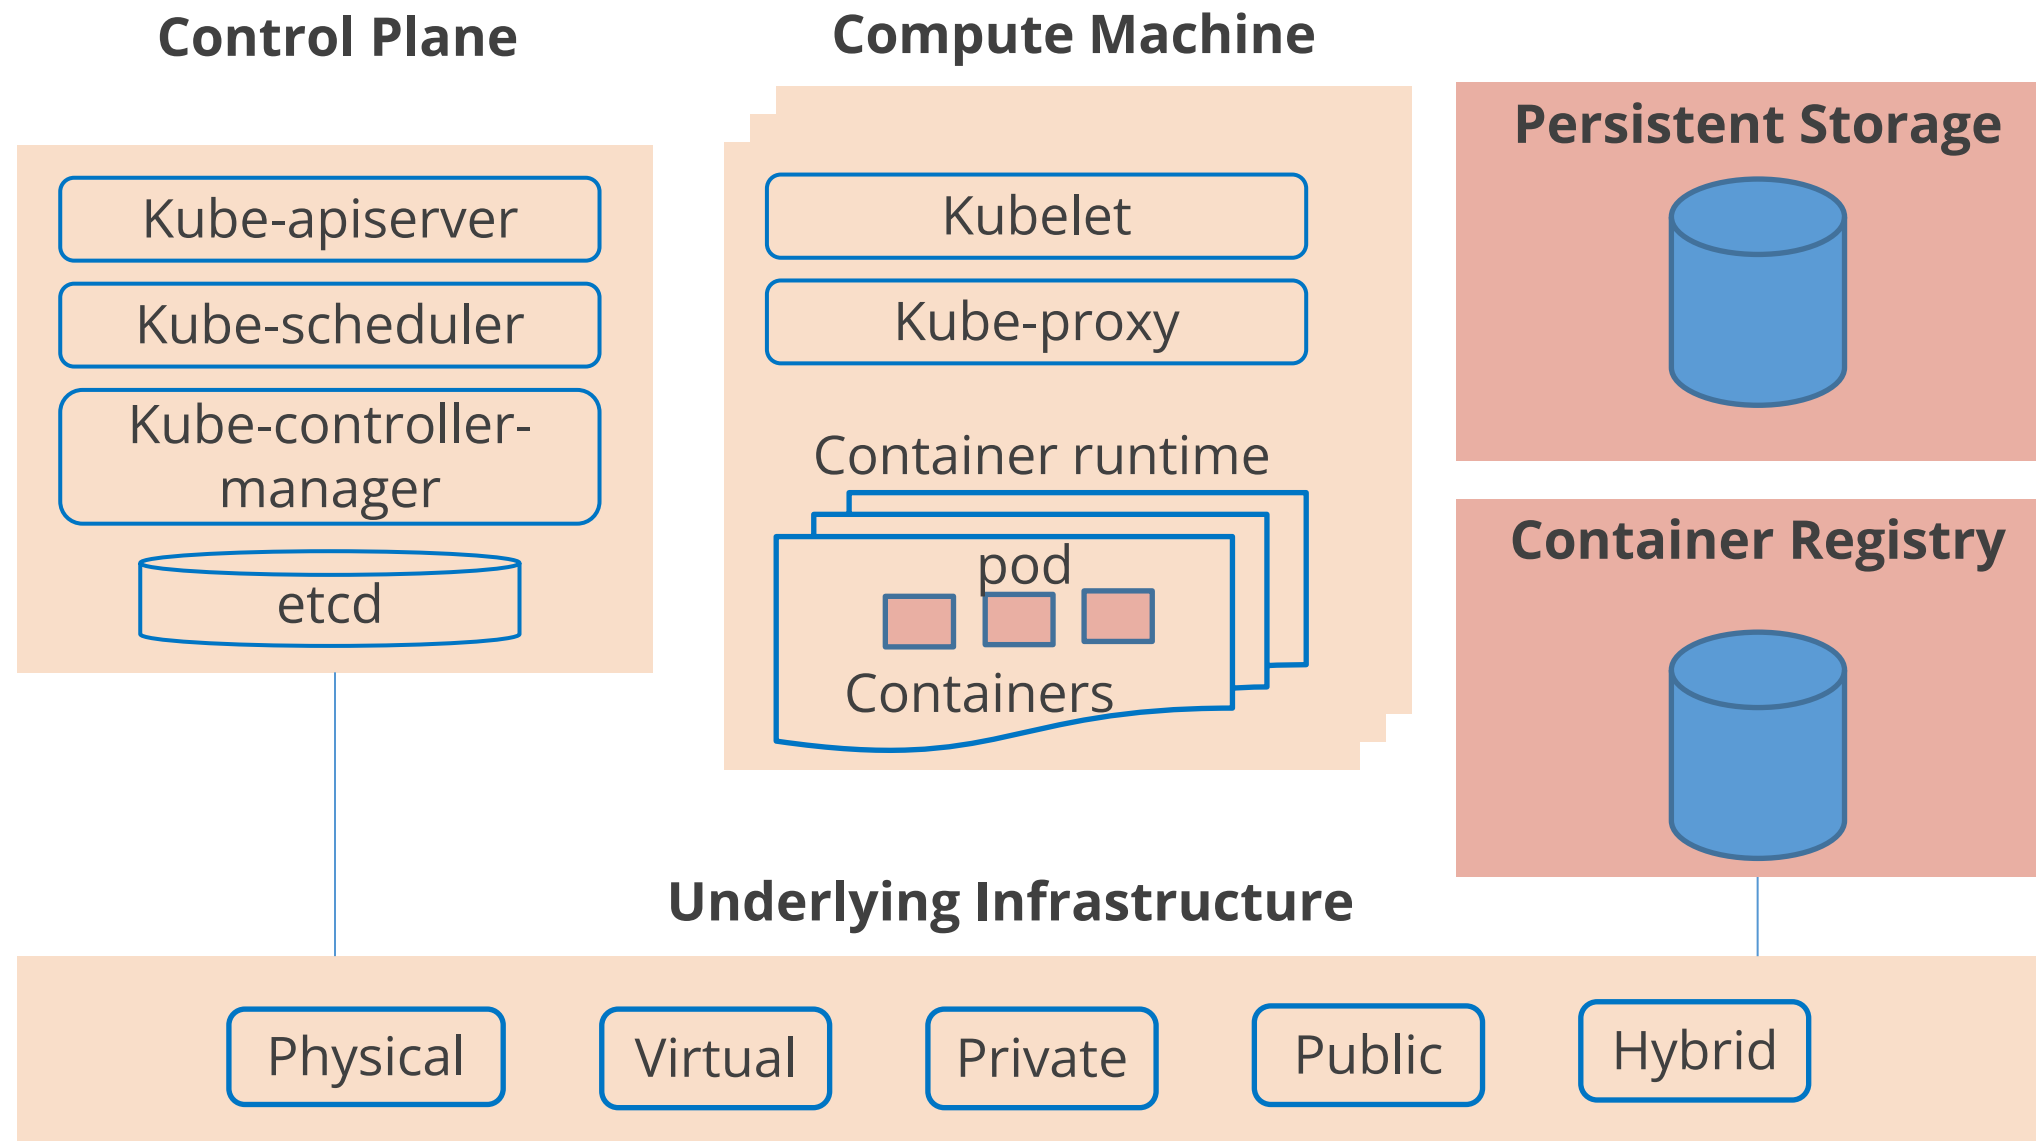
It consists of a master (control plane), a distributed storage system (etcd) for maintaining cluster state consistency, and several cluster nodes.





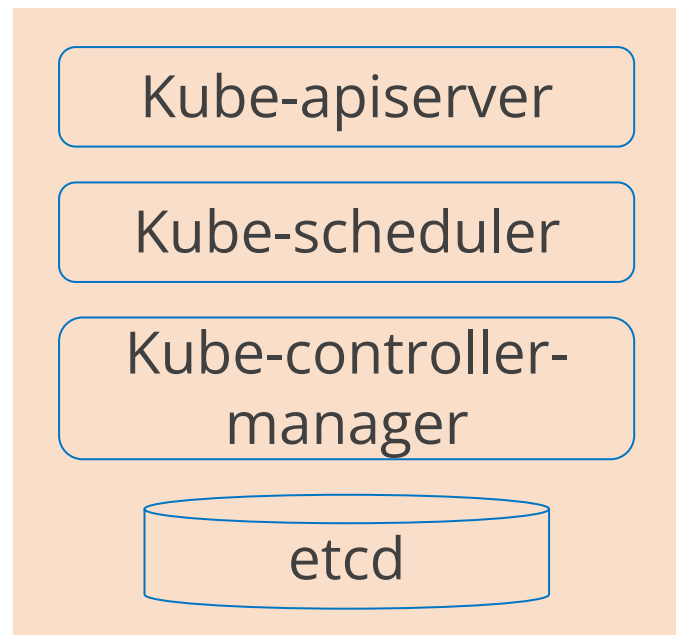
# Kubernetes Cluster

It consists of a master node and a set of worker nodes.



# Control Plane (Master)

It is in constant contact with the compute machines and ensures that the containers are running on the necessary resources.



**Kube-apiserver** handles internal and external requests.


**Kube-scheduler** schedules the pod to an appropriate node.

**Kube-controller-manager** helps run the cluster.

**etcd** stores configuration data and information about the state of cluster lives.

# Compute Machine (Worker)

A Kubernetes cluster requires at least one compute node. pods are scheduled and arranged to run on nodes. The control plane is not schedulable by default because of the default taint value.



**Kubelet** communicates with the control plane.


**Kube-proxy** is a network proxy that facilitates Kubernetes network services.

**Container runtime** manages and runs the components required to run containers.

**Pod** represents a single instance of an application.


# Storage and Registry

## Persistent Storage



**Persistent Storage** manages application data attached to a cluster.

## Container Registry

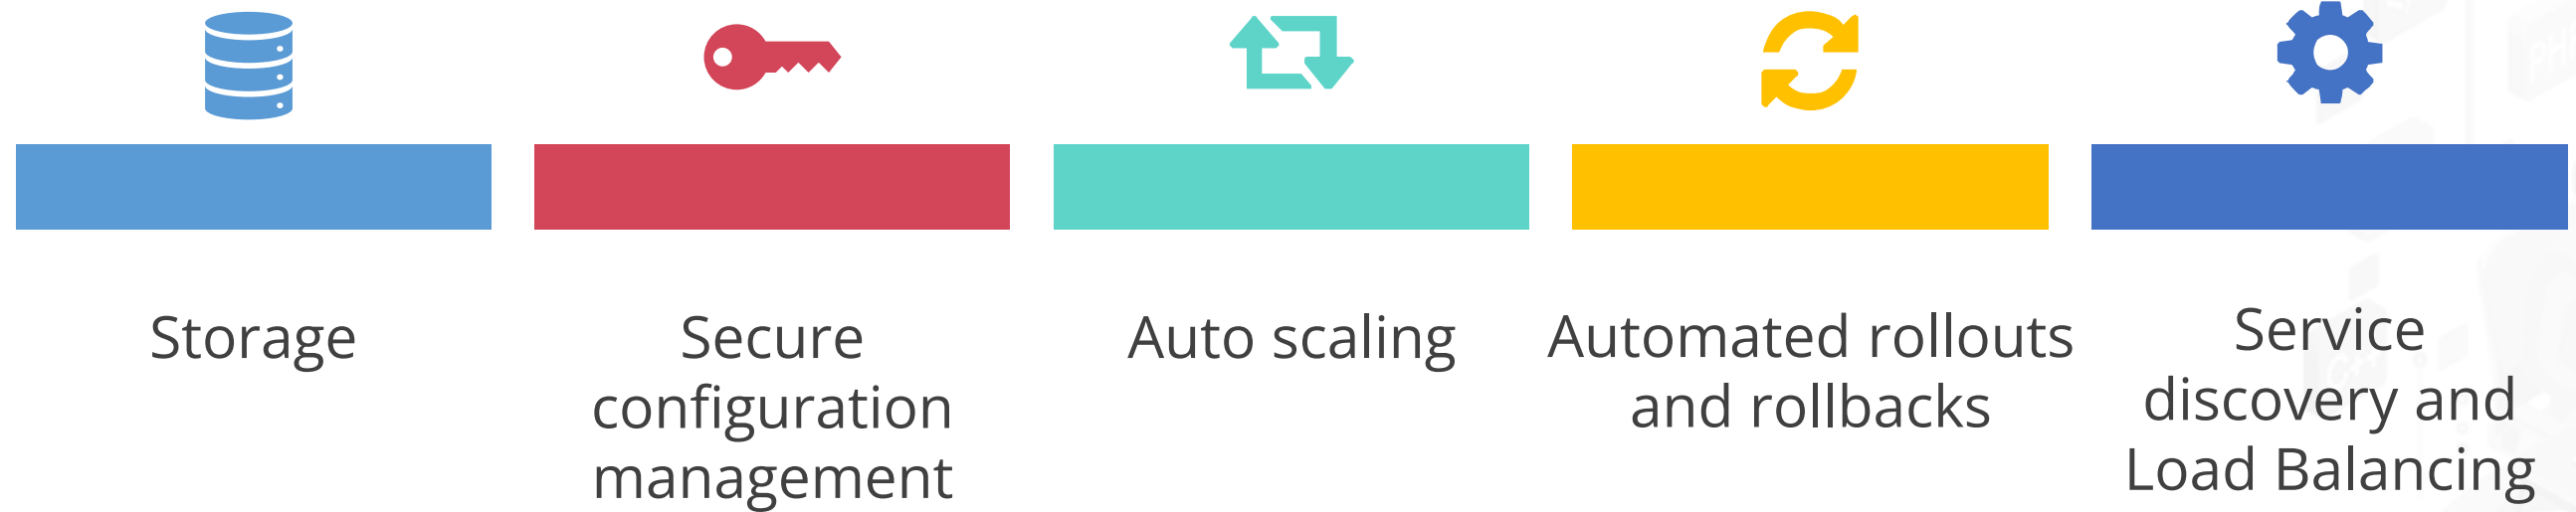


**Container Registry** stores the container images that Kubernetes relies on.



# Features of Kubernetes


Kubernetes offers the following features to its users:



# Container

It is a standard unit of software that aids in the packaging of both application source code and dependencies.

It runs in four modes:




Prepackaging allows the software to run fast and reliably from one computing environment to another. It can run on any compatible infrastructure.

# Container Runtime

It is the software that is responsible for running containers.


Commonly used container runtimes with Kubernetes are:



Container runtime interface (CRI) is an API for container runtime to integrate with Kubelet.



# Docker vs. Containerd

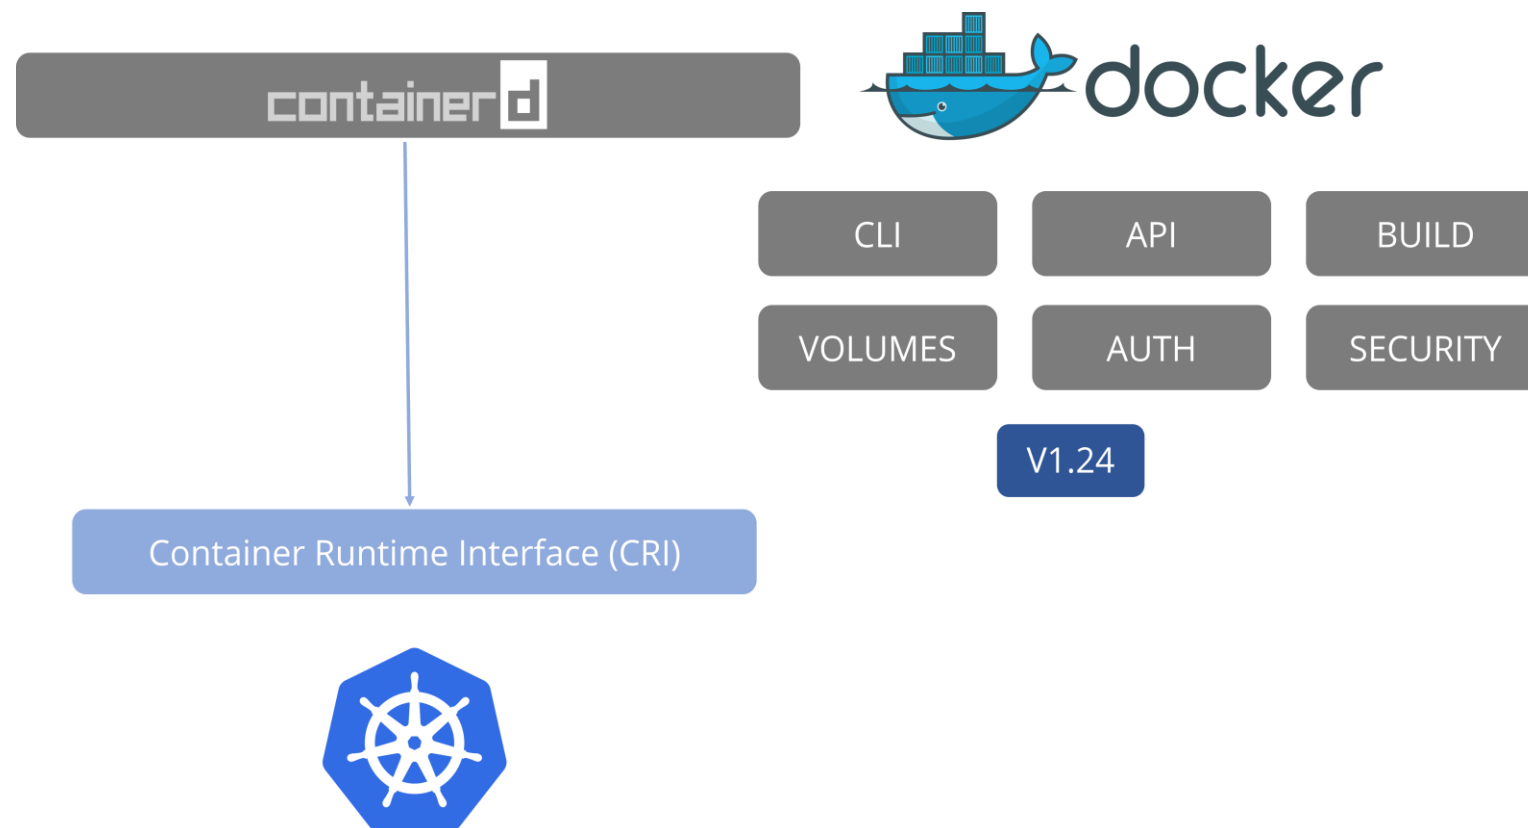


Kubernetes releases before v1.24 included a direct integration with Docker Engine, using a component named dockershim.

That special direct integration is no longer part of Kubernetes (this removal was announced as part of the v1.20 release).



# Docker vs. Containerd

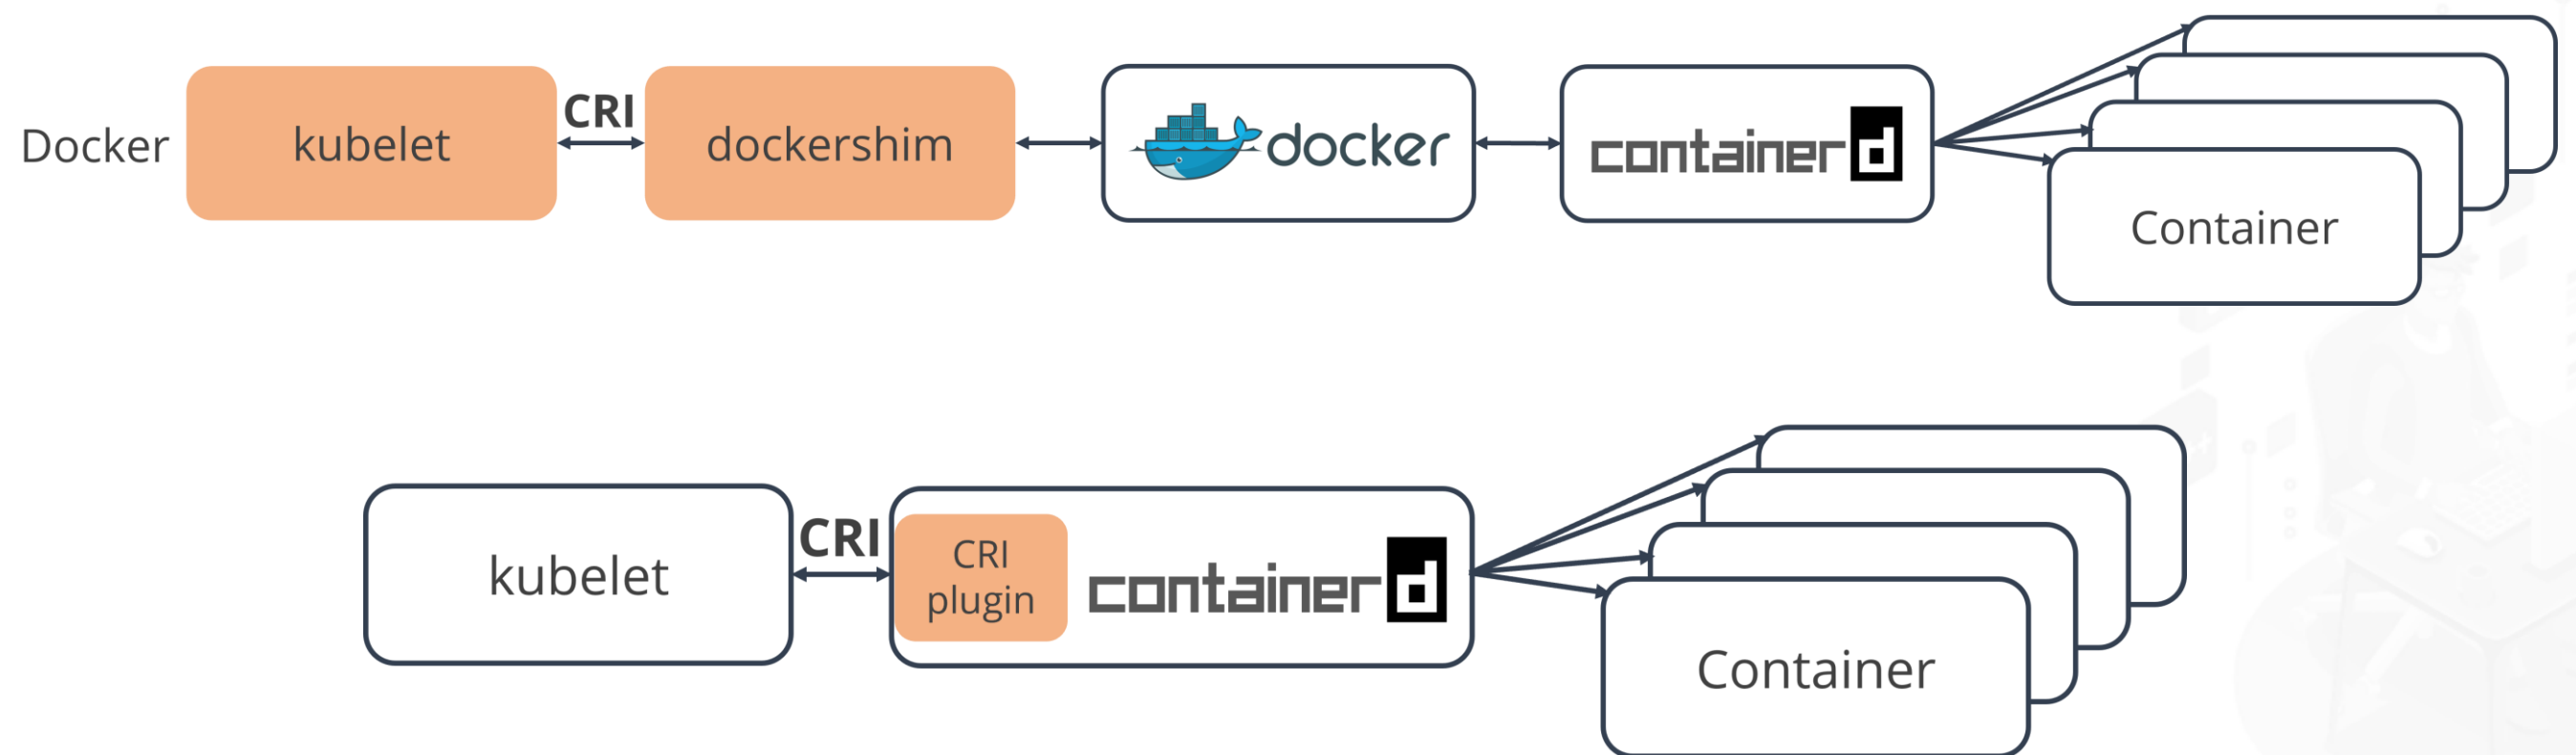


Containerd is CRI-compatible and can directly integrate with Kubernetes, like all other runtimes. Consequently, one can use containerd as a standalone runtime, separate from Docker.

Containerd initially started as an integral part of Docker but was later donated to the Cloud Native Computing Foundation (CNCF).

# Docker vs. Containerd

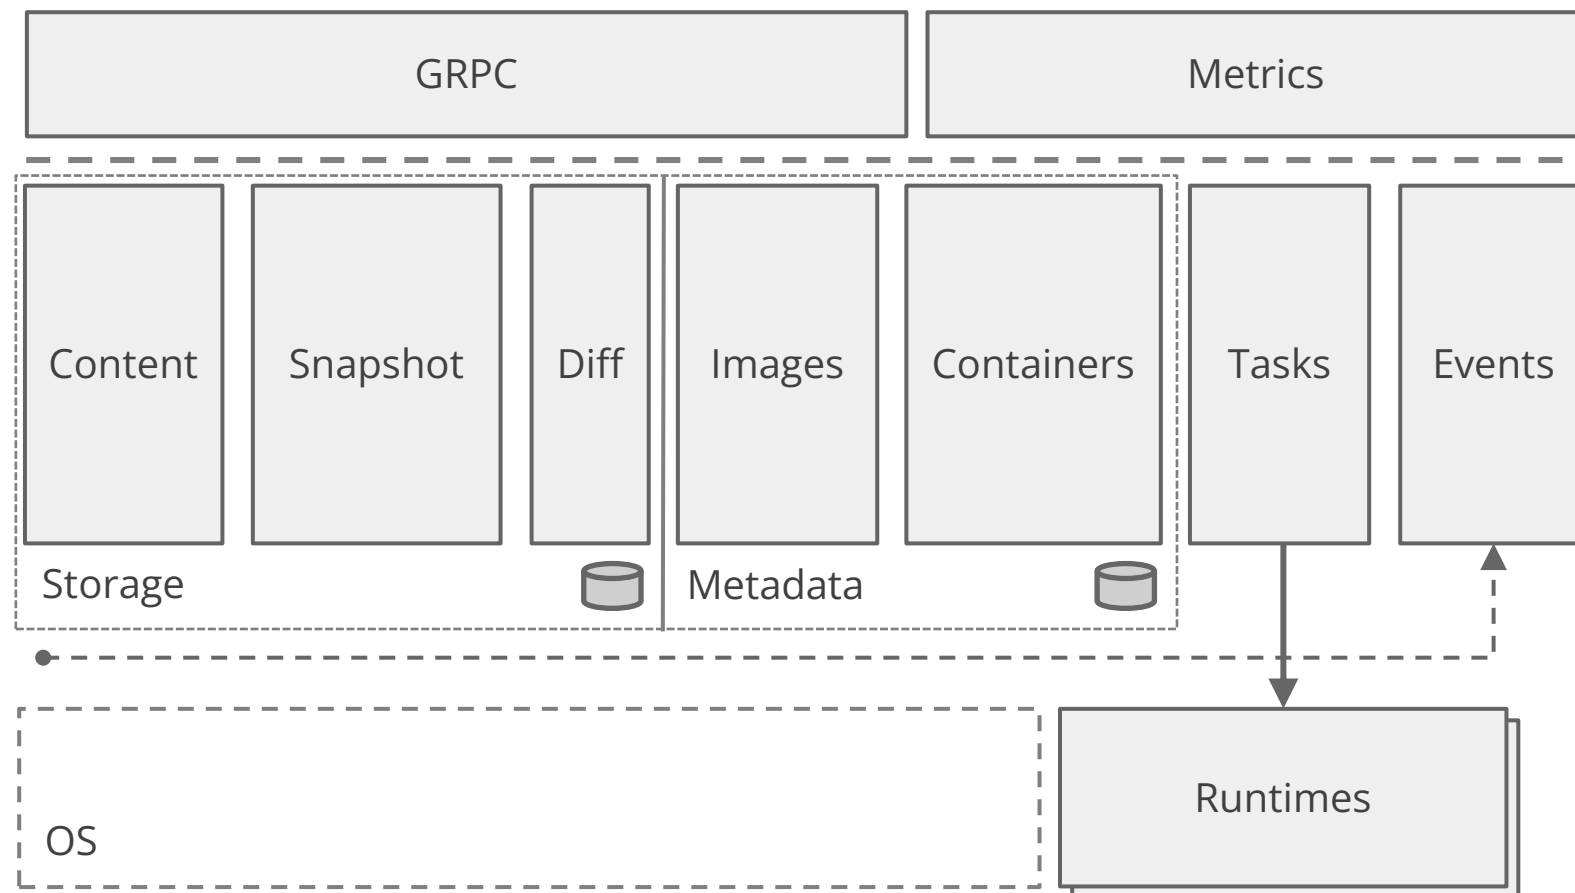
Before Kubernetes v1.24, Kubelet utilized Docker with containerd as an internal component to create containers.



Starting from Kubernetes v1.24, Kubelet directly employs containerd for creating containers.

# Containerd

It is an industry-standard container runtime that prioritizes simplicity, robustness, and portability.



It serves as a Linux daemon capable of managing the entire container lifecycle of its host system.

This includes tasks such as image transfer and storage, container execution and supervision, low-level storage, and network attachments.

# Containerd CLI

*Ctr* and *nerdctl* lack user-friendliness, while *crictl* commands closely resemble Docker commands.



ctr

nerdctl



kubernetes

Container Runtime Interface (CRI)

crictl

Purpose	Debugging	General Purpose	Debugging
Community	ContainerD	ContainerD	Kubernetes
Works With	ContainerD	ContainerD	All CRI Compatible Runtimes

The *crictl* command also possesses pod awareness, allowing you to list pods by executing the *crictl pods* command, a capability Docker did not have.



# Crictl

Crictl is a command-line interface for container runtimes that are compatible with CRI. It allows you to inspect and debug container runtimes and applications on a Kubernetes node.

The following commands are to retrieve debugging information:

docker cli	crictl	Description	Unsupported features
attach	attach	Attach to a running container	--detach-keys, --sig-proxy
exec	exec	Run a command in a running container	--privileged, --user, --detach-keys
images	images	List images	
info	info	Display system-wide information	

# Crictl

docker cli	crictl	Description	Unsupported features
Inspect	inspect	Return low-level information on a container, image or task	
logs	logs	Fetch the logs of a container	--details
ps	ps	List containers	
stats	stats	Display a live stream of container(s) resource usage statistics	Column: NET/BLOCK I/O, PIDs
version	version	Show the runtime (Docker, ContainerD, or others) version information	

# Crictl

The following commands are to perform changes:

docker cli	crictl	Description	Unsupported features
create	create	Create a new container	
kill	stop (timeout = 0)	Kill one or more running container	--signal
pull	pull	Pull an image or a repository from a registry	--all-tags, --disable-content-trust
rm	rm	Remove one or more containers	

# Crictl

docker cli	crictl	Description	Unsupported features
rmi	rmi	Remove one or more images	
run	run	Run a command in a new container	
start	start	Start one or more stopped containers	--detach-keys
stop	stop	Stop one or more running containers	
update	update	Update configuration of one or more containers	--restart, --blkio-weight and some other resource limit not supported by CRI



# Crictl


The following commands are only supported only in *crictl*, not in Docker:

crictl	Description
imagefsinfo	Return image filesystem info
inspectp	Display the status of one or more pods
port-forward	Forward local port to a pod
pods	List pods
runp	Run a new pod
rmp	Remove one or more pods
stopp	Stop one or more running pods



# Container Orchestration

It automates and simplifies the provisioning, deployment, and management of containerized applications.



# Kubernetes: Overview

The aspects of Kubernetes are as follows:

Kubernetes components

Kubernetes objects





Kubernetes API



# Kubernetes Components

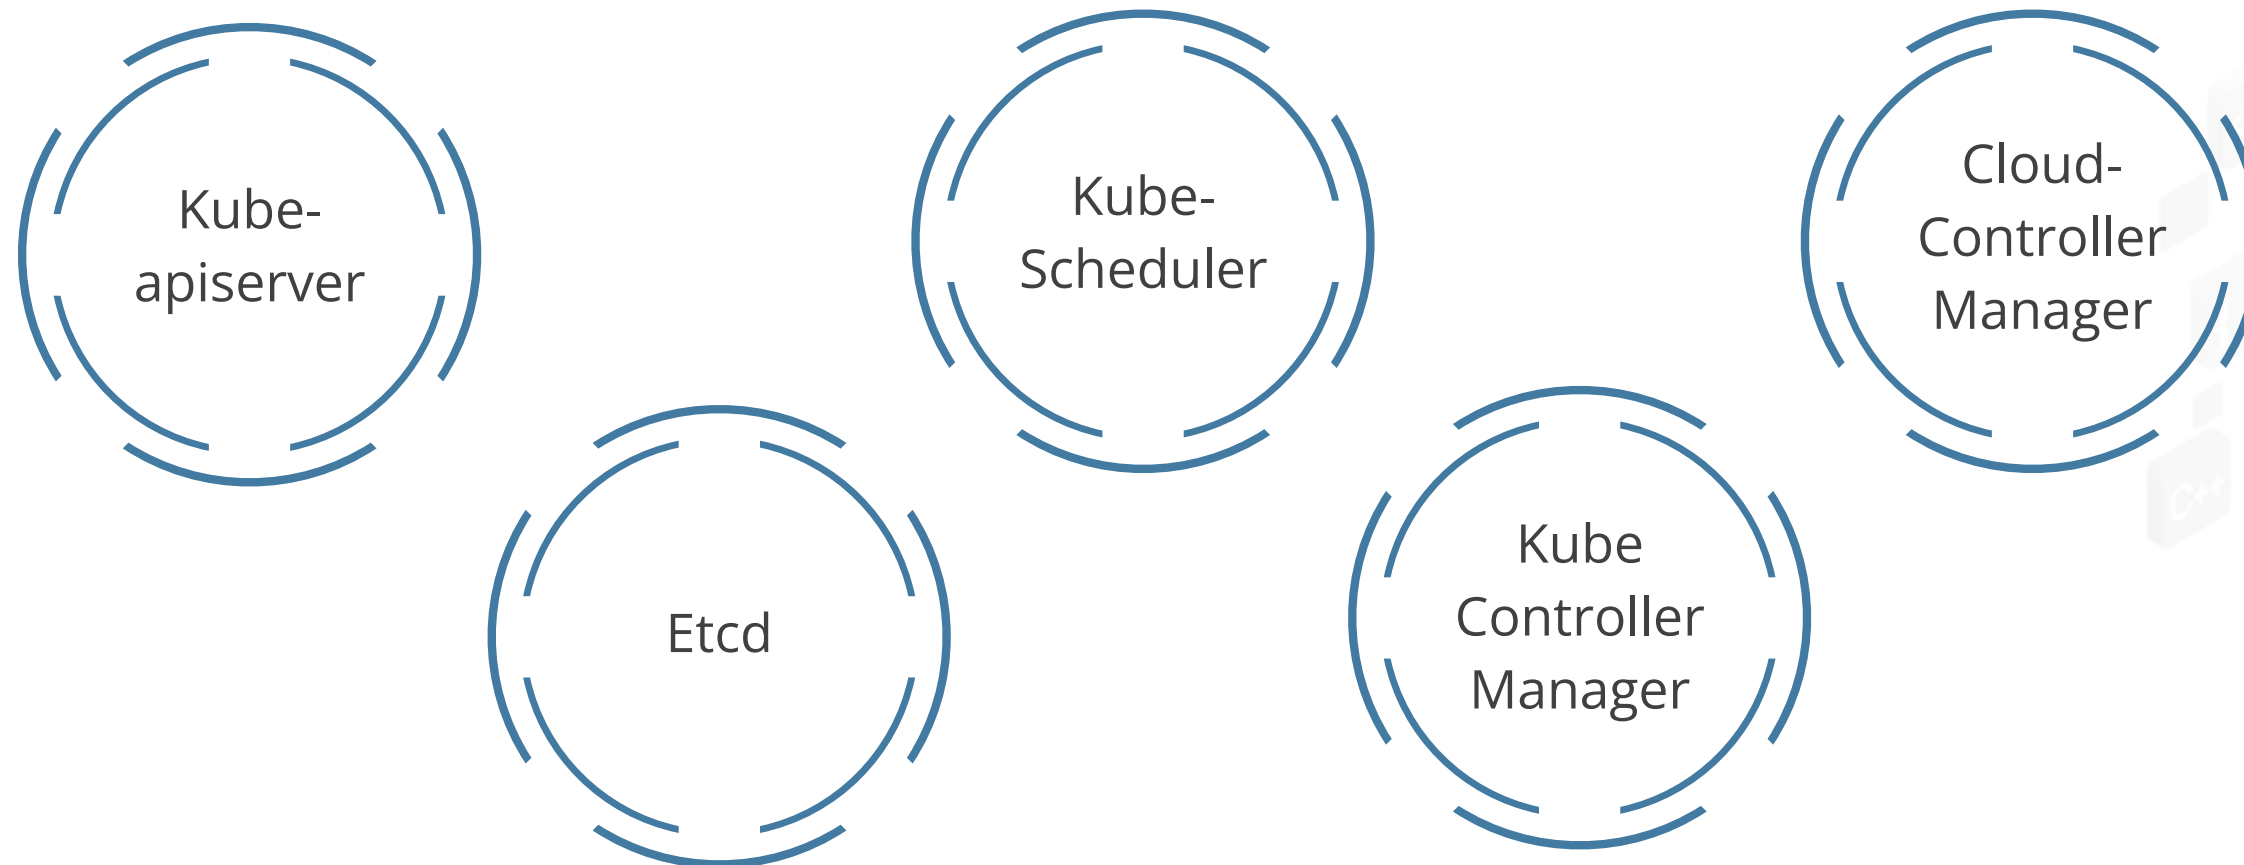
A Kubernetes cluster consists of components that represent the control plane and a set of machines called nodes.

They are divided into:



# Master Components

They are designed to monitor the cluster and respond to cluster events. These components are running as a pod on master in the kube-system namespace.

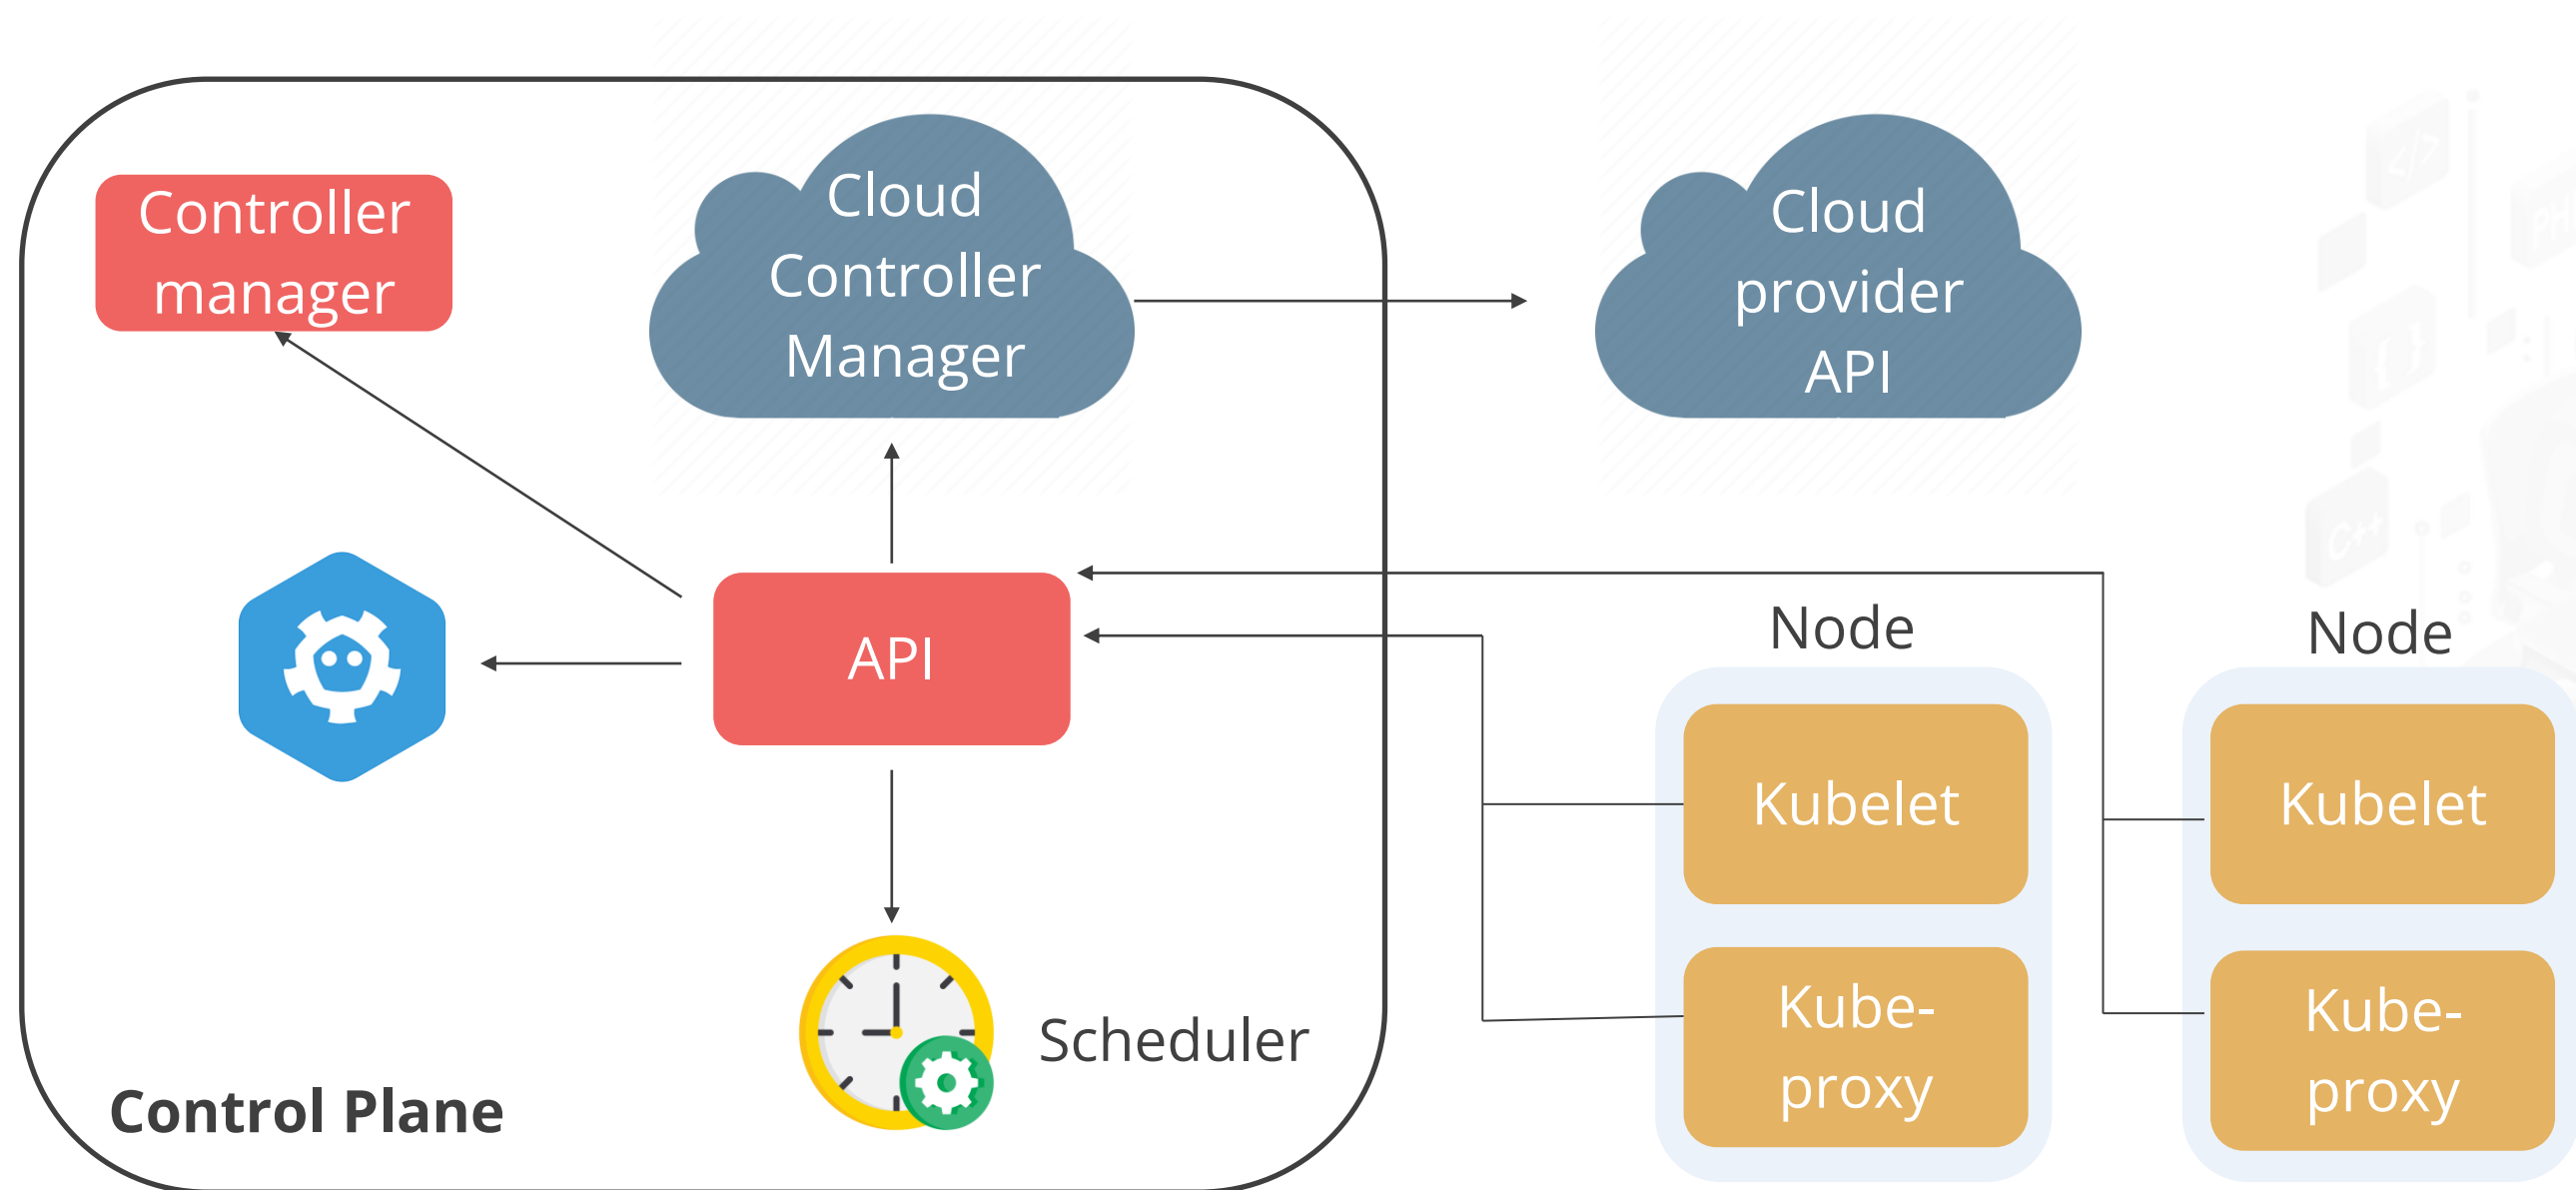


Default config file is **/etc/kubernetes/admin.conf**



# Cloud-Controller Manager

It helps to link clusters into the Cloud provider's API and separates components that interact with the Cloud platform.



# Kubernetes API

It facilitates querying and manipulating the state of objects. The nucleus of the Control Plane in Kubernetes is the API server.

It helps to manipulate the state of API objects, such as:

Pods

Namespaces

ConfigMaps and Events



# Kubernetes Objects

They are persistent entities in the Kubernetes system.

These objects can be described as:



What containerized applications are running



The resources available to those applications

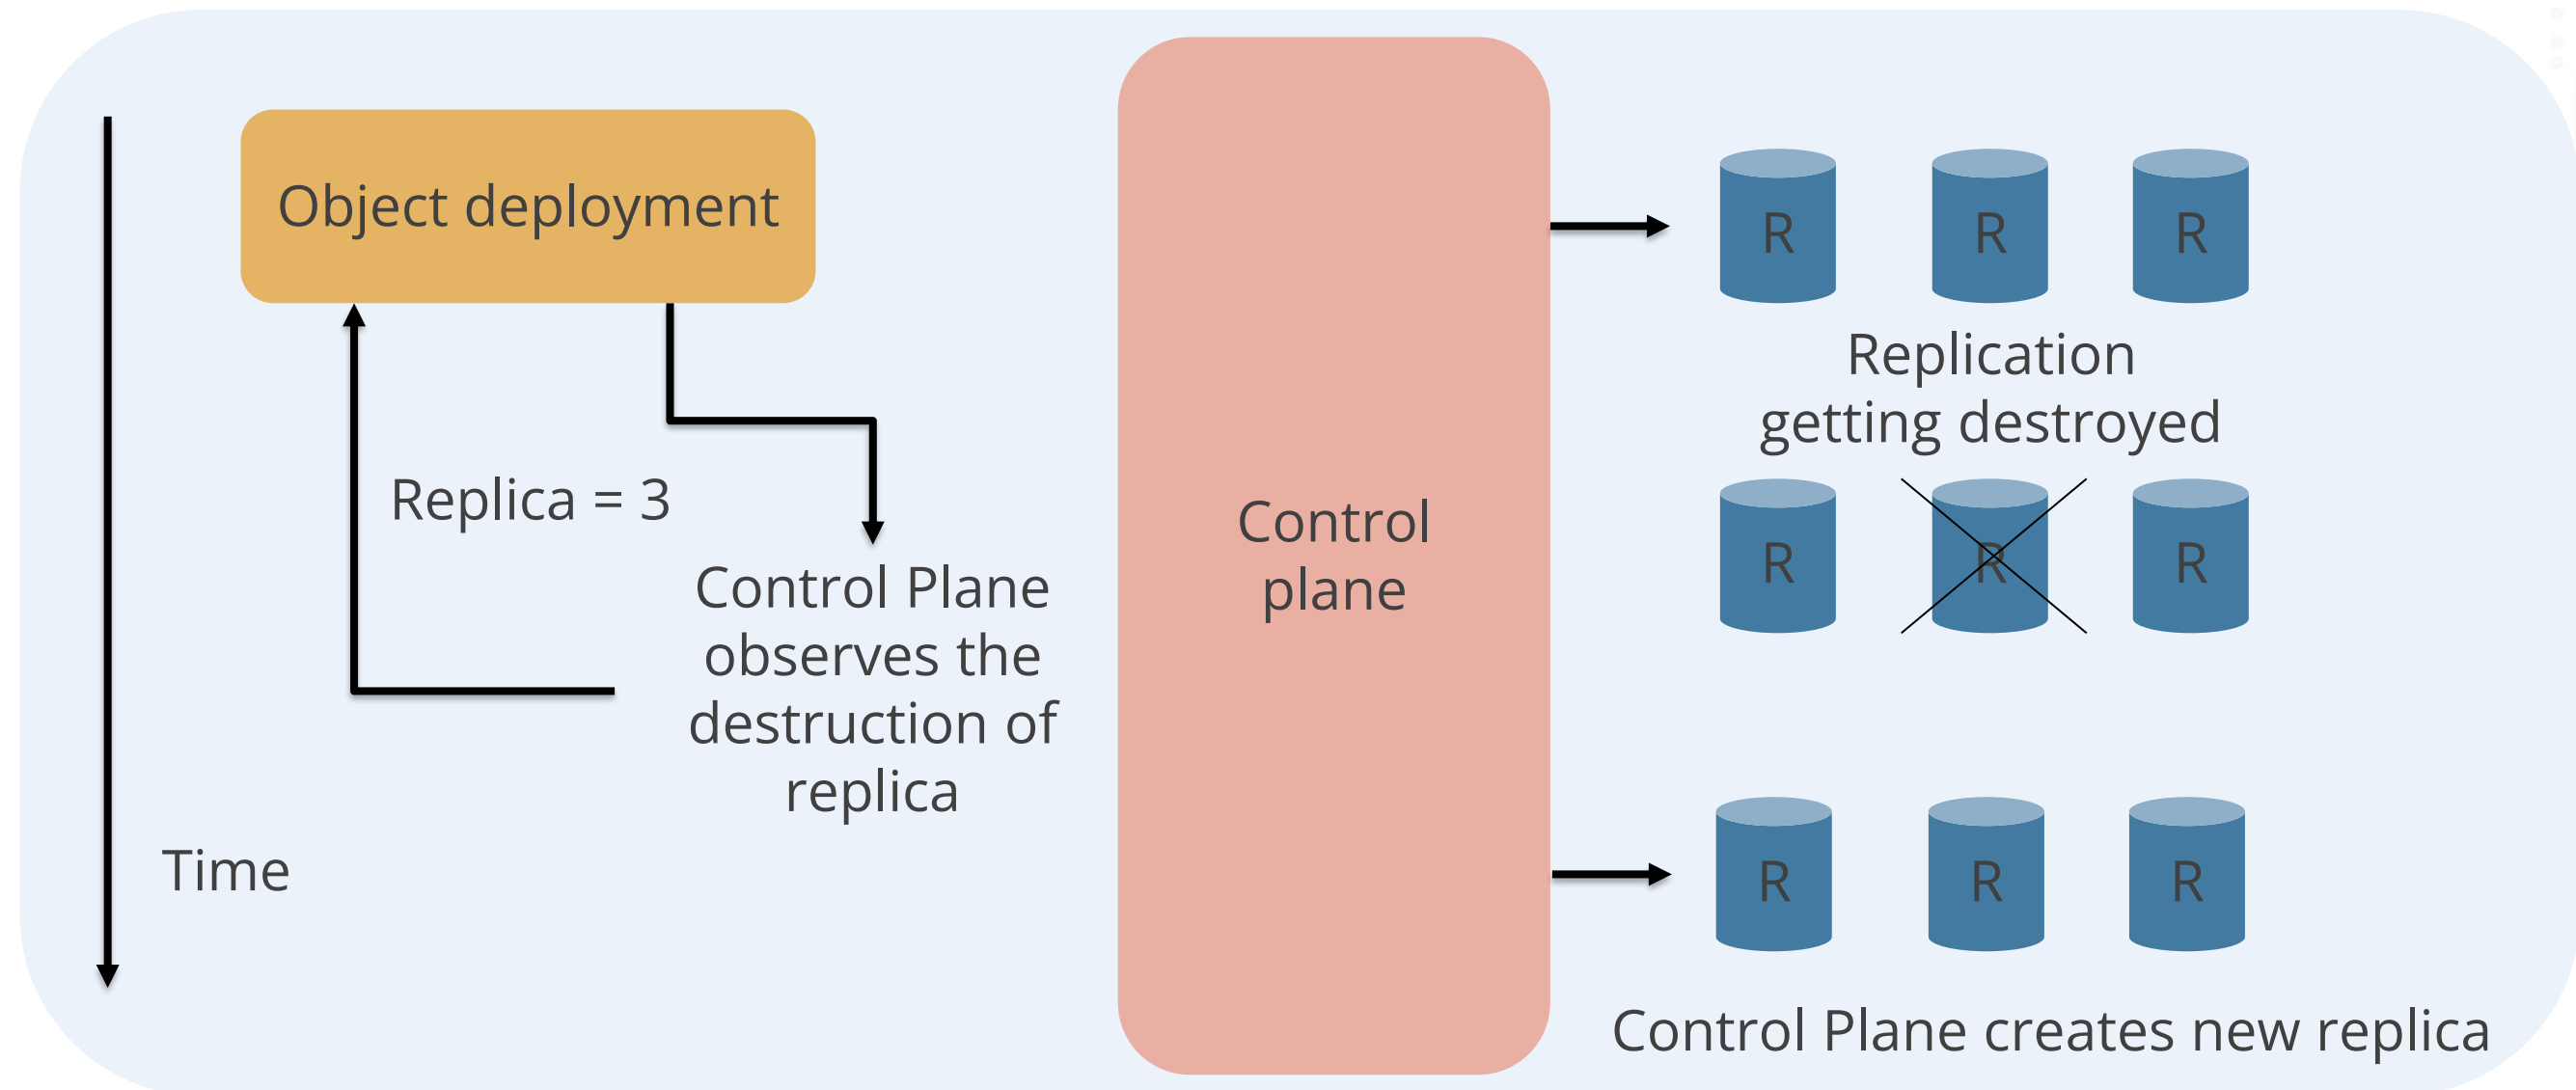


The policies around how applications behave,  
such as restart policies, upgrades, and  
fault-tolerance



# Object Fields

Every Kubernetes object includes two nested object fields that govern the object's configuration, namely, object **spec** and **status**.



# Describing Kubernetes Object


The object spec should be given while creating a Kubernetes object, which outlines the object's desired state as well as some basic information about it, such as a name.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-test-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx-1-17
          image: k8s-master:31320/nginx:1.17 ←
          ports:
            - containerPort: 80
```



# Required Fields in .yaml File

In the .yaml file, a set of values creates a Kubernetes object for the following fields:



# Creating and Configuring a Kubernetes Cluster



**Duration: 20 mins**

## **Problem Statement:**

You have been asked to create a Kubernetes cluster and add the nodes to it.

ASSISTED PRACTICE



# Assisted Practice: Guidelines

Steps to be followed:

1. Change the hostnames of all machines
2. Set up the master node
3. Join the worker nodes in the cluster



# Demonstrating Crictl Commands for Container Runtime Operations



**Duration: 20 mins**

## **Problem Statement:**

You have been asked to debug Kubernetes nodes with crictl CLI commands.

ASSISTED PRACTICE

# Assisted Practice: Guidelines

Steps to be followed:

1. Configure and manage container runtime environment



# TECHNOLOGY

**Etcd**

# Etcd

It is an open-source distributed key-value store for storing and managing important data that distributed systems require to function.


Distributed  
key-value store  
that manages  
critical  
information

Used in production  
environment by  
cloud service  
providers

Used as a  
distributed  
datastore

Used to manage  
configuration data,  
state data, and  
metadata

# Features of Etcd



# Working of Etcd

It is defined through three key concepts (Raft-based system):

Leaders


Handle all client requests that need consensus among clusters

Elections

Is used when the leader is down for any reason

Terms

Appear for multiple candidates





# TECHNOLOGY

## Controller

# Controllers in Kubernetes

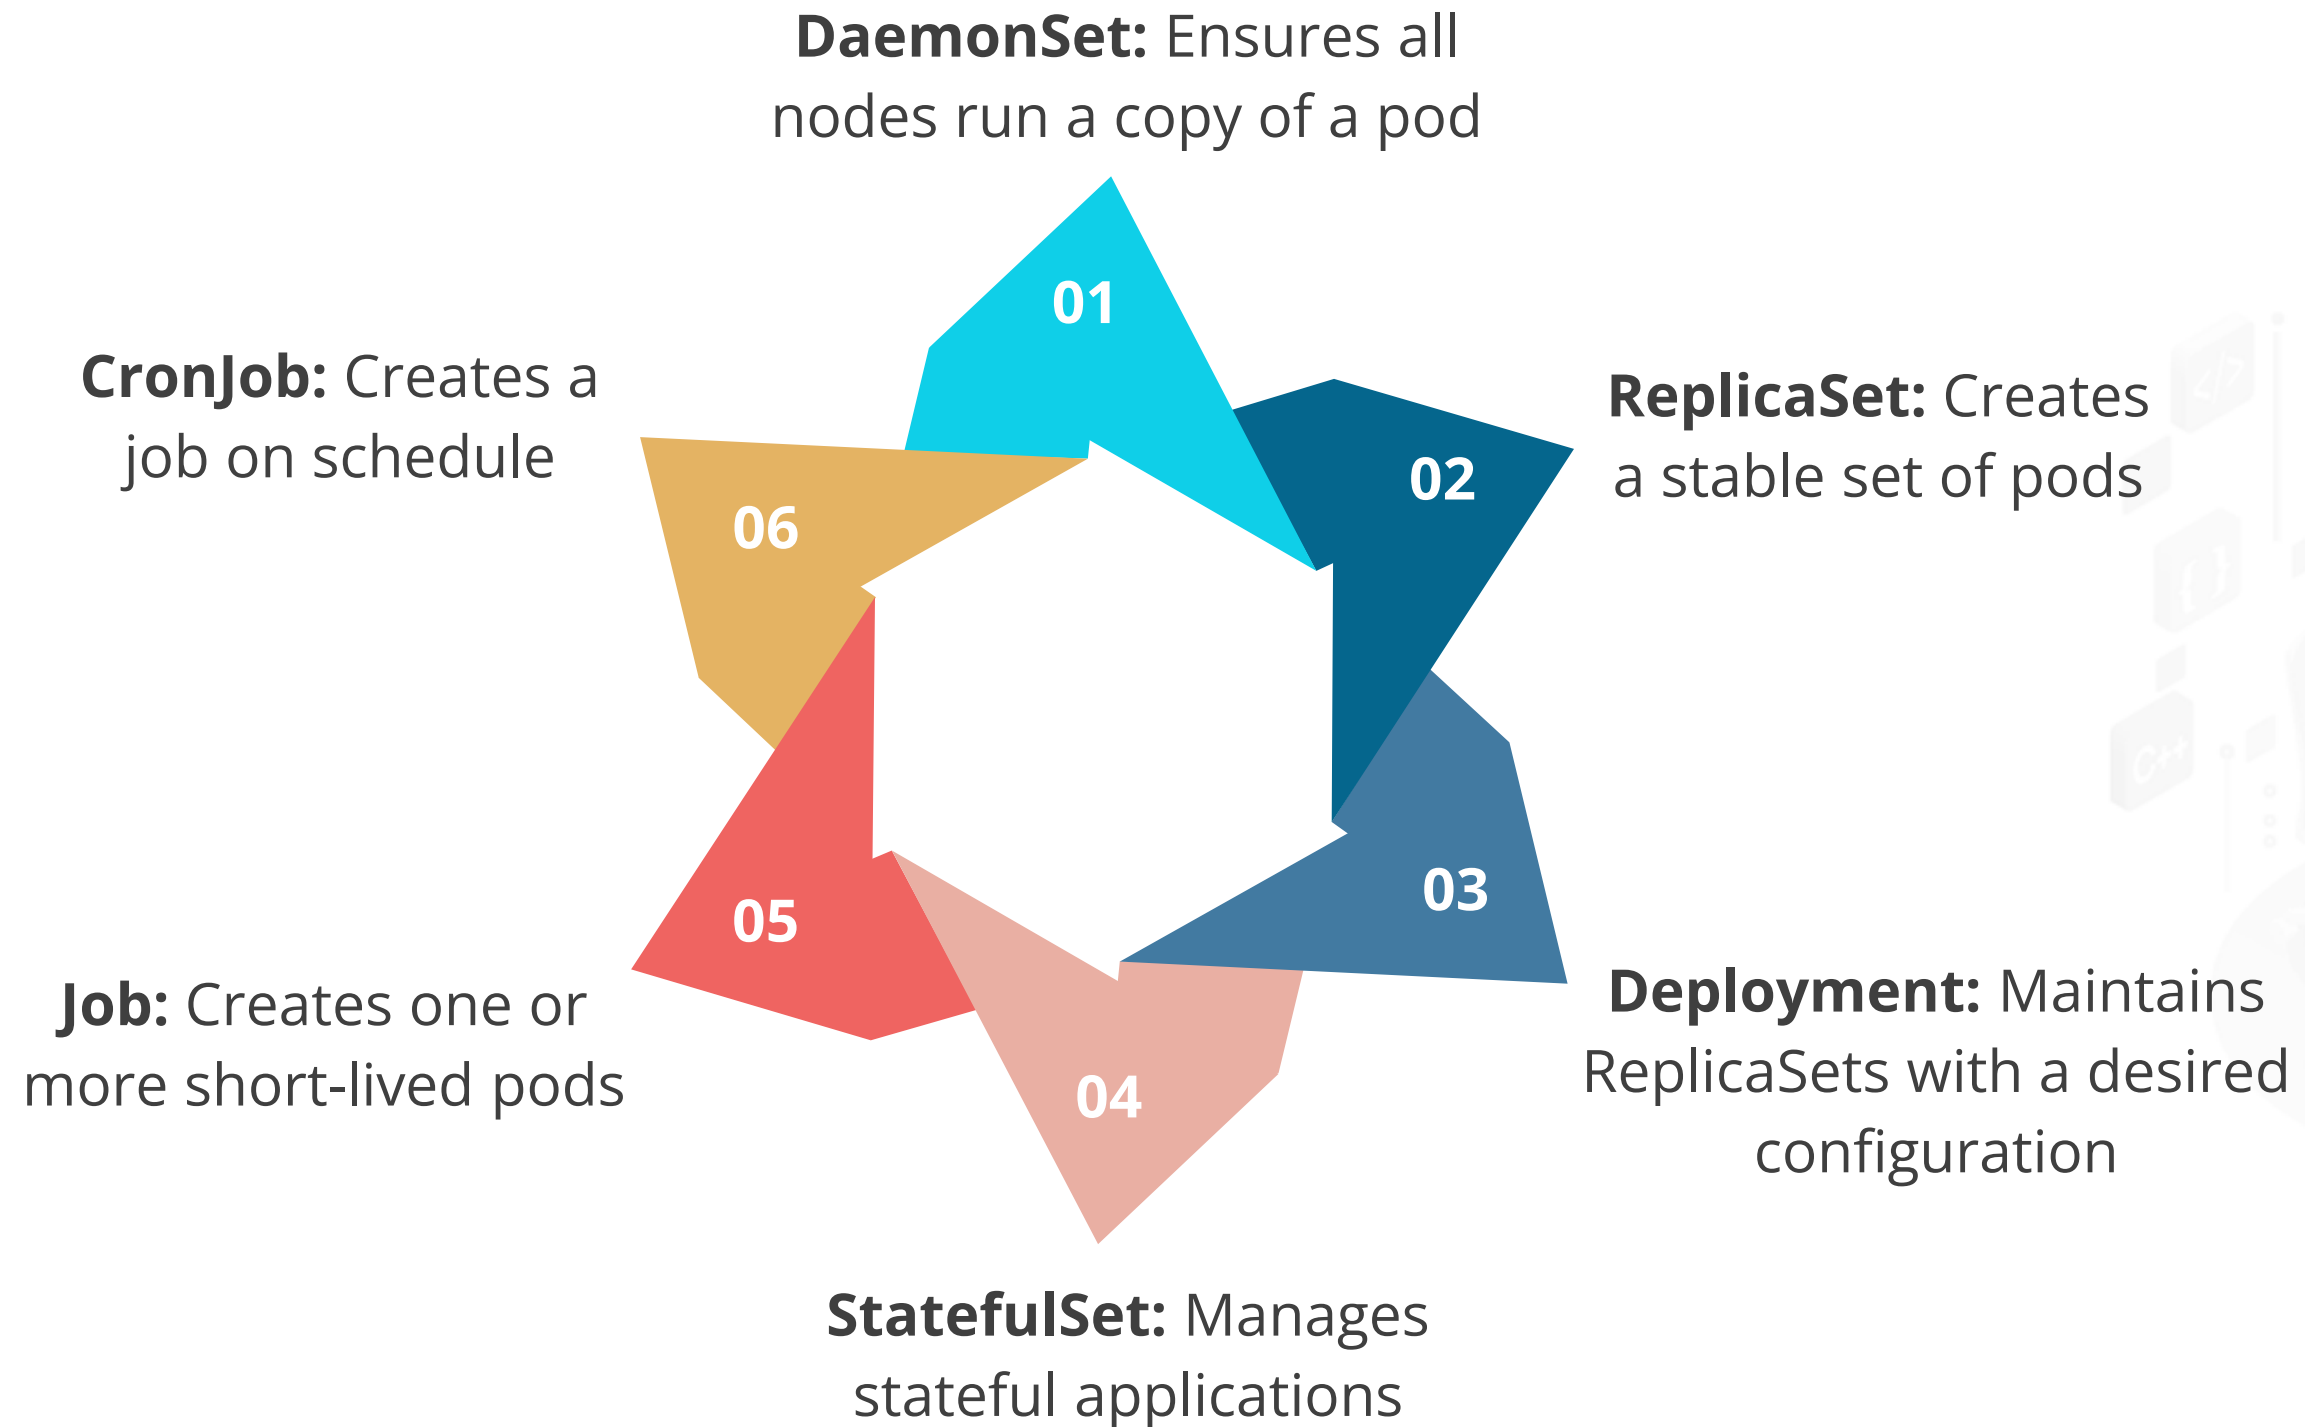
Built-in Controllers manage the state by interacting with the cluster API server.



Job Controller is a good example of a Kubernetes built-in controller.



# Types of Controllers



# Controller Design

Kubernetes uses many controllers, each of which manages a specific feature of the cluster state.



A specific control loop uses one kind of resource as its desired state.




Kubernetes is designed to handle controller failures.



# Working of Controllers

Kubernetes comes with a set of built-in controllers that run inside the Kube-controller-manager. These built-in controllers provide important core behaviors.




To extend Kubernetes, controllers in the system run outside the control plane.

## Scheduler

# Kubernetes Scheduler

Scheduling refers to ensuring that pods are matched to nodes so that a Kubelet can run them. A Scheduler watches for newly created pods that are not assigned to any nodes.





# Kube-scheduler

It is the default scheduler of any Kubernetes system and runs as a part of the control plane.



Helps in writing scheduling components and using them



Finds workable or feasible nodes for a pod and then runs a set of functions to score feasible nodes



Provides optimal node for newly created pods



# Node Selection in Kube-scheduler

The Kube-scheduler selects a node for the pod in a two-step operation:

Filtering

Finds the set of nodes where it is feasible to schedule the pod

Scoring

Ranks the nodes that remain to select the most suitable pod placement



# Configuring Filtering and Scoring Behavior

There are two supported ways to configure the filtering and scoring behavior of a scheduler:

01


Scheduling policies

02


Scheduling profiles



# Node Affinity and Anti-Affinity



Node affinity allows flexible decisions.  
YAML file configuration represents  
specific requirements.



Node anti-affinity is created to allow  
flexible decision-making processes.

## Kubelet

# Kubelet

It is a tiny application that communicates with the control plane. It makes sure that the containers are running in a pod.

It works in terms of a specification called PodSpec.

A PodSpec is a YAML or JSON object that describes a pod.



# Providing Container Manifest to Kubelet

Apart from PodSpec, there are three ways to provide a manifest to Kubelet:

## File

The path passed as a flag on the command line

## HTTP Server

Listens for HTTP and responds to a simple API



## HTTP endpoint

Endpoint passed as a parameter on the command line

## Kube-proxy



# Kube-proxy

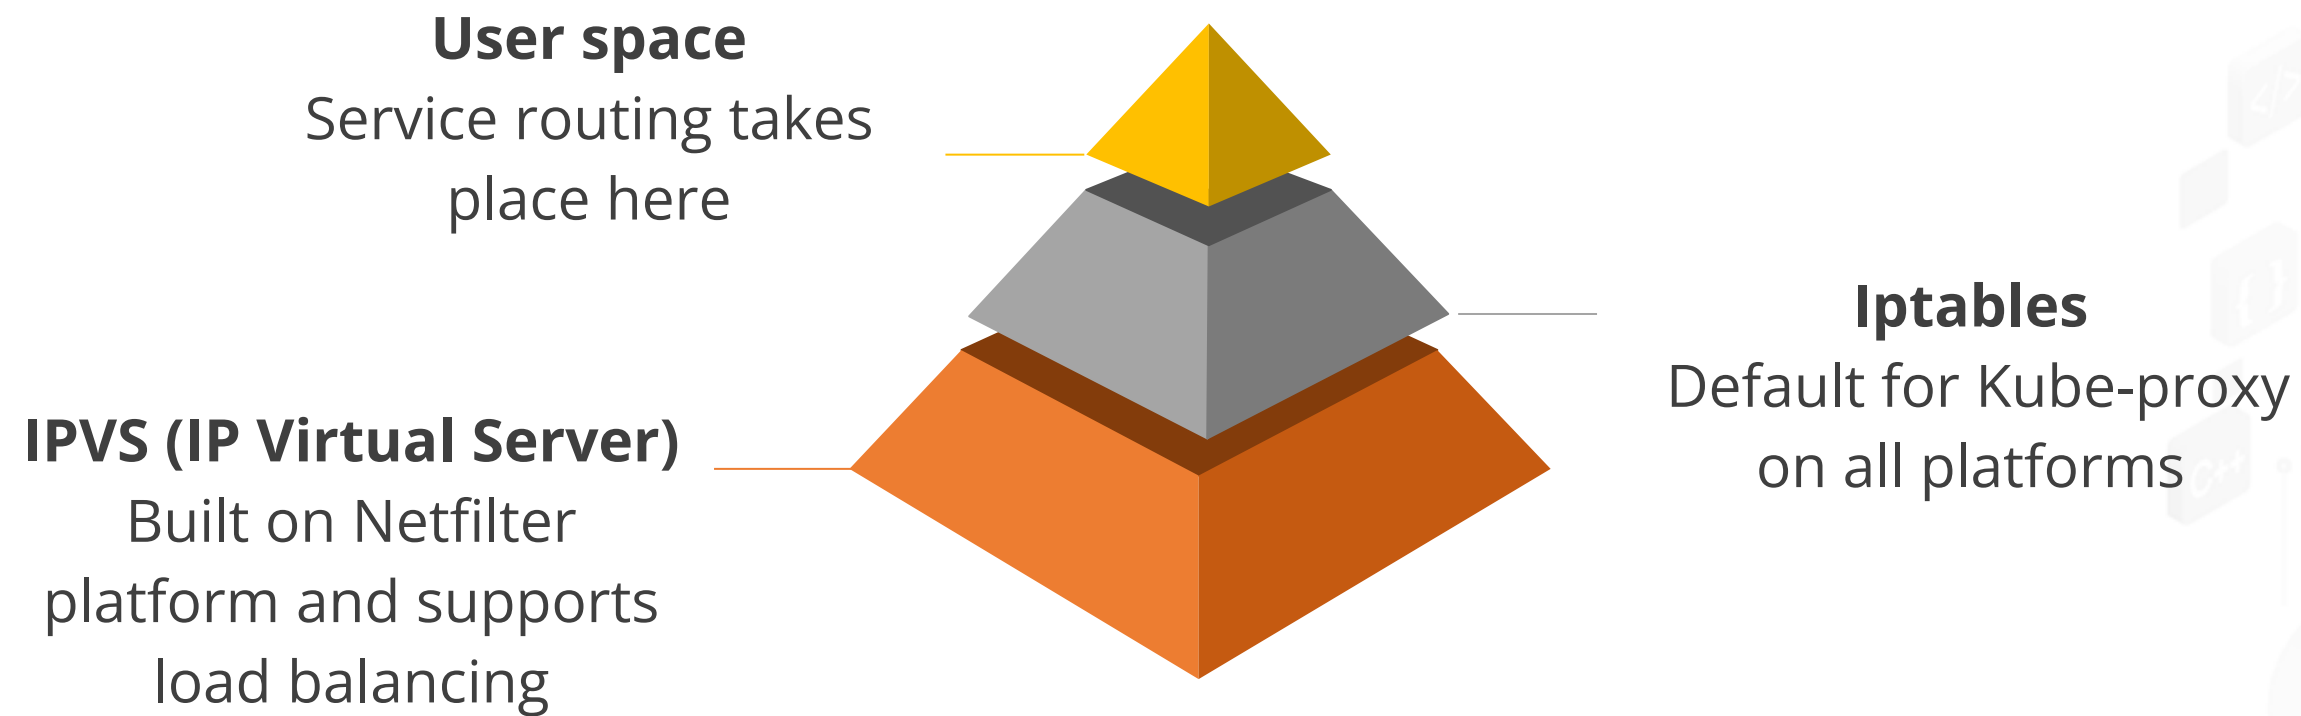
It is a network proxy that runs on every node in a cluster, implementing the Kubernetes Service concept.

Maintains network rules on nodes

Manages forwarding of traffic addressed to the virtual IP addresses of the clusters

# Operation Modes



Each node has a Kube-proxy container process. The Kube-proxy currently supports three different operation modes:



## Containers

# Introduction

A container image is a software package that is ready to use and contains everything needed to run an application:



# Containerized Application

They can be deployed without regard to the underlying infrastructure.




Containerized applications are isolated from each other, like virtual machines, increasing reliability and reducing problems resulting from inter-application interactions.

# Benefits of Containers

Containers consume fewer resources than virtual machines or other outdated program architectures. As they are:




Lightweight



Scalable



Portable and  
consistent



Agile



# Problems Containers Solve

Containers are versatile and solve a broad range of IT problems throughout an application's life cycle.

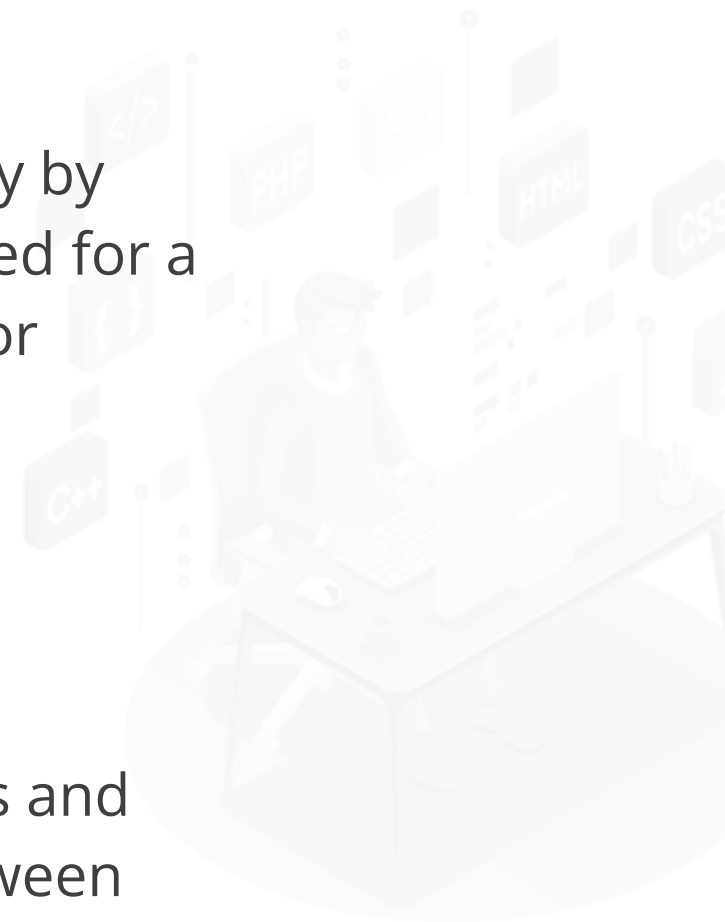
Ensures that software runs properly when moved between computing environments

Facilitates micro-service deployments



Increases efficiency by eliminating the need for a separate hypervisor

Eliminates conflicts and dependencies between multiple applications



# Use Cases for IT Operations

Improve application security by isolating it from other applications

Facilitate seamless migration of applications

Improve IT efficiency by enabling multiple application containers to run on a single OS instance

Offer on-demand scalability





# How Do Containers Work?

Containers isolate applications from one another.

A registry or repository transfers container images and the application container engine which turns the images into executable code.

Container repositories facilitate reusing commonly used container images.

Containers are created using the process of packaging applications.


# Containers: Features

**Namespaces** provide access to the underlying operating system.

**Union File Systems** prevent data duplication.



**Control Groups** implement resource accounting and resource limitation.



# TECHNOLOGY

## Pods

# Understanding Pods

Pods are the smallest deployable unit of computing, which can be created and managed in Kubernetes.



Pods in a Kubernetes Cluster are mainly used in two ways:

- 1 Pods that run a single container; the most common Kubernetes use case is the "one-container-per-pod" model.
- 2 Pods that run multiple containers, which should work in conjunction with each other.




# Understanding Pods

Pods also contain shared networking and storage resources for their containers:



# How Pods Manage Multiple Containers?


Pods are designed to support multiple cooperating processes (as containers) that form a cohesive unit of service.




# Pods and Controllers

Workload resources create and manage one or more pods.


Examples of workload resources:



Deployment



StatefulSet



DemonSet



# Pod Template

Pod templates are specifications for creating pods and are included in workload resources, such as deployments, Jobs, and DaemonSets.

Sample pod template:


```
apiVersion: v1
kind: pod
metadata:
  name: nginx
spec:
  containers:
  - name: nginx
    image: nginx:1.14.2
    ports:
    - containerPort: 80
```





# Pod Update and Replacement

The controller does not update or patch existing pods when the pod template for a workload resource is updated or changed.



The controller creates new pods based on the updated template.

# Pod Update and Replacement

Pod update operations like **patch** and **replace** have some limitations:

The metadata about a pod is immutable.

If the **metadata.deletionTimestamp** is set, no new entry can be added to the **metadata.finalizers** list.

Pod updates may not change fields.

When updating the **spec.activeDeadline** seconds field, two types of updates are allowed.

# Resource Sharing and Communication

Pods enable data sharing and communication among their constituent containers using two methods:

**Storage in pods:** All containers in a pod have access to the shared volumes, allowing those containers to share data.

**Pod Networking:** When containers within a pod communicate with entities outside the pod, they must coordinate how they use the shared network resources.

# Privileged Mode for Containers

Any container in a pod can get the Privileged Mode into working by utilizing the privileged flag on the security context of the container spec.

Privileged Mode is used for containers that use the operating system's administrative capabilities.

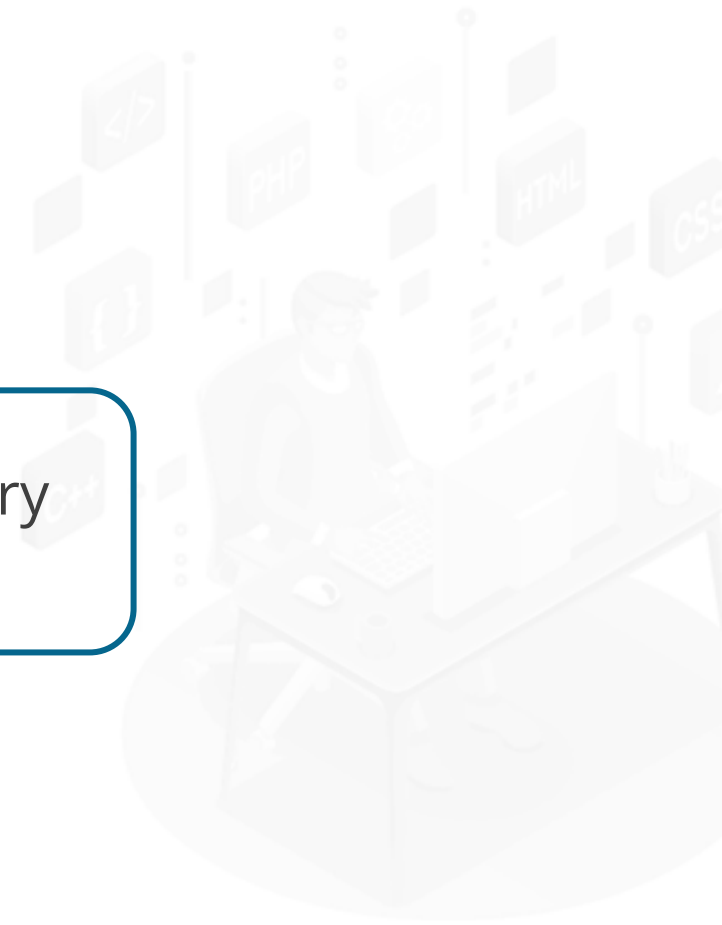
The processes in Privileged Mode have the same privileges as the processes outside a container.

# Static Pods

The Kubelet daemon manages static pods directly on a specific node, without being observed by the API server.



The control plane manages most pods. The Kubelet supervises every static pod directly.



# Configuring Pods in the Kubernetes Cluster



**Duration: 15 mins**

## **Problem Statement:**

You have been assigned a task to create, configure pods, and execute the Apache services.

ASSISTED PRACTICE

# Assisted Practice: Guidelines

Steps to be followed:

1. Configure and set up the pod files
2. Configure and set up the Service file
3. Execute the Apache services



## ReplicaSets




# ReplicaSets

It maintain a stable set of replica pods running at any given time.



# Operators to Use with ReplicaSets

There are three important operators that play a crucial role in managing and configuring ReplicaSets within a Kubernetes cluster:



Ensure that the selectors of one ReplicaSet do not match another's.

# ReplicaSet Manifest



The following is an example of a ReplicaSet manifest:

```
apiVersion: apps/v1 # our API version
kind: ReplicaSet    # The kind we are creating
Metadata: # Specify all Metadata like name, labels
  name: some-name
  labels:
    app: some-App
    tier: some-Tier
Spec:
  replicas: 3 # Here is where we tell k8s how many replicas we want
  Selector: # This is our label selector field.
    matchLabels:
      tier: some-Tier
    matchExpressions:
      - {key: tier, operator: In, values: [some-Tier]} # we are using the set-based
operators
  template:
    metadata:
      labels:
        app: some-App
        tier: someTier
    Spec: # This spec section should look like spec in a pod definition
      Containers:
```



# Working of ReplicaSet

A ReplicaSet is defined with fields, including a selector that specifies:



It ensures that a specified number of pod replicas are running at any given time.

## Deployments


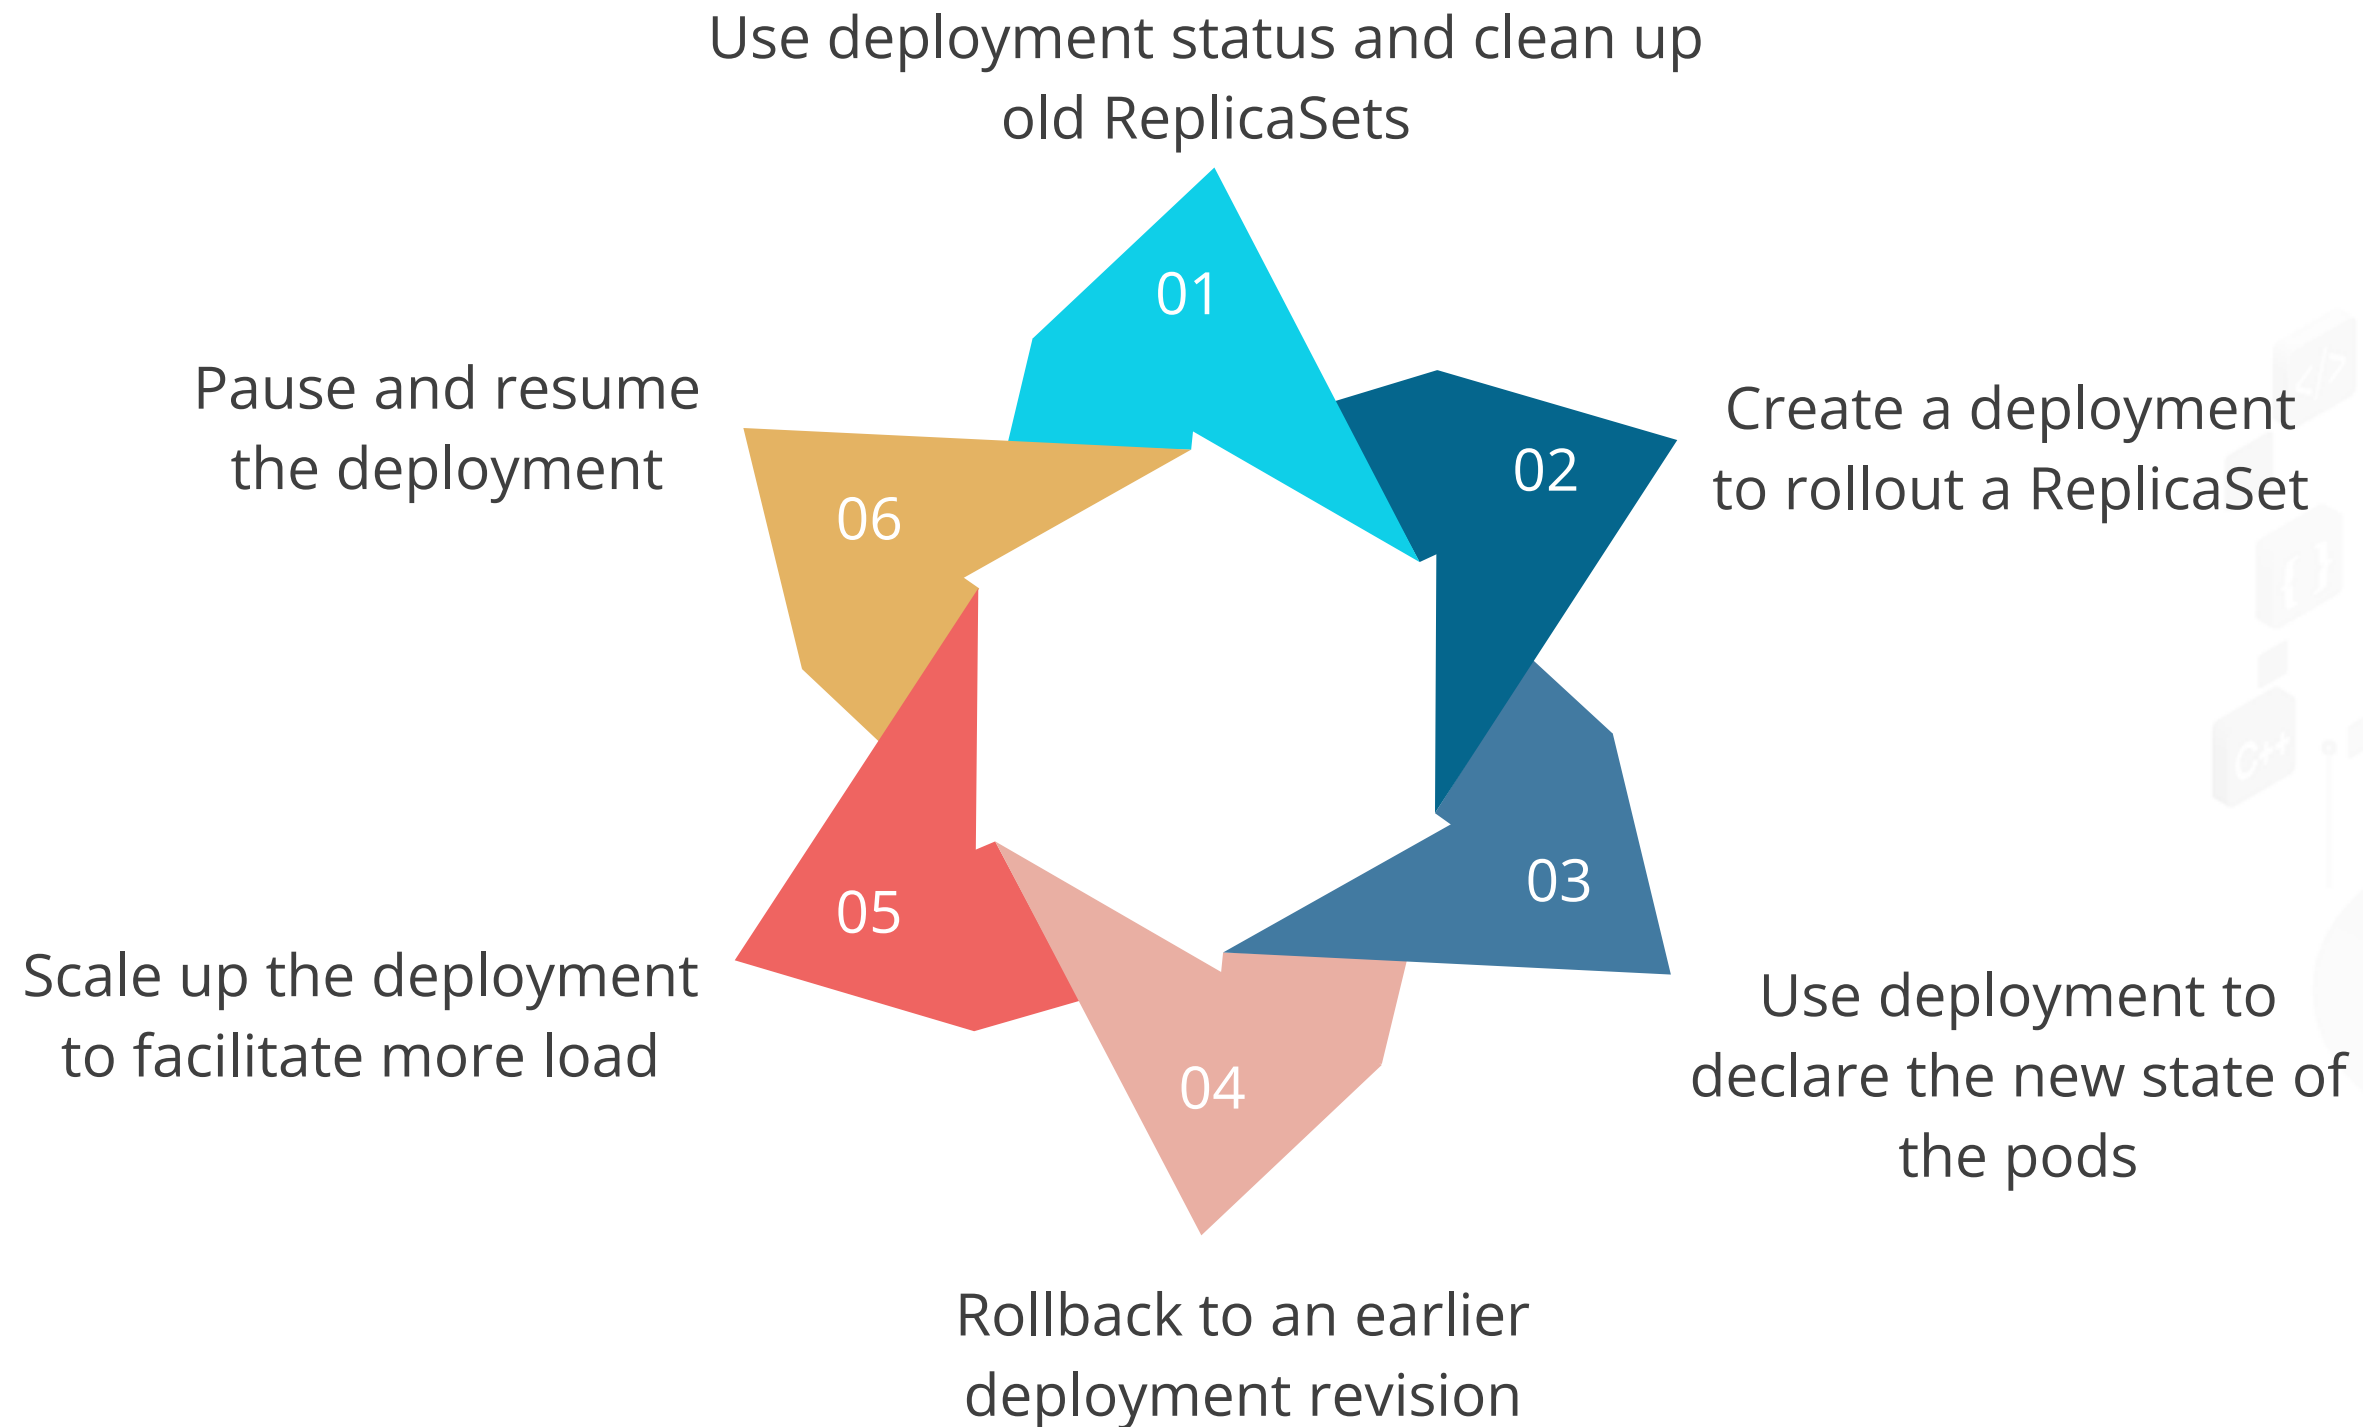
# Deployment

It provides declarative updates for pods and ReplicaSets.



It can be defined to create new ReplicaSets or remove existing Deployments.

# Use Cases of Deployment



# Creating a Deployment

The following is an example of a Deployment:

## Example

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:1.14.2
        ports:
        - containerPort: 80
```





# Updating and Rolling Back Deployment

Deployments can be updated by making changes to the pod template spec in the Deployment; it automatically generates an update rollout.

## Rolling Back Deployment

```
kubectl rollout undo [deployment_name]
```

```
#Adding the argument
```

```
-to-revision=
```

```
#will roll back to that specific  
version of the deployment
```



# Scaling a Deployment

Deployments are useful for scaling the number of replicas as demand increases for a particular application.

## Example:

```
# to scale a deployment up to 20 replicas  
  
kubectl scale [deployment-name] -replicas 20
```



# Pause and Resume

Multiple fixes can be applied between pausing and resuming without triggering unnecessary rollouts.

## Example:

```
#Pause a deployment
```

```
kubectl rollout pause deployment.v1.apps/nginx-deployment
```

```
#Resuming a deployment
```

```
kubectl rollout resume deployment.v1.apps/nginx-deployment
```



# Creating and Configuring the Deployment



**Duration: 20 mins**

## **Problem Statement:**

You have been assigned a task to create and configure deployment for an application.

ASSISTED PRACTICE

# Assisted Practice: Guidelines

Steps to be followed:

1. Create the deployment
2. Access the pod



## Services, Load Balancing, and Networking

# Services, Load Balancing, and Networking

Services and Load Balancing are the most important parts of Kubernetes networking, and they address four main concerns.

Containers in a pod use networking to communicate via loopback.

Cluster networking facilitates communication between various pods.


The Service resource lets exposing an application running in pods to be accessible from outside the cluster.

Services are used to publish services meant for consumption inside the cluster only.



# Kubernetes Pod Network

A Kubernetes pod network connects several interrelated components including:








# Networking in Kubernetes

Traffic that flows between nodes can also flow to and from nodes and an external physical machine or a VM.

There are four ways of getting external traffic into a Kubernetes cluster:



# Using Basic Commands of Kubernetes



**Duration: 15 mins**

## **Problem Statement:**

You have been asked to execute the basic commands used in Kubernetes.

ASSISTED PRACTICE

# Assisted Practice: Guidelines

Steps to be followed:

1. Create the deployment
2. Create the namespaces
3. Scale and delete the deployment



# TECHNOLOGY

## Policies

# Overview

Policies define what end users can do on the cluster and possible ways to ensure that clusters comply.

Policies are applicable to network, volume, resource usage, resource consumption, access control, and security.

A constraint is a declaration that expects a system to meet a set of requirements.

Policy enablement helps organizations take control of Kubernetes operations.



# Key Benefits of Policies



Simplified  
operations



Ease of policy  
enforcement



Automated discovery of  
violations and conflicts

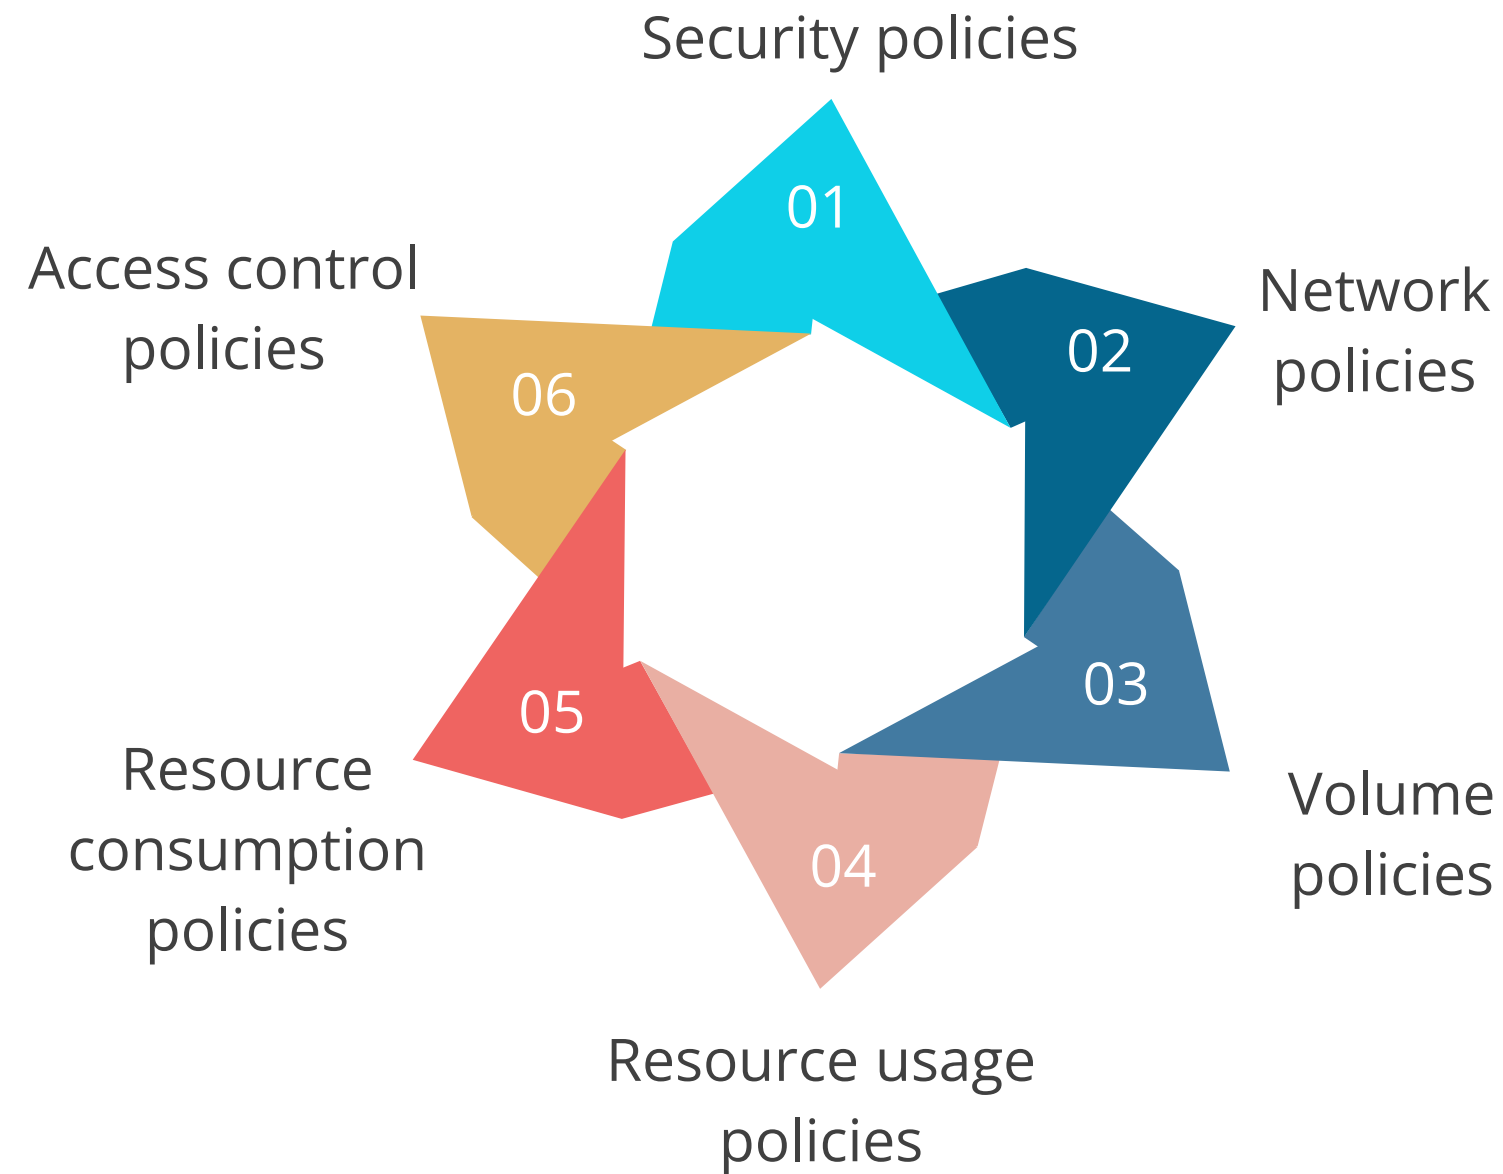


Better flexibility to  
changing requirements



# Policy Restrictions

On a Kubernetes cluster, containers run with unbounded compute resources by default. To limit or restrict, appropriate policies must be implemented in the following ways:



## Key Takeaways

- Containers are lightweight, standalone, executable software packages that include everything required to run an application: code, runtime, system tools, system libraries, and settings.
- Etcctl is a reliable and highly available key-value store that serves as the backup store for all cluster data in Kubernetes.
- Kube-proxy is a network proxy that runs on every node in a cluster, implementing the Kubernetes Service concept.
- Policies define what end users can do on the cluster and ways to ensure that clusters comply.





## Fetching Cluster Specific Configuration

Duration: 25 Min.

**Description:** Your team lead has asked you to access the Kubernetes cluster and report the following cluster-specific details:

- Available nodes and their IP addresses
- Supported API versions on the server
- Status of the control plane and CoreDNS
- Status of pods with the kube-system namespace

### Steps to Perform:

1. List the available nodes and their IP addresses
2. Identify supported API versions
3. Examine the control plane and CoreDNS status
4. Review the status of the pods in the kube-system namespace



# TECHNOLOGY

**Thank You**