

TECHNOLOGY



Container Orchestration Using Kubernetes

Services, Load Balancing, and Networking



A Day in the Life of a DevOps Engineer

You are working as a DevOps engineer and having trouble comprehending Kubernetes services, networking policies, and how load balancing is done.

The goal is to understand how services are configured, networking policies are created, and load balancing is achieved. Your organization needs someone to deploy and manage applications that route traffic to several containers.

To achieve these and some additional features, you will learn a few concepts in this lesson that will assist you in solving the given scenario.



Learning Objectives

By the end of this lesson, you will be able to:

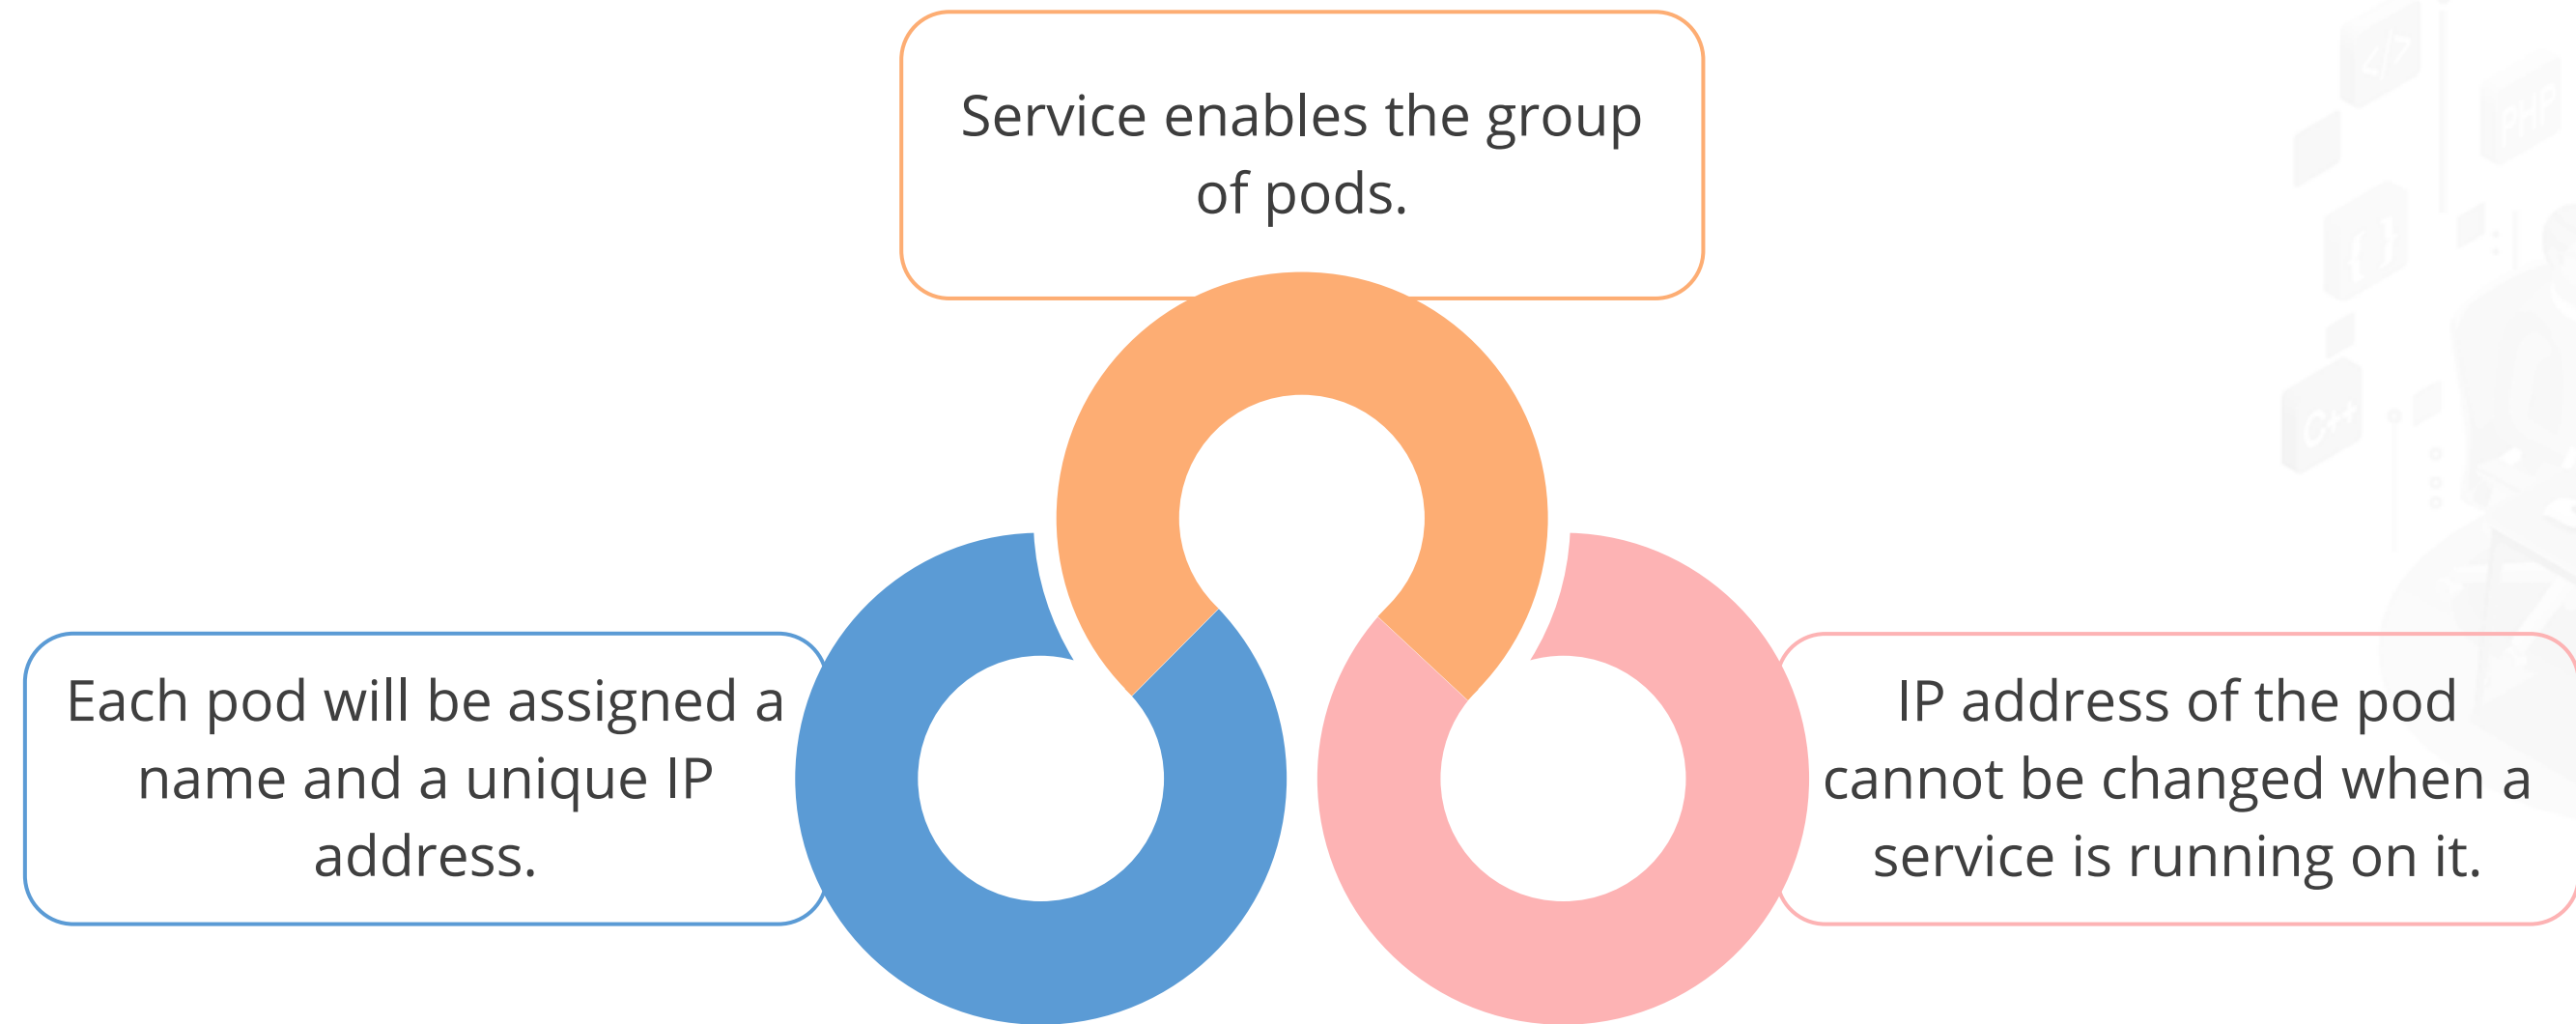
- Demonstrate networking solutions with Kubernetes to manage containerized applications
- Demonstrate the topology and DNS to make Kubernetes services discoverable
- Work with EndpointSlices and Ingress controllers to enable object tracking
- Create and apply an IPv4/IPv6 dual-stack network policy for the autonomous transition of addresses
- List how cluster networking is implemented to run containers across multiple machines



Overview of Services

Services

Kubernetes service is a logical abstraction for a deployed group of pods in a cluster.



Networking

Kubernetes networking addresses four concerns:

Loopback is used by the containers within a pod for communication.

Communication between pods happens through cluster networking.

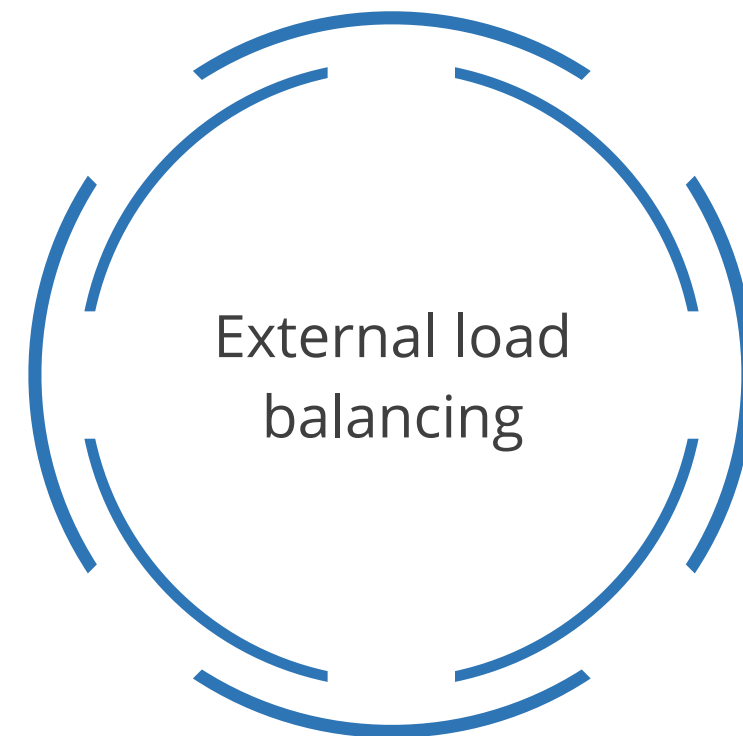
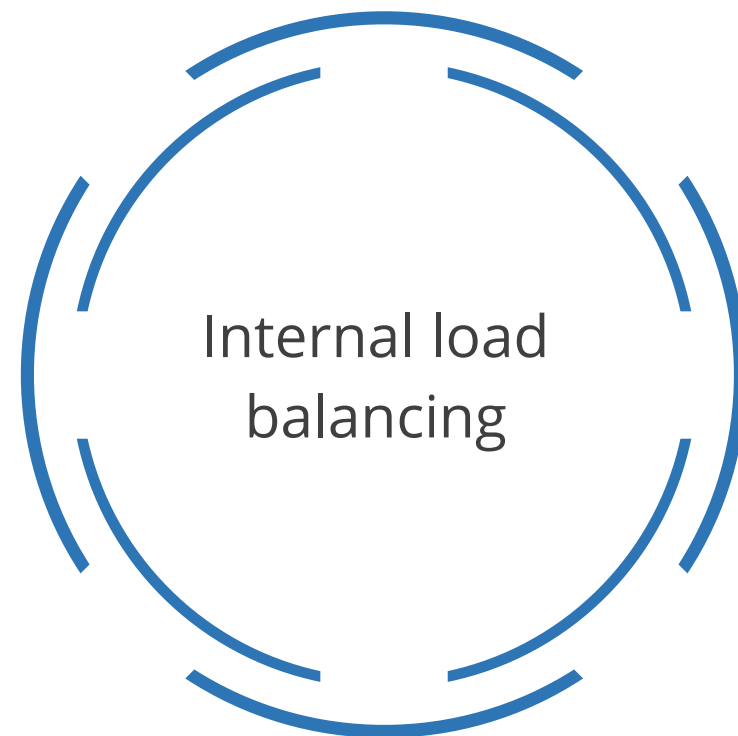
The application running in a pod can be made reachable from outside the cluster using the service resource.

Services can be published only for consumption within the cluster.



Load Balancing

Load balancing helps in the balanced distribution of network traffic or client requests across multiple servers. It can be:



Cluster Networking

Overview of Networking

Networking is a central part of Kubernetes, but it can be challenging to understand exactly how it is expected to work.

There are four distinct networking problems to address:



Highly-coupled container-to-container communications



Pod-to-Pod communications



Pod-to-Service communications



External-to-Service communications



Kubernetes Network Model

Every pod is assigned an IP address. Kubernetes imposes fundamental requirements on any networking implementation.

Kubernetes allows communication between:

1

Pods in the host network of a node and pods on all nodes without NAT

2

Agents and all the pods on a node



Implementing Kubernetes Network Model

There are several ways in which this network model can be implemented:

1 | AOS from Apstra

2 | AWS VPC CNI for Kubernetes

3 | Azure CNI for Kubernetes

4 | Cilium

5 | CNI-Genie from Huawei

6 | cni-ipvlan-vpc-k8s

7 | Coil

8 | Contiv

9 | Contrail/Tungsten Fabric

10 | DANM

11 | Knitter

12 | Kube-OVN

13 | Kube-router

14 | Multus

15 | NSX-T

16 | Nuage Networks VCS



Implementing Kubernetes Network Model

There are several ways in which this network model can be implemented:

17 | OpenVSwitch

18 | ACI

19 | Flannel

20 | Jaguar

21 | Romana

22 | Weave Net from Weaveworks

23 | OVN

24 | Antrea

25 | Google Compute Engine

26 | k-vswitch

27 | Big Cloud Fabric from Big Switch Networks



TECHNOLOGY

Services

What Is a Service?

A service in Kubernetes is an abstraction that exposes a set of pods to be accessed. It also provides a policy for accessing them. The policy is referred to as a microservice.

Kubernetes makes the pods discoverable by providing them with their IP addresses and a single DNS name for a set of pods.

Kubernetes also balances the load between pods.

Define a Service

A service is a REST object. POST a service definition to the API server, like any other REST object, to create a new instance.

Demo:

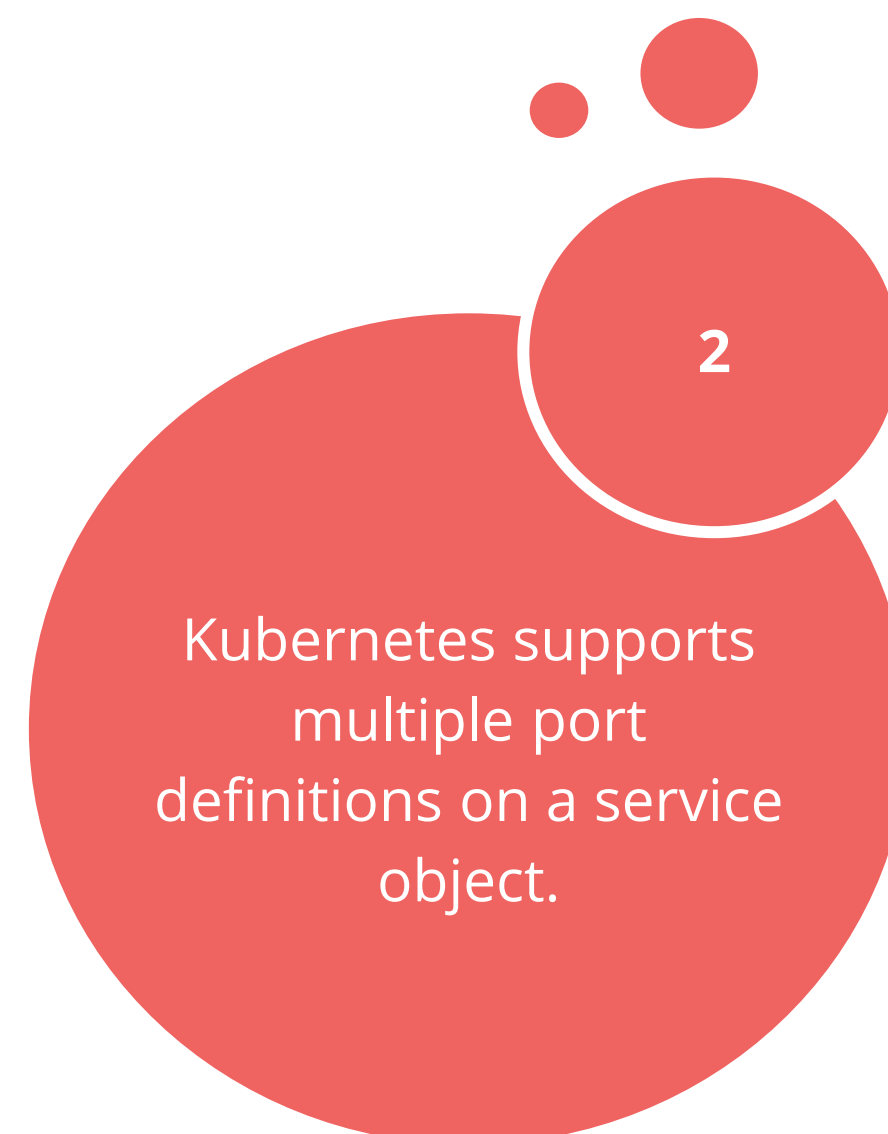
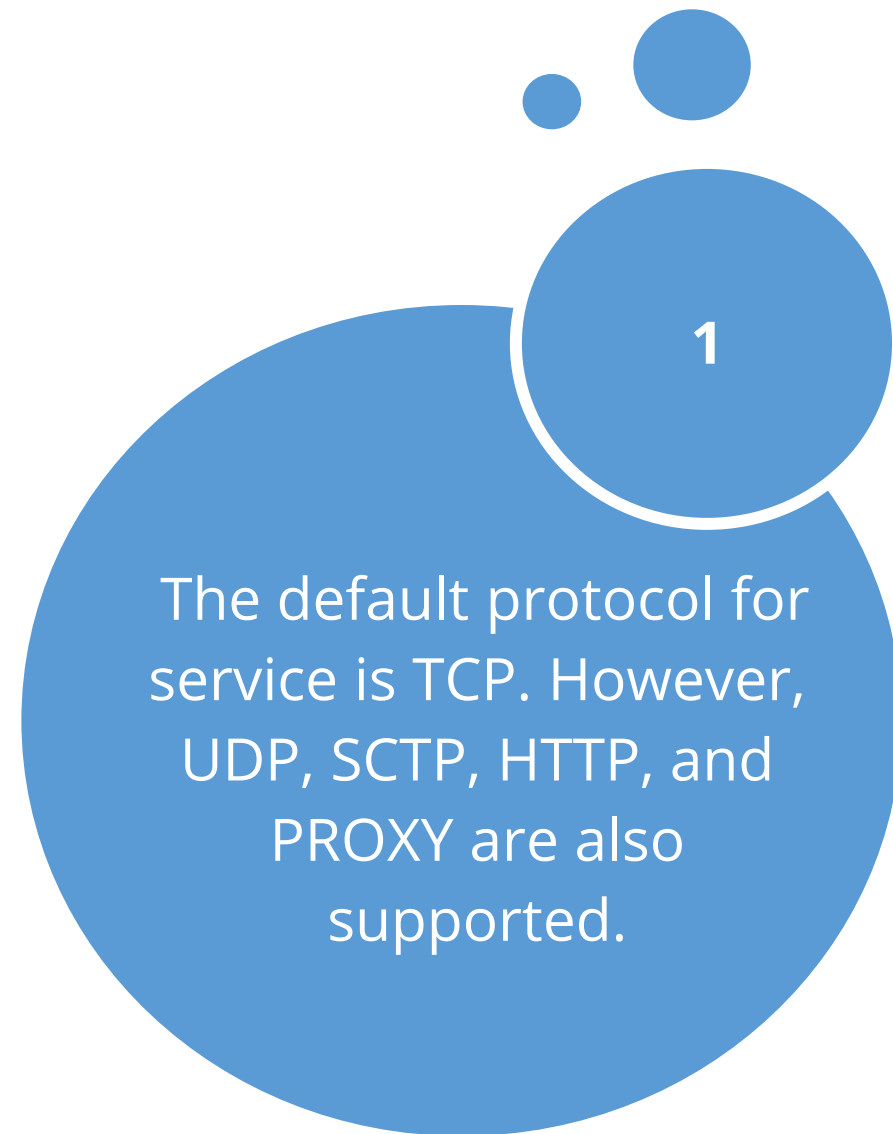
```
apiVersion: v1
Kind: service
Metadata:
  name: my-service
spec:
  selector:
    app: MyApp
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
```

This specification creates a new service object called **my-service**, which targets TCP port **9376** on any pod with the app labeled **MyApp**.



Service

A service can map any incoming port to a **targetPort**. By default, the **targetPort** is set to the same value as the port field.



Service Without Selectors

Services mostly abstract access to Kubernetes pods although they can also abstract other kinds of backends.

To define a service without a pod selector, add the following in the YAML file:

Demo:

```
apiVersion: v1
Kind: service
Metadata:
  name: my-service
spec:
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
```



Service Without Selectors

A service can be mapped manually to the network address and the port where it is running by adding an endpoints object.

Demo:

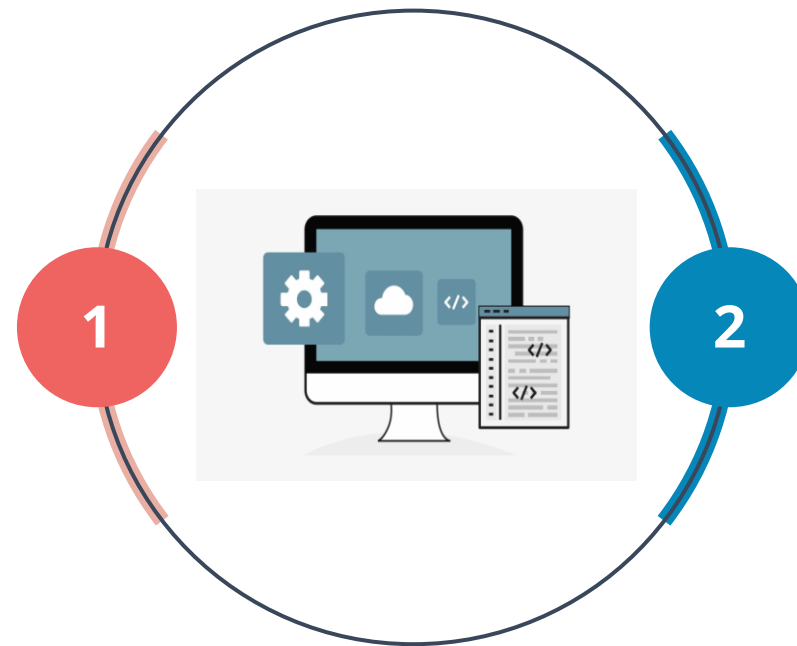
```
apiVersion: v1
Kind: Endpoints
Metadata:
  name: my-service
sybsets:
- addresses:
  -ip: 192.0.2.42
  ports:
  - port: 9376
```



Application Protocol

Kubernetes services expose one or more ports. A port exposed by an application can serve a specific application protocol such as HTTP, TCP, gRPC, and so on.

The **appProtocol** field provides a way to specify an application protocol for each service port.



It follows standard Kubernetes label syntax.

The values should either be IANA standard service names or domain prefixed names.

Virtual IPs

When users create a service, a new virtual IP (also known as a **clusterIP**) is created on their behalf.

1 Every node in a Kubernetes cluster runs a **kube-proxy**.

2 **kube-proxy** is responsible for implementing a type of virtual IP for services other than **ExternalName**.



Multi-Port Services

Kubernetes helps to configure multiple port definitions on a service object. The configuration below shows the use of the **port** attribute:

Demo:

```
apiVersion: v1
Kind: service
Metadata:
  name: my-service
spec:
  selector:
    app: MyApp
  ports:
    - name: http
      protocol: TCP
      port: 80
    - name: http
      protocol: TCP
      port: 443
      targetPort: 9377
```



Choose IP Addresses

Cluster IP address can be specified as a part of a service creation request.



Note

The chosen IP address must be a valid IPv4 or IPv6 address from within the service-cluster-ip-range CIDR range that is configured for the API server.



Environment Variables

The kubelet adds a set of environment variables for each active service when a pod is running on a node.

```
REDIS_MASTER_SERVICE_HOST = 10.0.0.11
REDIS_MASTER_SERVICE_PORT= 6379
REDIS_MASTER_PORT = tcp://10.0.0.11:6379
REDIS_MASTER_PORT_6379_TCP = tcp://10.0.0.11:6379
REDIS_MASTER_PORT_6379_TCP_PROTO = tcp
REDIS_MASTER_PORT_6379_TCP_PORT = 6379
REDIS_MASTER_PORT_6379_TCP_ADDR = 10.0.0.11
```

The example above shows the environment variables for the service **redis-master**, which exposes TCP port **6379** and has been allocated cluster IP address **10.0.0.11**.



DNS

A cluster-aware DNS server, such as CoreDNS, watches the Kubernetes API for new services and creates a set of DNS records for each.

The administrator should always set up a DNS service for the Kubernetes cluster using an add-on.

The Kubernetes DNS server is the only way to access ExternalName services.

DNS

Example:

For a service called **my-service** in a Kubernetes namespace **my-ns**, the control plane and the DNS service together create a DNS record for **my-service.my-ns**.

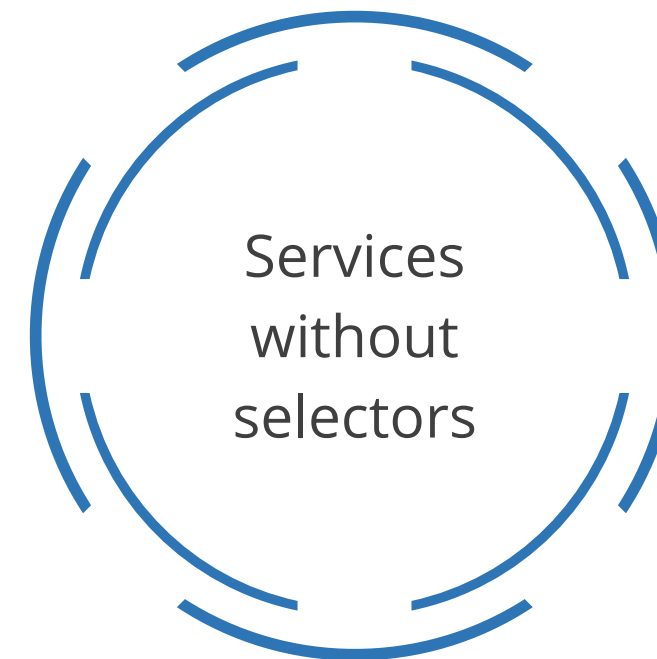
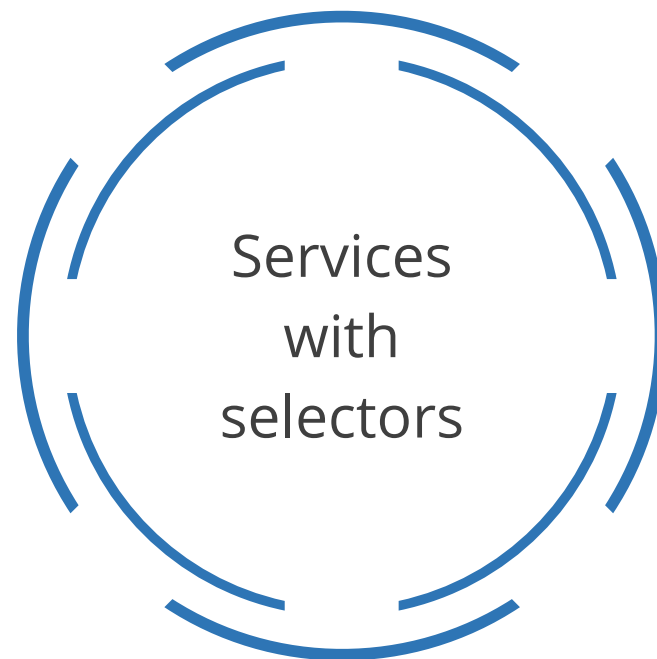
Pods in the **my-ns** namespace should be able to find the service by doing a name lookup for **my-service**.

Pods in other namespaces must qualify the name as **my-service.my-ns**.

Headless Service

A headless service is a service that has no service IP. They are used to interface with other service discovery mechanisms without being tied to the implementation of Kubernetes.

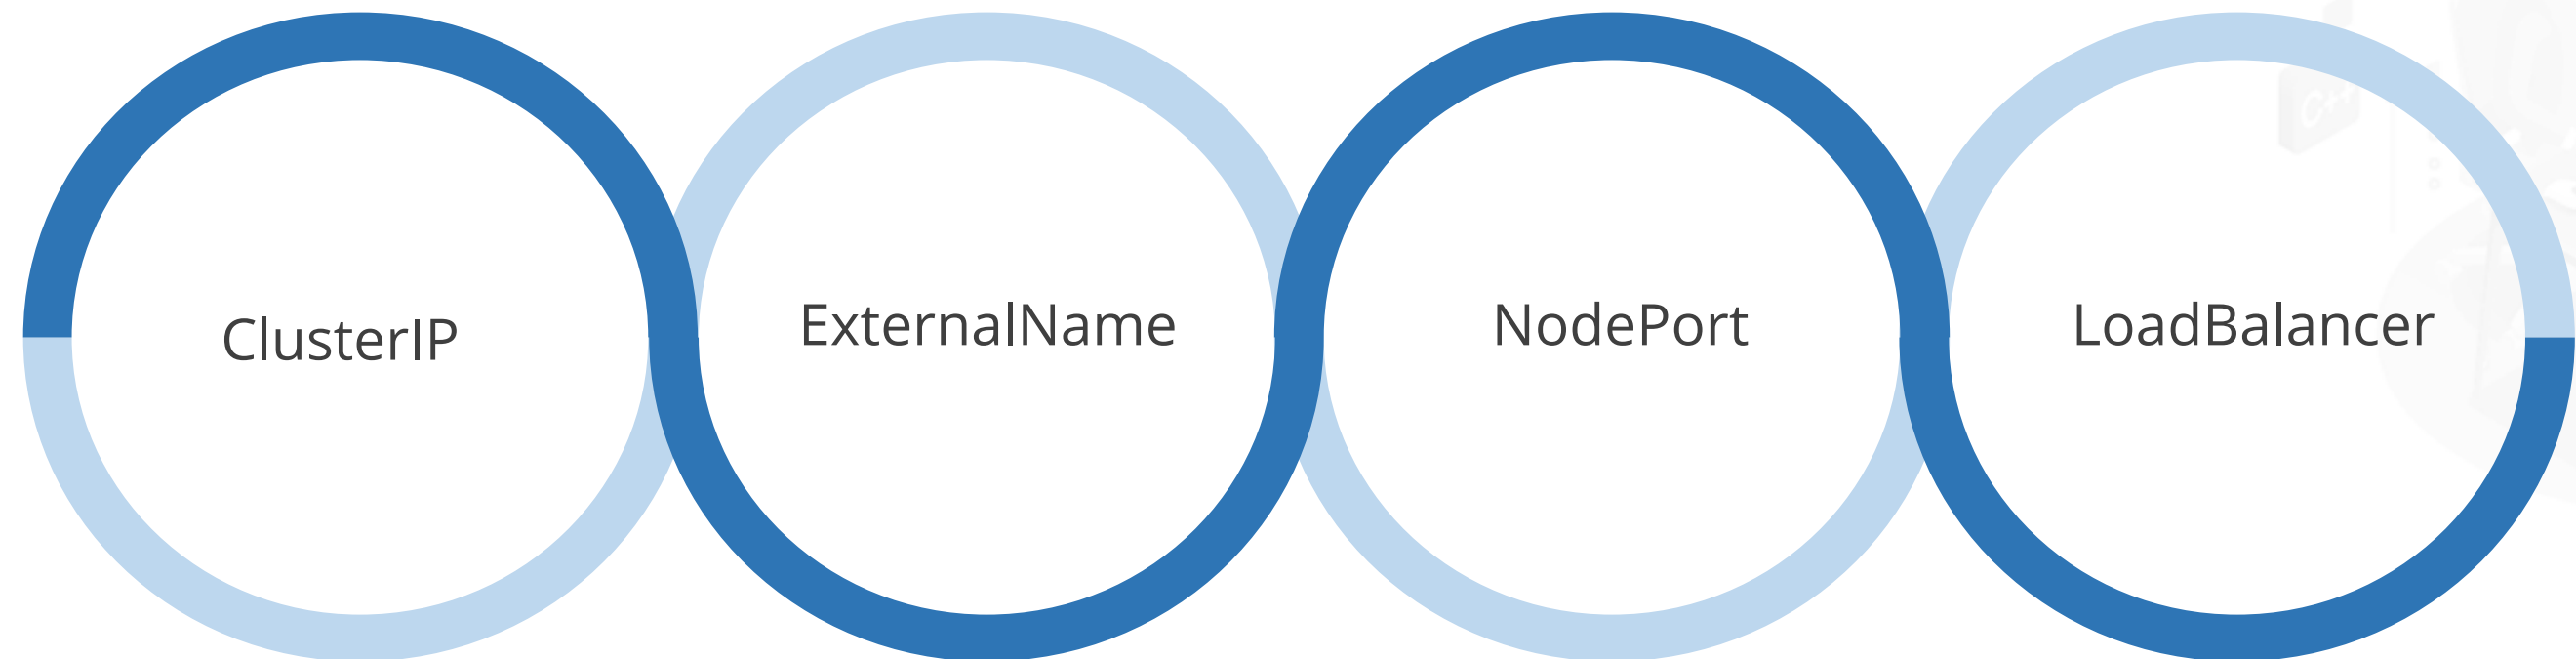
Headless services can be classified into two categories:



Publishing Services (ServiceTypes)

Kubernetes **ServiceTypes** help to specify the required kind of service. The default is ClusterIP.

Type values and their behaviors include:



Type NodePort

The NodePort service type allows setting up a custom load balancing solution for configuring environments that are not fully supported by Kubernetes.

The NodePort type can be configured as shown below:

Demo

```
apiVersion: v1
kind: service
metadata:
  name: my-service
spec:
  type: NodePort
  selector:
    app: MyApp
  ports:
    # By default, the 'targetPort' is set to the same value as the 'port' field.
    - port: 80
      targetPort: 80
    # By default, the kubernetes control plane will allocate a port from a range (30000-32767)
    NodePort: 30007
```


Type LoadBalancer

On cloud providers support external load balancers by setting the type field to **LoadBalancer** which provisions a load balancer for the service.

Demo

```
apiVersion: v1
Kind: service
Metadata:
  name: my-service
spec:
  selector:
    app: MyApp
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
  clusterIP: 10.0.171.239
  type: LoadBalancer
status:
  LoadBalancer: 9377
  ingress:
    - ip: 192.0.2.127
```

To direct traffic from the external load balancer to the backend pods, use the configuration shown.

The **.status.loadBalancer** field contains information about the provisioned balancer.

Load Balancers with Mixed Protocol Types

When multiple ports are configured for load balancer services, all ports must have the same protocol by default.

The protocol must be supported by the cloud provider.

If the feature gate **MixedProtocolLBService** is enabled for the **kube-apiserver**, it can use different protocols when multiple ports are specified.

Disable Load Balancer NodePort Allocation

The node port allocation for a service of **Type=LoadBalancer** can be optionally disabled.

This can be done by setting the field **spec.allocateLoadBalancerNodePorts** to false.



By default, **spec.allocateLoadBalancerNodePorts** is true and type LoadBalancer services will continue to allocate node ports.



If **spec.allocateLoadBalancerNodePorts** is set to false on an existing service with allocated node ports, those node ports will not be de-allocated automatically.



The **ServiceLBNodePortControl** feature gate must be enabled to use this field.

Connecting Applications with Services

Kubernetes Model for Connecting Containers

Now that users have a continuously running and replicated application, they can expose it on a network.

1

Kubernetes assumes that pods can communicate with other pods, regardless of the host on which they land.

2

Kubernetes gives every pod its own cluster-private IP address. As a result, there is no need to make explicit linkages between pods or map container ports to host ports.

Expose Pods to a Cluster

When pods are exposed to a cluster, it makes them accessible from any node in the cluster.
This is done as follows:

Demo

```
apiVersion: apps/ v1
Kind: Deployment
metadata:
  name: my-nginx
spec:
  selector:
    matchLabels:
      run: my-nginx
  replicas: 2
  template:
    metadata:
      labels:
        run: my-nginx
    spec:
      containers:
        - name: my-nginx
          image: nginx
          ports:
            - containerPort: 80
```



Expose Pods to a Cluster

The command to verify the nodes on which the pod is running is as follows:

Demo

Command:

```
kubectl apply -f ./run-my-nginx.yaml  
kubectl get pods -l run=my-nginx -o wide
```

Result:

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
my-nginx-3800858182-jr4a2	1/1	Running	0	13s	10.244.3.4	kubernetes-minion-905m
my-nginx-3800858182-kna2y	1/1	Running	0	13s	10.244.2.5	kubernetes-minion-ljyd

Create a Service

A Kubernetes service refers to the abstraction that defines a logical set of pods running in a cluster with the same functionality. The configuration below helps to create a service for **nginx** replica:

Demo

```
#Create a Service for your 2 nginx replicas  
with kubectl expose:
```

Command:

```
kubectl expose deployment/my-nginx
```

Result:

```
service/my-nginx exposed
```

Demo

```
#This is equivalent to kubectl apply -f the  
following yaml:
```

```
apiVersion: v1  
Kind: Service  
metadata:  
  name: my-nginx  
  labels:  
    run: my-nginx  
spec:  
  ports:  
    - port: 80  
      protocol: TCP  
  selector:  
    run: my-nginx
```

Create a Service

This specification will create a service that targets TCP port **80** on any pod with the **run my-nginx label**. It exposes it on an abstracted service port.

Demo

Command:

```
kubectl get svc my-nginx
```

Result:

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
my-nginx	ClusterIP	10.0.162.149	<none>	80/TCP	21s



Create a Service

Service is backed by a group of pods. These pods are exposed through **endpoints**.

The service's selector is evaluated continuously, and the results are posted to an endpoints object, also named **my-nginx**.

When a pod dies, it is automatically removed from the endpoints, and new pods matching the service's selector automatically gets added to the endpoints.

Create a Service

Check the endpoints and note that the IPs are the same as the pods created in the first step:

Demo

Command:

```
kubectl describe svc my-nginx
```

Result:

```
Name:          my-nginx
Namespace:     default
Labels:        run=my-nginx
Annotations:    <none>
Selector:      run=my-nginx
Type:          ClusterIP
IP:            10.0.162.149
Port:          <unset> 80/TCP
Endpoints:     10.244.2.5:80,10.244.3.4:80
Session Affinity: None
Events:        <none>
```

Demo

Command:

```
kubectl get ep my-nginx
```

Result:

NAME	ENDPOINTS	AGE
my-nginx	10.244.2.5:80,10.244.3.4:80	1m

Access a Service Using Environment Variables

When a pod runs on a node, the kubelet adds a set of environment variables for each active service.

Demo

Command:

```
kubectl exec my-nginx-3800858182-jr4a2 -- printenv | grep SERVICE
```

Result:

```
KUBERNETES_SERVICE_HOST= 10.0.0.1  
KUBERNETES_SERVICE_PORT= 443  
KUBERNETES_SERVICE_PORT_HTTPS= 443
```



Access a Service Using DNS

Kubernetes offers a DNS cluster addon service that automatically assigns DNS names to other services.

Demo

Command:

```
kubectl get services kube-dns --namespace=kube-system
```

Result:

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kube-dns	ClusterIP	10.0.0.10	<none>	53/UDP,53/TCP	8m



Secure a Service

These are the requirements of a secure communication channel:



Secure a Service

The following commands are used to create a Nginx HTTPS service useful in verifying proof of **concept, keys, secrets, configmap**, and **end-to-end HTTPS service** creation in Kubernetes.

Demo

Command:

```
make keys KEY=/tmp/nginx.key CERT=/tmp/nginx.crt
kubectl create secret tls nginxsecret --key /tmp/nginx.key --cert /tmp/nginx.crt
```

Result:

```
secret/nginxsecret created
```

Command:

```
kubectl get secrets
```

Result:

NAME	TYPE	DATA	AGE
default-token-il9rc	kubernetes.io/service-account-token	1	1d
nginxsecret	kubernetes.io/tls	2	1m



Secure a Service

Here is an example of how to create a ConfigMap:

Demo

Command:

```
kubectl create configmap nginxconfigmap --from-file=default.conf
```

Result:

```
configmap/nginxconfigmap created
```

Command:

```
kubectl get configmaps
```

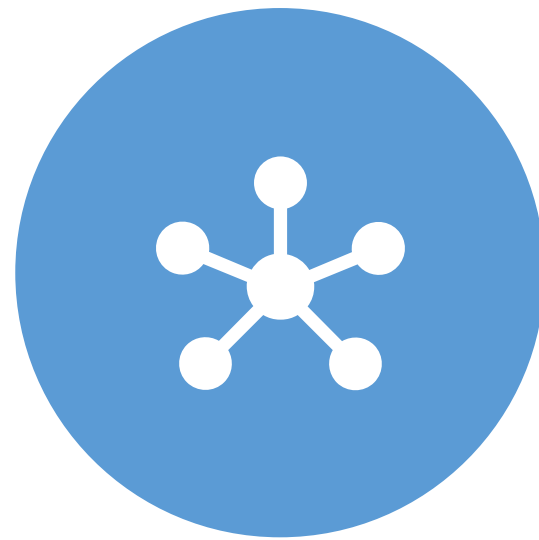
Result:

NAME	DATA	AGE
nginxconfigmap	1	114s



Expose a Service

Kubernetes supports two ways of exposing a service:



NodePorts



LoadBalancers



Expose a Service

The YAML configuration given below creates a Nginx HTTPS replica to serve the traffic on the internet if the node has a public IP:

Demo

```
kubectl get svc my-nginx -o yaml | grep nodePort -C 5
uid: 07191fb3-f61a-11e5-8ae5-42010af00002
spec:
  clusterIP: 10.0.162.149
  ports:
  - name: http
    nodePort: 31704
    port: 8080
    protocol: TCP
    targetPort: 80
  - name: https
    nodePort: 32453
    port: 443
    protocol: TCP
    targetPort: 443
```

Demo

```
kubectl get nodes -o yaml | grep ExternalIP -C 1
- address: 104.197.41.11
  type: ExternalIP
allocatable:
--
- address: 23.251.152.56
  type: ExternalIP
allocatable:
...

$ curl https://<EXTERNAL-IP>:<NODE-PORT> -k
...
<h1>Welcome to nginx!</h1>
```

Expose a Service

The IP address in the **EXTERNAL-IP** column is the one that is available on the public internet.
The **CLUSTER-IP** is only available inside the cluster or private cloud network.

Command:

```
kubectl edit svc my-nginx  
kubectl get svc my-nginx
```

Result:

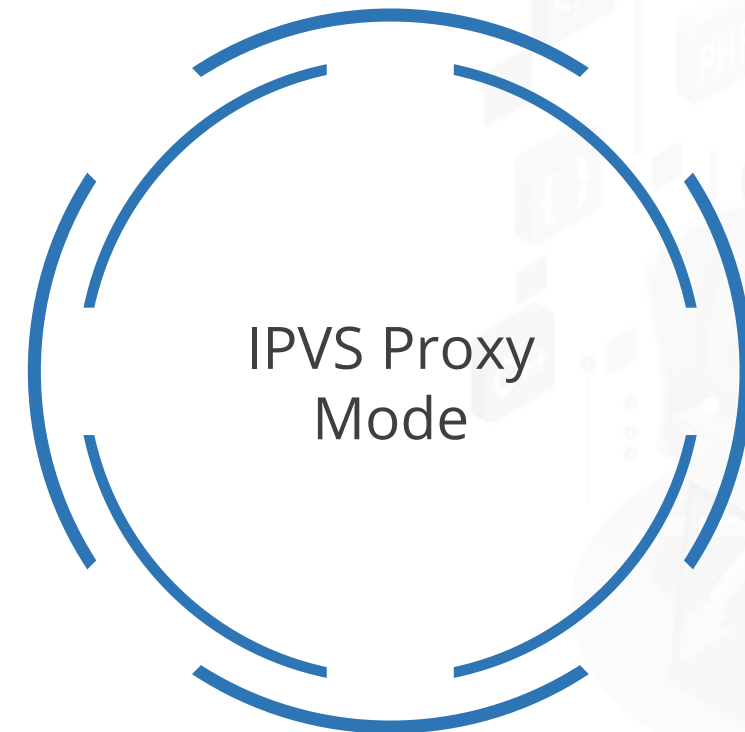
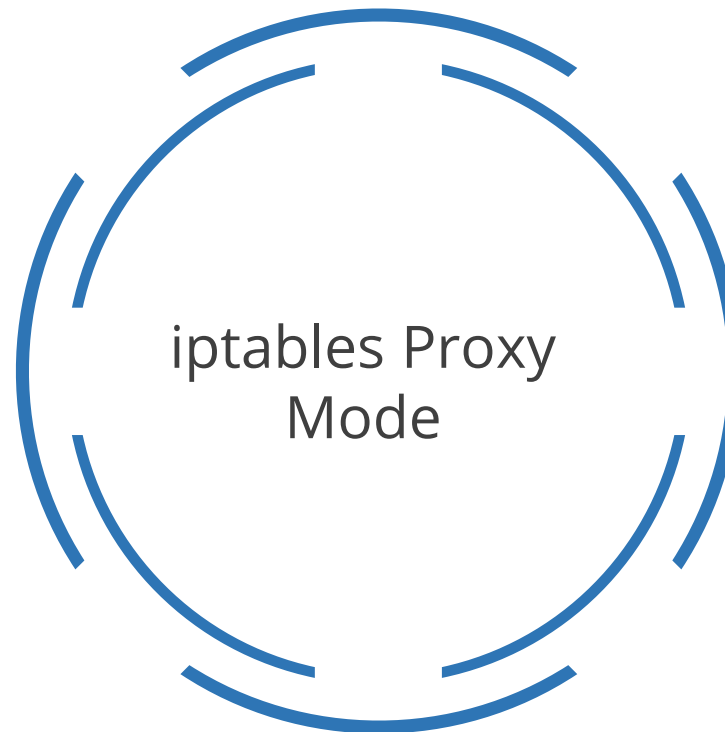
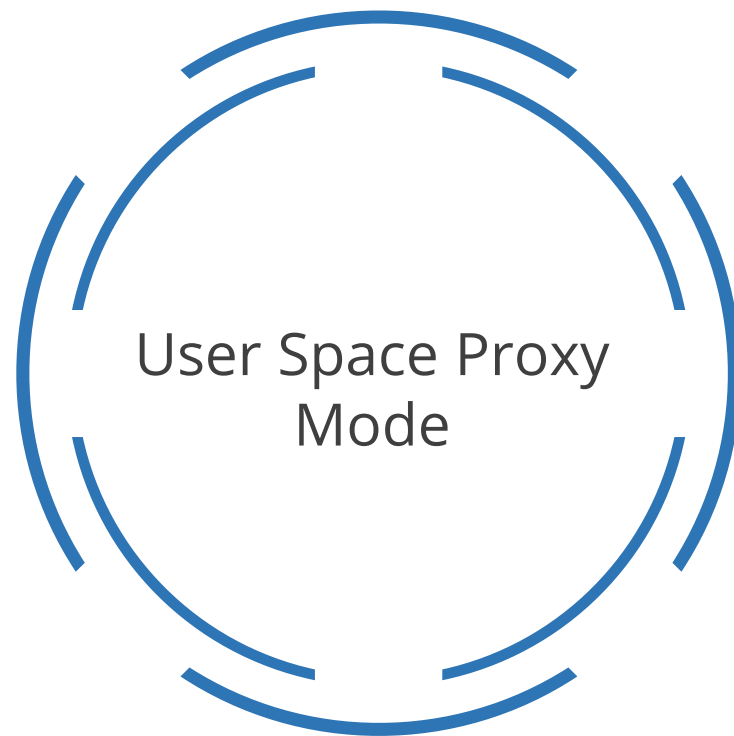
NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
my-nginx	LoadBalancer	10.0.162.149	xx.xxx.xxx.xxx	8080:30163/TCP	21

```
curl https://<EXTERNAL-IP> -k  
...  
<title>Welcome to nginx!</title>
```



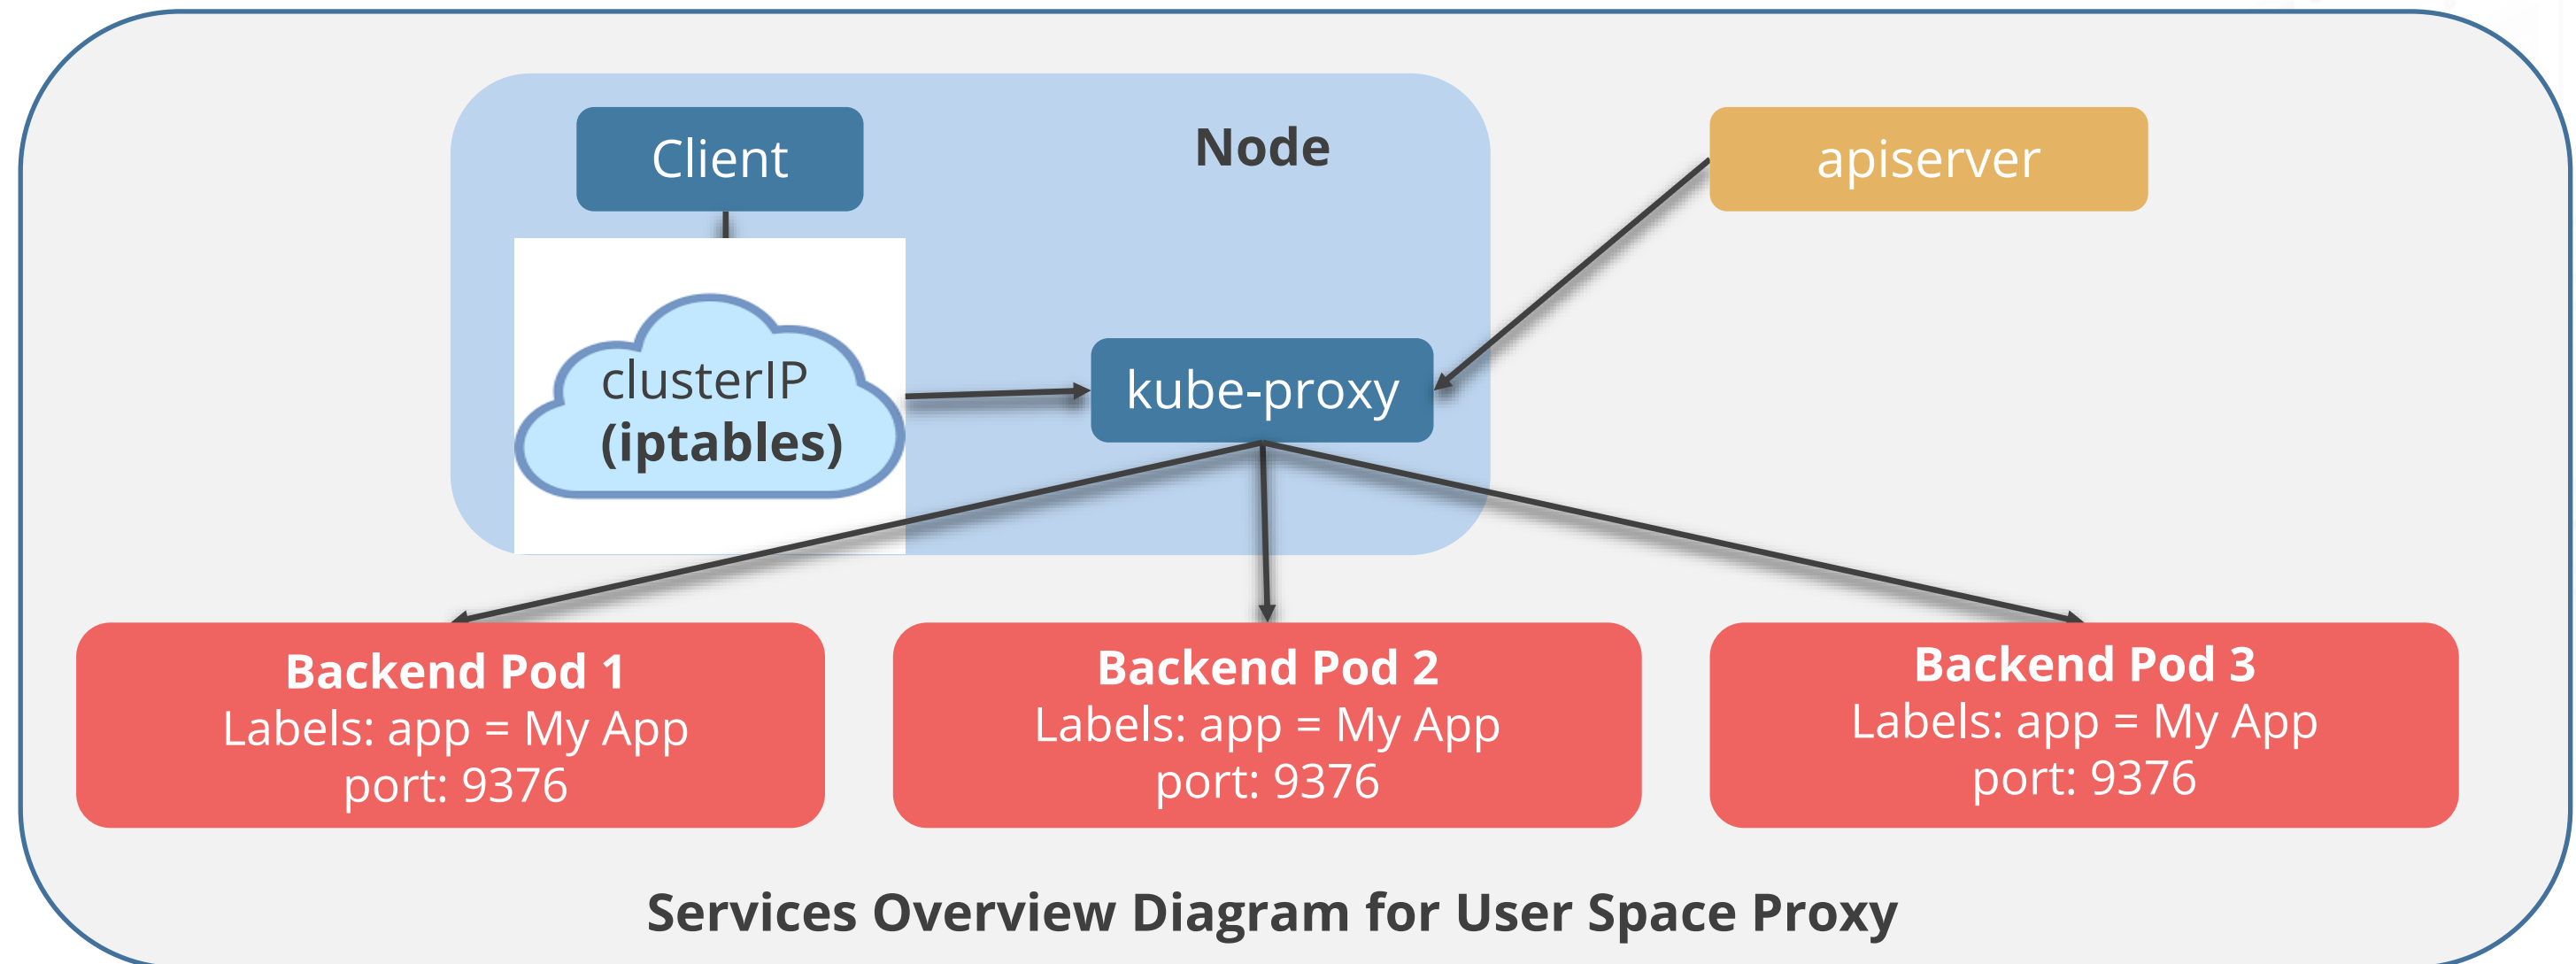
Proxy Modes

Three types of proxy modes supported in Kubernetes are:



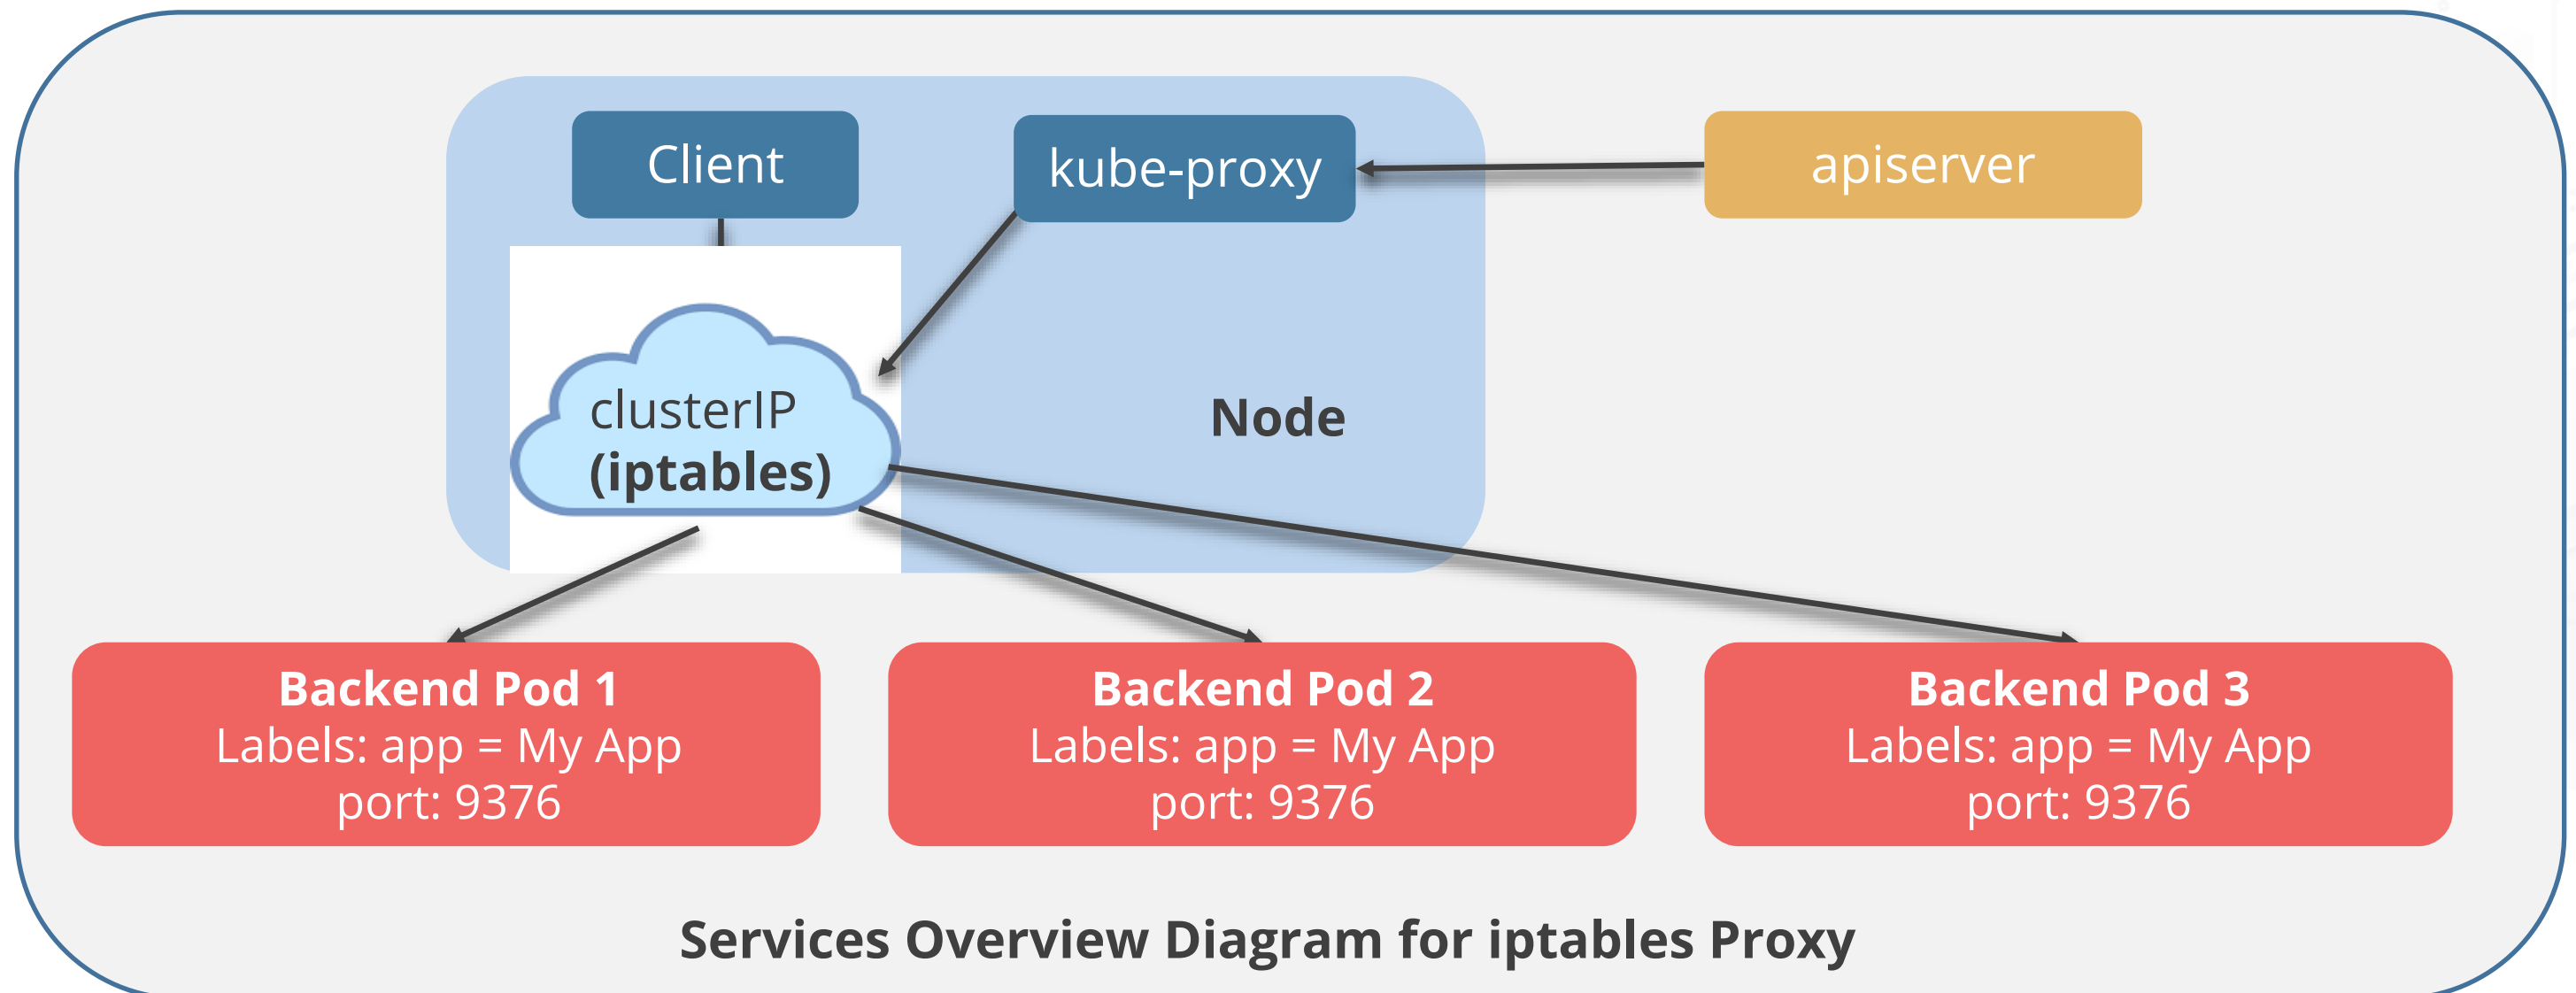
User Space Proxy Mode

In this (legacy) mode, kube-proxy watches the Kubernetes control plane for the addition and removal of service and endpoint objects. By default, kube-proxy in userspace mode chooses a backend via a round-robin algorithm.



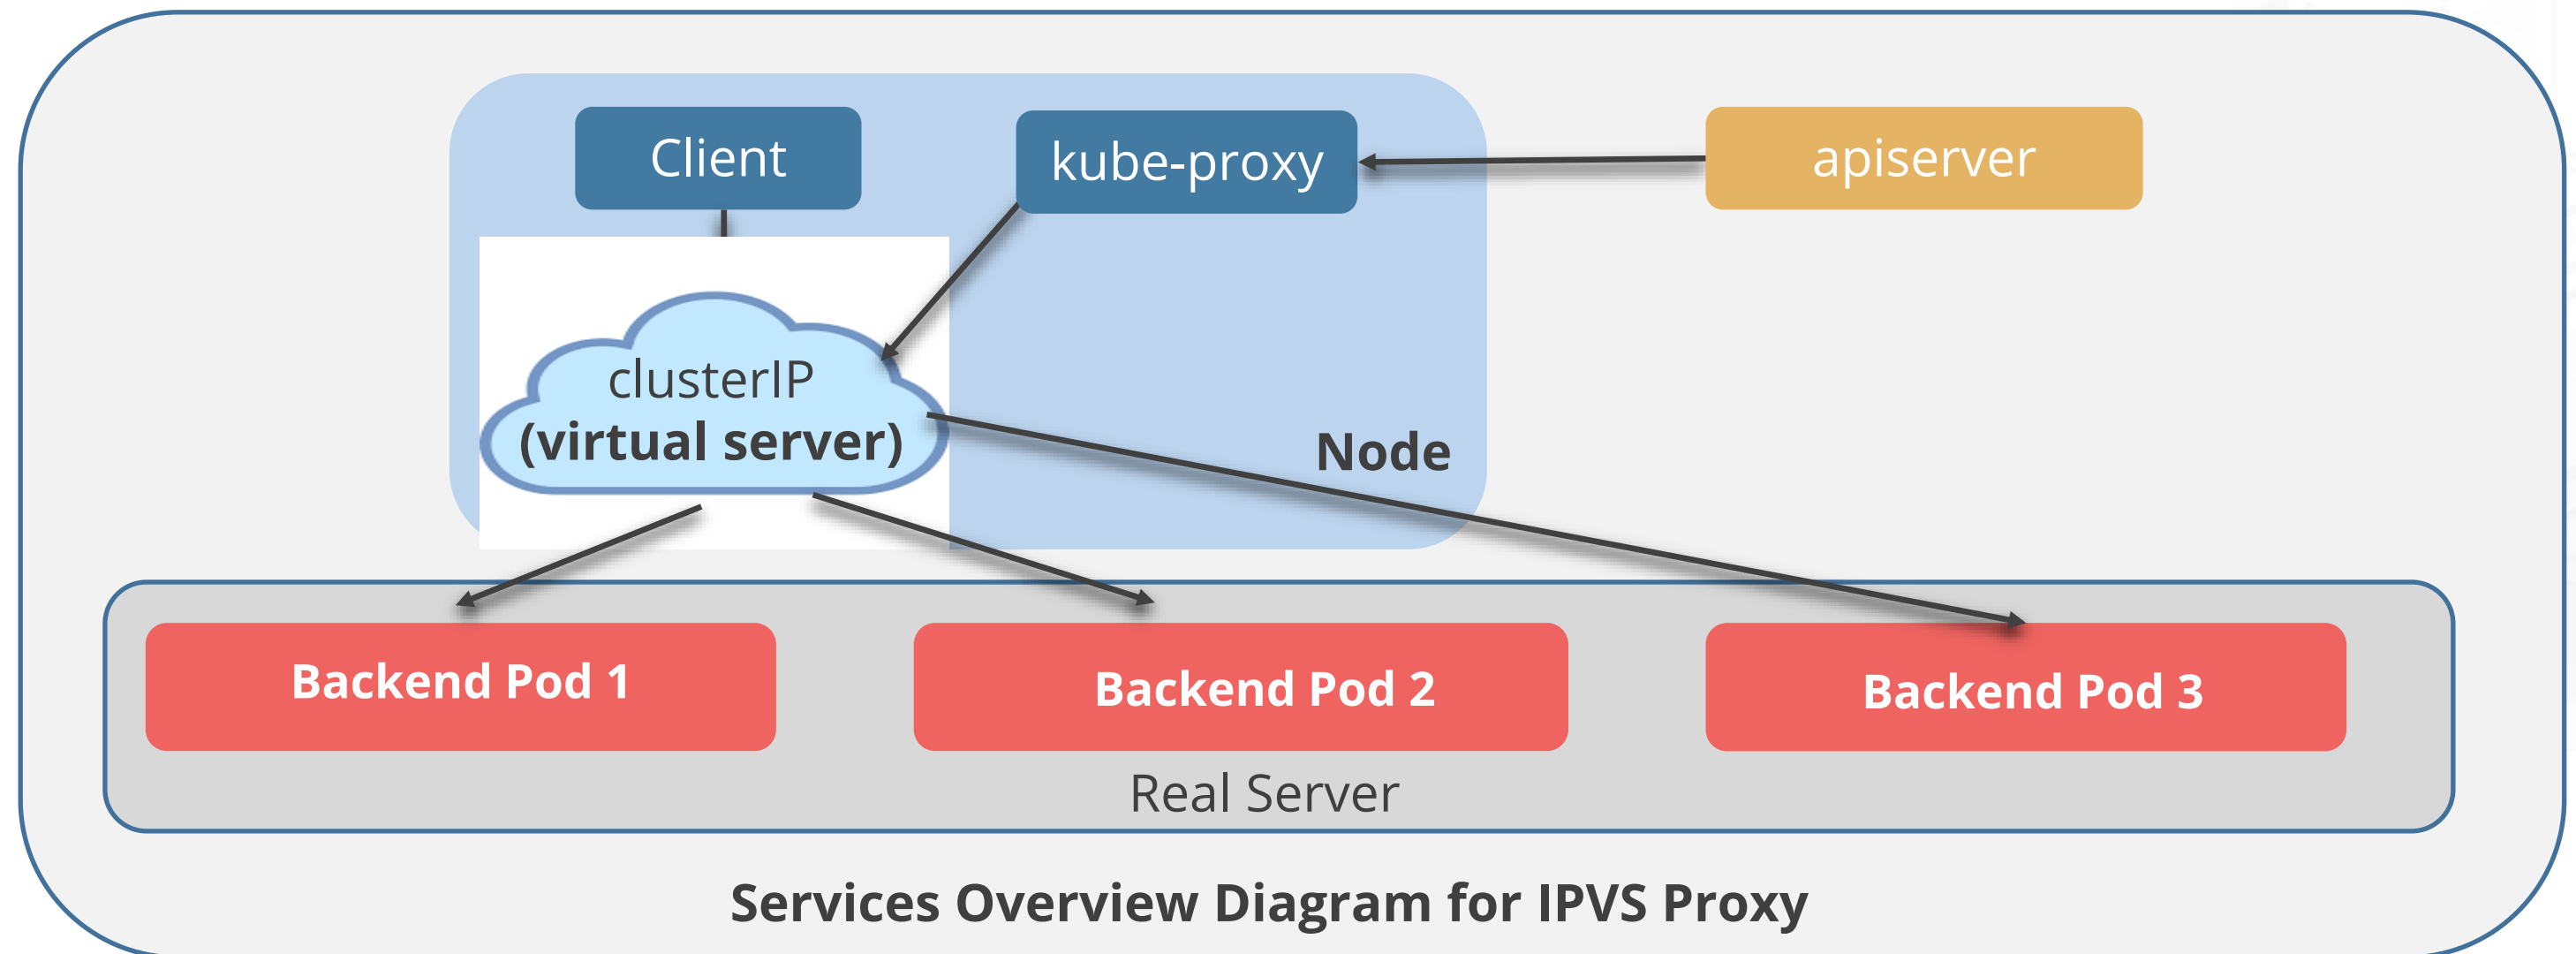
iptables Proxy Mode

In this mode, kube-proxy watches the Kubernetes control plane for the addition and removal of service and endpoint objects. By default, kube-proxy in iptables mode chooses a backend at random.



IPVS Proxy Mode

In IPVS proxy mode, kube-proxy watches Kubernetes services and endpoints and calls net link interface to create IPVS rules accordingly. It synchronizes IPVS rules with Kubernetes services and endpoints periodically.



Deploying a Multi-Port Service Pod



Duration: 10 mins

Problem Statement:

You've been asked to deploy a Kubernetes pod using a multi-port service for accessing the deployment through multiple ports

ASSISTED PRACTICE

Assisted Practice: Guidelines

Steps to be followed:

1. Create a deployment
2. Define a dervice
3. Access the deployment from multiple ports



Topology

Topology-Aware Traffic Routing

Topology enables a service to route traffic based on the node topology of the cluster.

By default, traffic sent to a **ClusterIP** or **NodePort** service can be routed to any backend address for the service.

Label matching between the source and destination lets the cluster operator designate sets of nodes that are closer and farther from one another.

Usage of Service Topology

Service traffic routing can be controlled if the **ServiceTopology** feature of the cluster is enabled.

Traffic will be directed to the node whose first label value matches the originating node's value.

Topology constraints will not be applied if the topologyKeys field is empty or unspecified.



A second label will be considered when there is no backend for the service on a matching node.

Constraints

1

Service topology is not compatible with **externalTrafficPolicy=Local**, and therefore, a service cannot use both of these features.

2

Valid topology keys are currently limited to **kubernetes.io/hostname**, **topology.kubernetes.io/zone**, and **topology.kubernetes.io/region**.

3

Topology keys must be valid label keys. At most, 16 keys may be specified.

4

If the catch-all value ***** is used, it must be the last value in the topology keys.

DNS for Services and Pods

Introduction to DNS

Kubernetes creates DNS records for services and pods.

A DNS query may return different results based on the namespace of the pod making it.

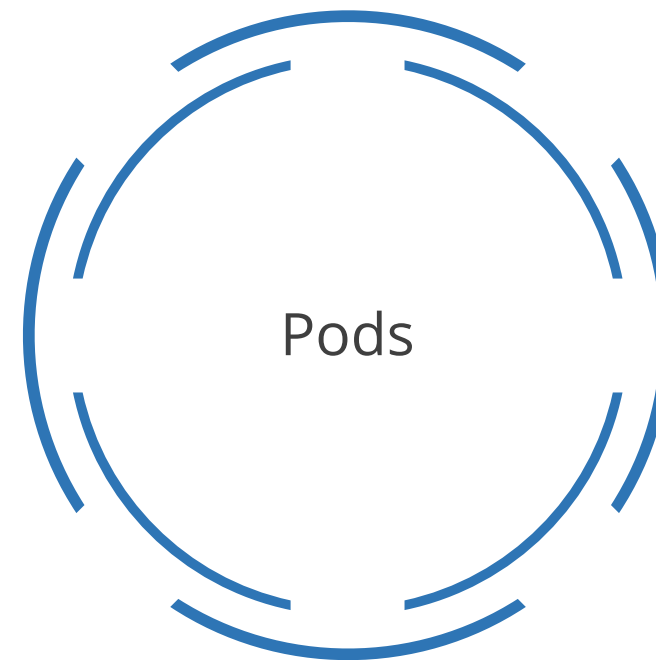
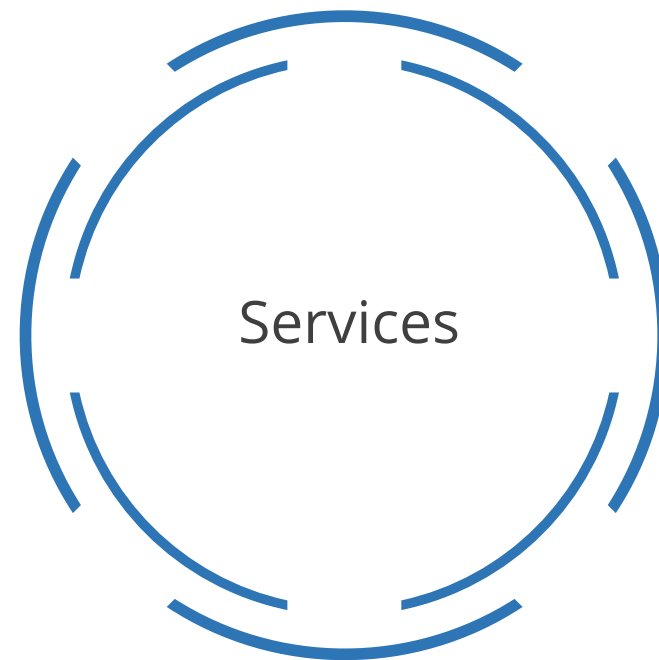


DNS queries may be expanded using the pod's **/etc/resolv.conf** file.

Kubernetes DNS schedules a DNS pod and service on the cluster and configures the Kubelets to tell individual containers to use the DNS service's IP to resolve DNS names.

DNS Records

The following objects get DNS records:



Services

Snapshots can be provisioned in two ways:

1

Normal services are assigned a DNS A or AAAA record, depending on the IP family of the service, for the name of the form **my-svc.my-namespace.svc.cluster-domain.example**.

2

SRV records are created for named ports that are part of normal or headless services.

Pods

A pod has the following DNS resolution:

```
Pod-ip-address.my-namespace.Pod.cluster-domain.example
```



The pod spec has an optional hostname field, which can be used to specify the pod's hostname.

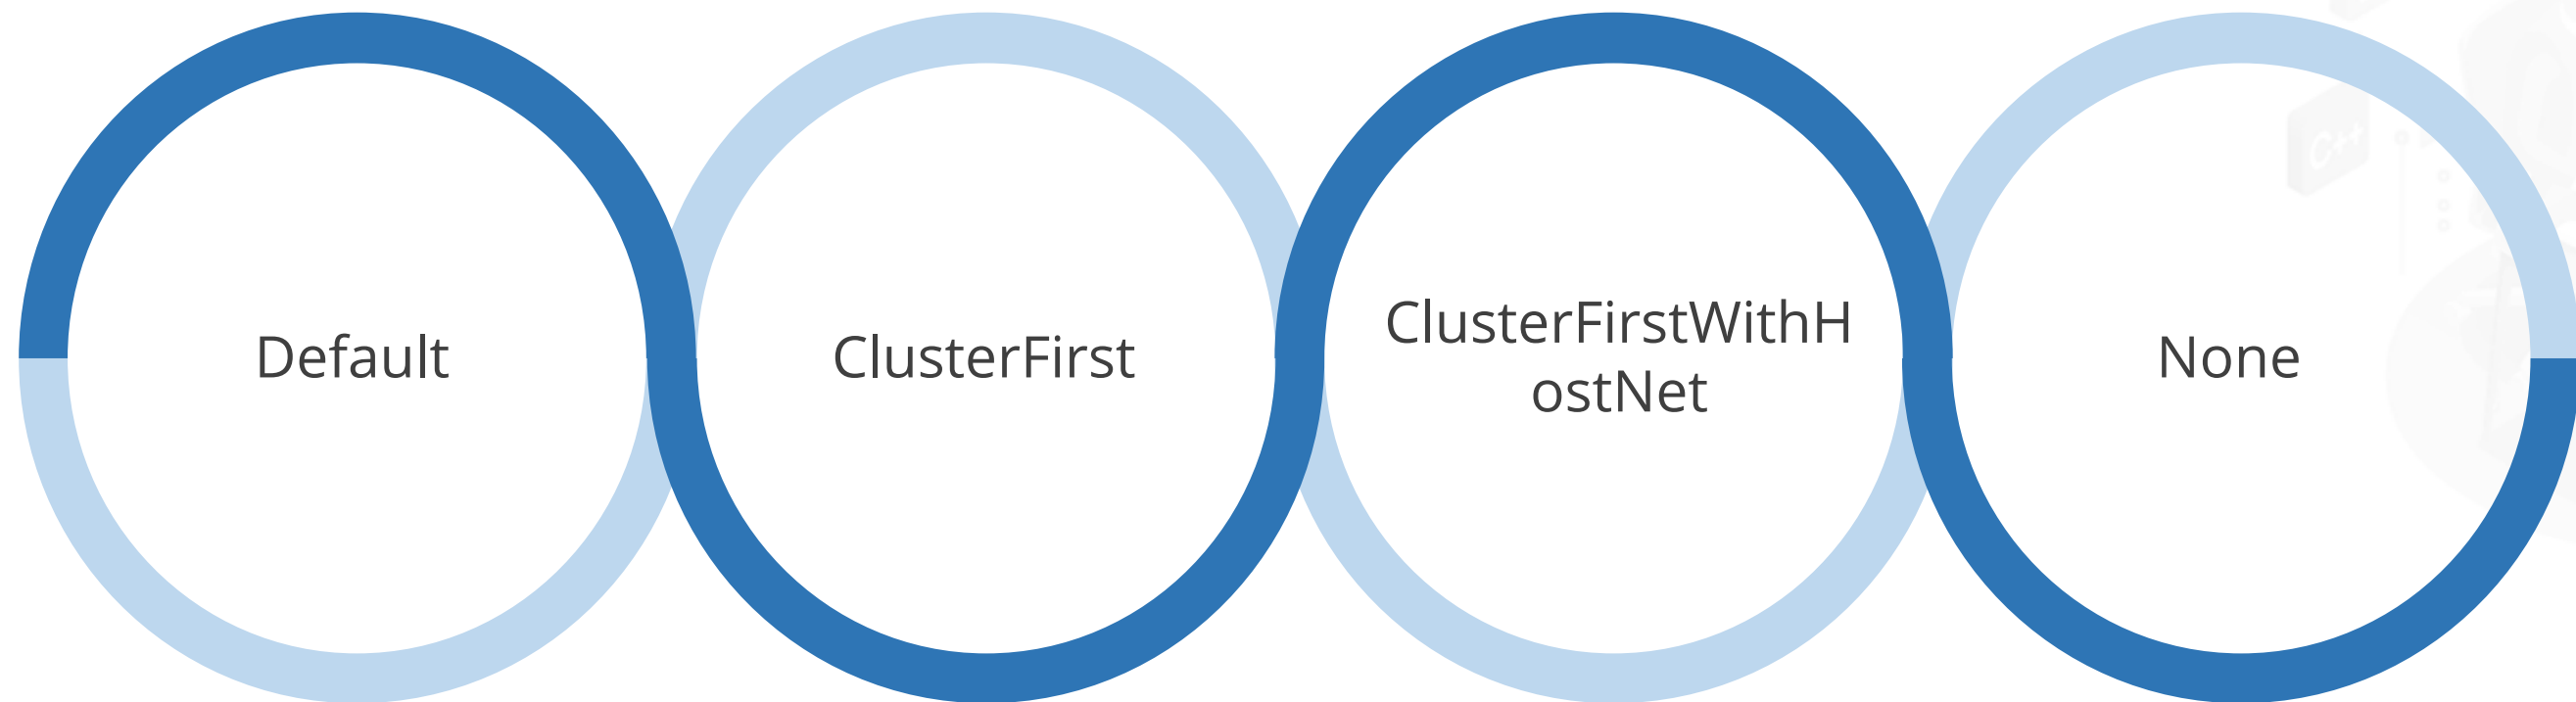


It also has an optional subdomain field that can be used to specify its subdomain.

Pod's DNS Policy

DNS policies can be set on a per-pod basis.

Kubernetes supports four pod-specific DNS policies:



Pod's DNS Policy

The given example shows a pod with its DNS policy set to **ClusterFirstWithHostNet** because it has **hostNetwork** set to **true**.

Demo

```
apiVersion: v1
Kind: Pod
metadata:
  name: busybox
  namespace: default
spec:
  containers:
  - image: busybox:1.28
    command:
      - sleep:
      - "3600"
    imagePullPolicy: IfNotPresent
    name: busybox
  restartPolicy: Always
  hostNetwork: true
  dnsPolicy: ClusterFirstWithHostNet
```



Pod's DNS Config

Pod's DNS config allows users to have more control over the DNS settings for a pod.

A user can specify three properties in the **dnsConfig** field:

nameservers

searches

options



Pod's DNS Config

Here is an example of a pod with custom DNS settings:

Demo

```
apiVersion: v1
Kind: Pod
metadata:
  namespace: default
  name: dns-example
spec:
  containers:
    - name: test
      image: nginx
  dnsPolicy: "None"
  dnsConfig:
    nameservers:
      - 1.2.3.4
    searches:
      - ns1.svc.cluater-domain.example
      - my.dns.search.suffix
    options:
      - name: ndots
        value: "2"
      - name: edns0
```



Configuring DNS for Kubernetes Services and Pods



Duration: 15 mins

Problem Statement:

You've been assigned a task To configure the Domain Name System (DNS) for Kubernetes services and pods to ensure proper network resolution and connectivity

ASSISTED PRACTICE

Assisted Practice: Guidelines

Steps to be followed:

1. Determine the default DNS in the cluster
2. Execute DNS query
3. Configure the DNS policy
4. Create a custom DNS configuration

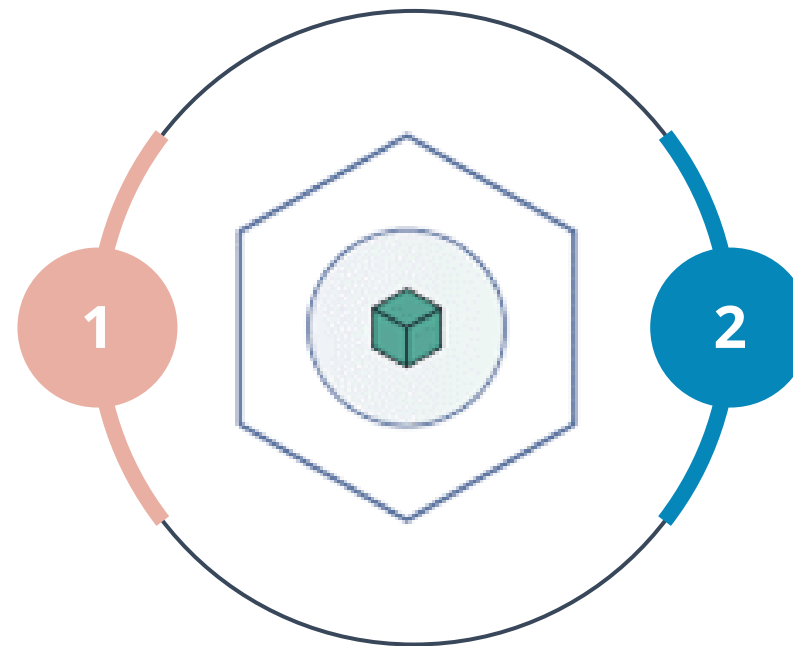


EndpointSlices

EndpointSlices

EndpointSlices provides a simple way to track network endpoints within a Kubernetes cluster.

They help to mitigate issues and provide an extensible platform for additional features such as topological routing.



They offer a more scalable and extensible alternative to endpoints.

EndpointSlice Resources

An EndpointSlice contains references to a set of network endpoints. Here is a sample EndpointSlice resource:

Demo

```
apiVersion: discovery.k8s.io/v1
kind: EndpointSlice
metadata:
  name: example-abc
  labels:
    kubernetes>io/servic-name: example
addressType: IPv4
ports:
  - name: http
    protocol: TCP
    port: 80
endpoints :
  - addresses:
    - "10.1.2.3"
    conditions:
      ready: true
    host name : pod-1
    nodeName: node-1
    zone: us-west2-a
```



Address Types and Conditions

EndpointSlices support three address types, namely, IPv4, IPv6, and Fully Qualified Domain Name (FQDN).

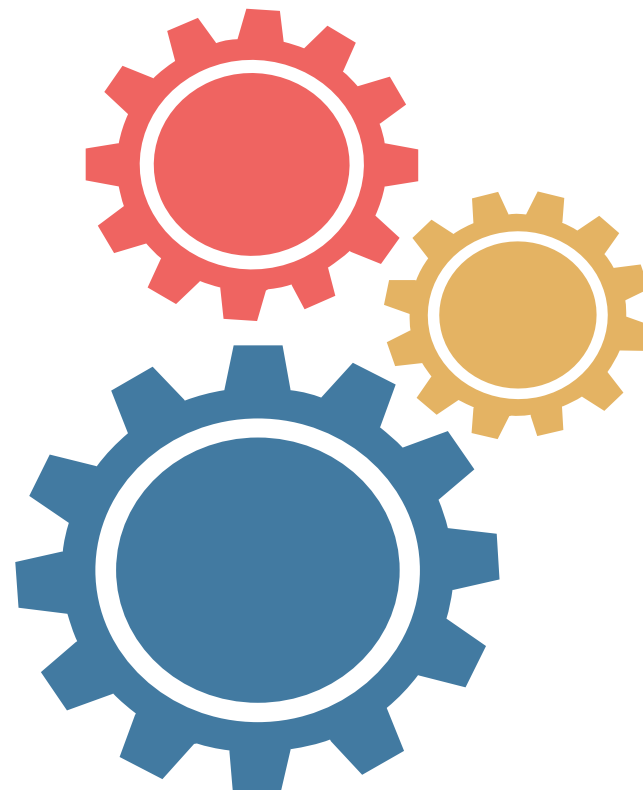
The EndpointSlice API stores three Endpoint conditions:

Ready

A condition that maps to a pod's ready condition

Terminating

A condition that indicates whether an endpoint is terminating



Serving

A condition that does not account for terminating states

Topology Information, Management, and Ownership

Topology enables a service to route traffic based on the node topology of the cluster.



Information

Topology information includes the location of the endpoint and information about the corresponding node and zone.



Management

EndpointSlice objects are often created and managed by the control plane (particularly, the EndpointSlice controller).



Ownership

EndpointSlices are owned by the service that the Endpointslice object tracks endpoints for.

EndpointSlice Mirroring

When applications create custom Endpoint resources, the cluster's control plane mirrors most endpoint resources to corresponding EndpointSlices. This helps to avoid concurrently writing to both endpoints and EndpointSlice resources.

The control plane mirrors endpoints resources unless:

- ✓ The endpoints resource has a **endpointslice.kubernetes.io/skip-mirror** label set to **true**.
- ✓ The endpoints resource has a **control-plane.alpha.kubernetes.io/leader** annotation.
- ✓ The corresponding service resource does not exist.
- ✓ The corresponding service resource has a non-nil Selector.

Distribution of EndpointSlices

To fill the EndpointSlices, the control plane does the following:

1

Iterate through existing EndpointSlices, remove endpoints that are no longer desired, and update matching endpoints that have changed

2

Iterate through EndpointSlices that have been modified in the first step and fill them up with any new endpoints needed

3

If there are new endpoints left to add, try to fit them into a previously unchanged slice or create new ones, or both

Duplicate Endpoints

Due to the nature of EndpointSlice changes, endpoints may be represented in more than one EndpointSlice at the same time.



This naturally occurs as changes to different EndpointSlice objects can arrive at the Kubernetes client watch or cache at different times.



Implementations using EndpointSlice must be able to have the endpoint appear in more than one slice.

Configuring EndpointSlice



Duration: 10 mins

Problem Statement:

You've been asked to configure the EndpointSlice to track network endpoints within a cluster

ASSISTED PRACTICE

Assisted Practice: Guidelines

Steps to be followed:

1. Create a deployment and identify its EndpointSlice
2. Create a YAML file for custom EndpointSlice configuration
3. Create a resource for the custom EndpointSlice configuration



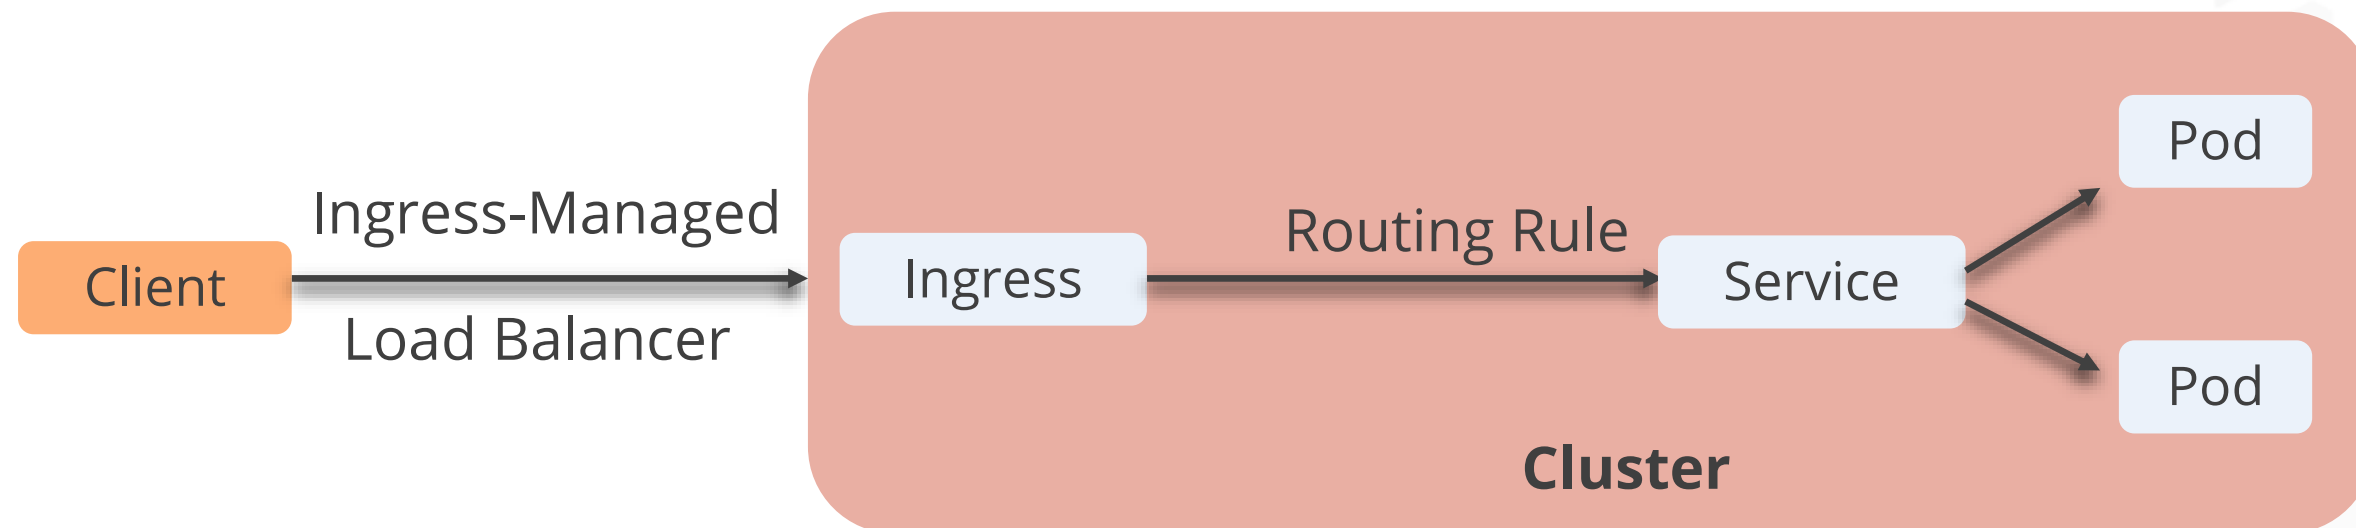
TECHNOLOGY

Ingress

What Is Ingress?

Ingress exposes HTTP and HTTPS routes from outside the cluster to services within the cluster.

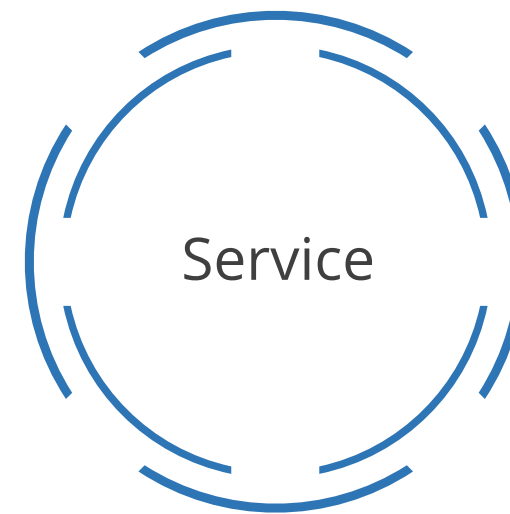
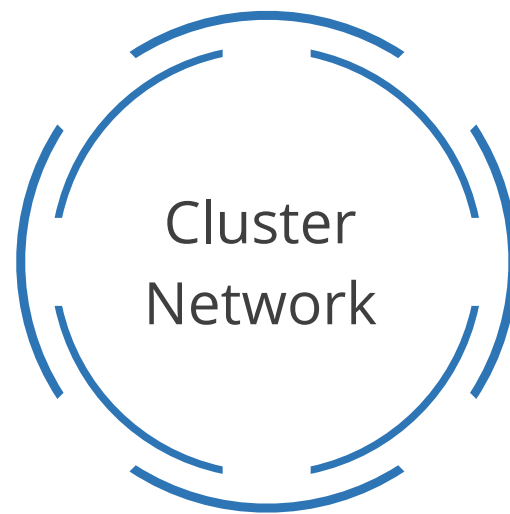
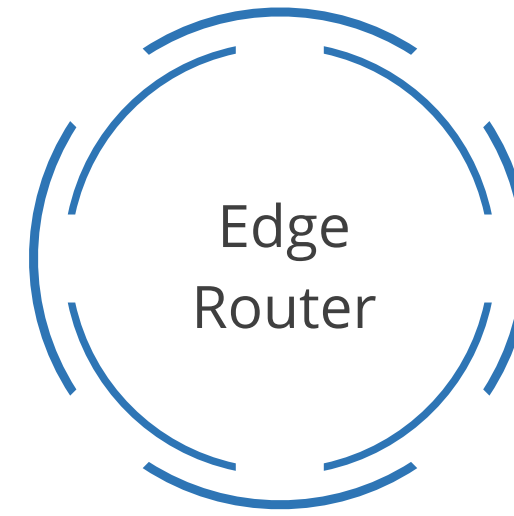
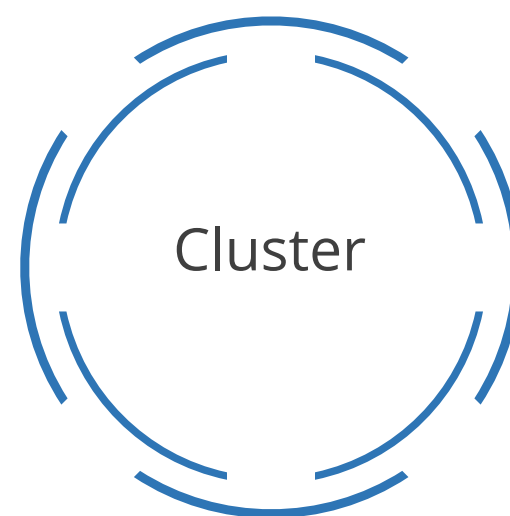
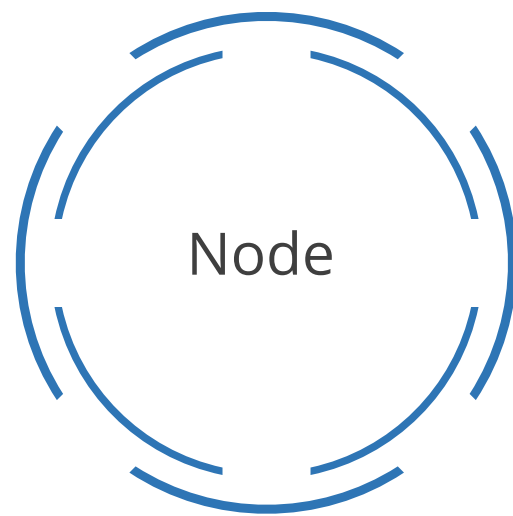
Here is a simple example where an Ingress sends all its traffic to one service:



An Ingress may be configured to give services externally reachable URLs and load balance traffic, terminate SSL/TLS, and offer name-based virtual hosting.

Terminologies Used in Ingress

Ingress uses the following terminologies:



Ingress Resource

An Ingress needs **apiVersion**, **kind**, **metadata** and **spec** fields.

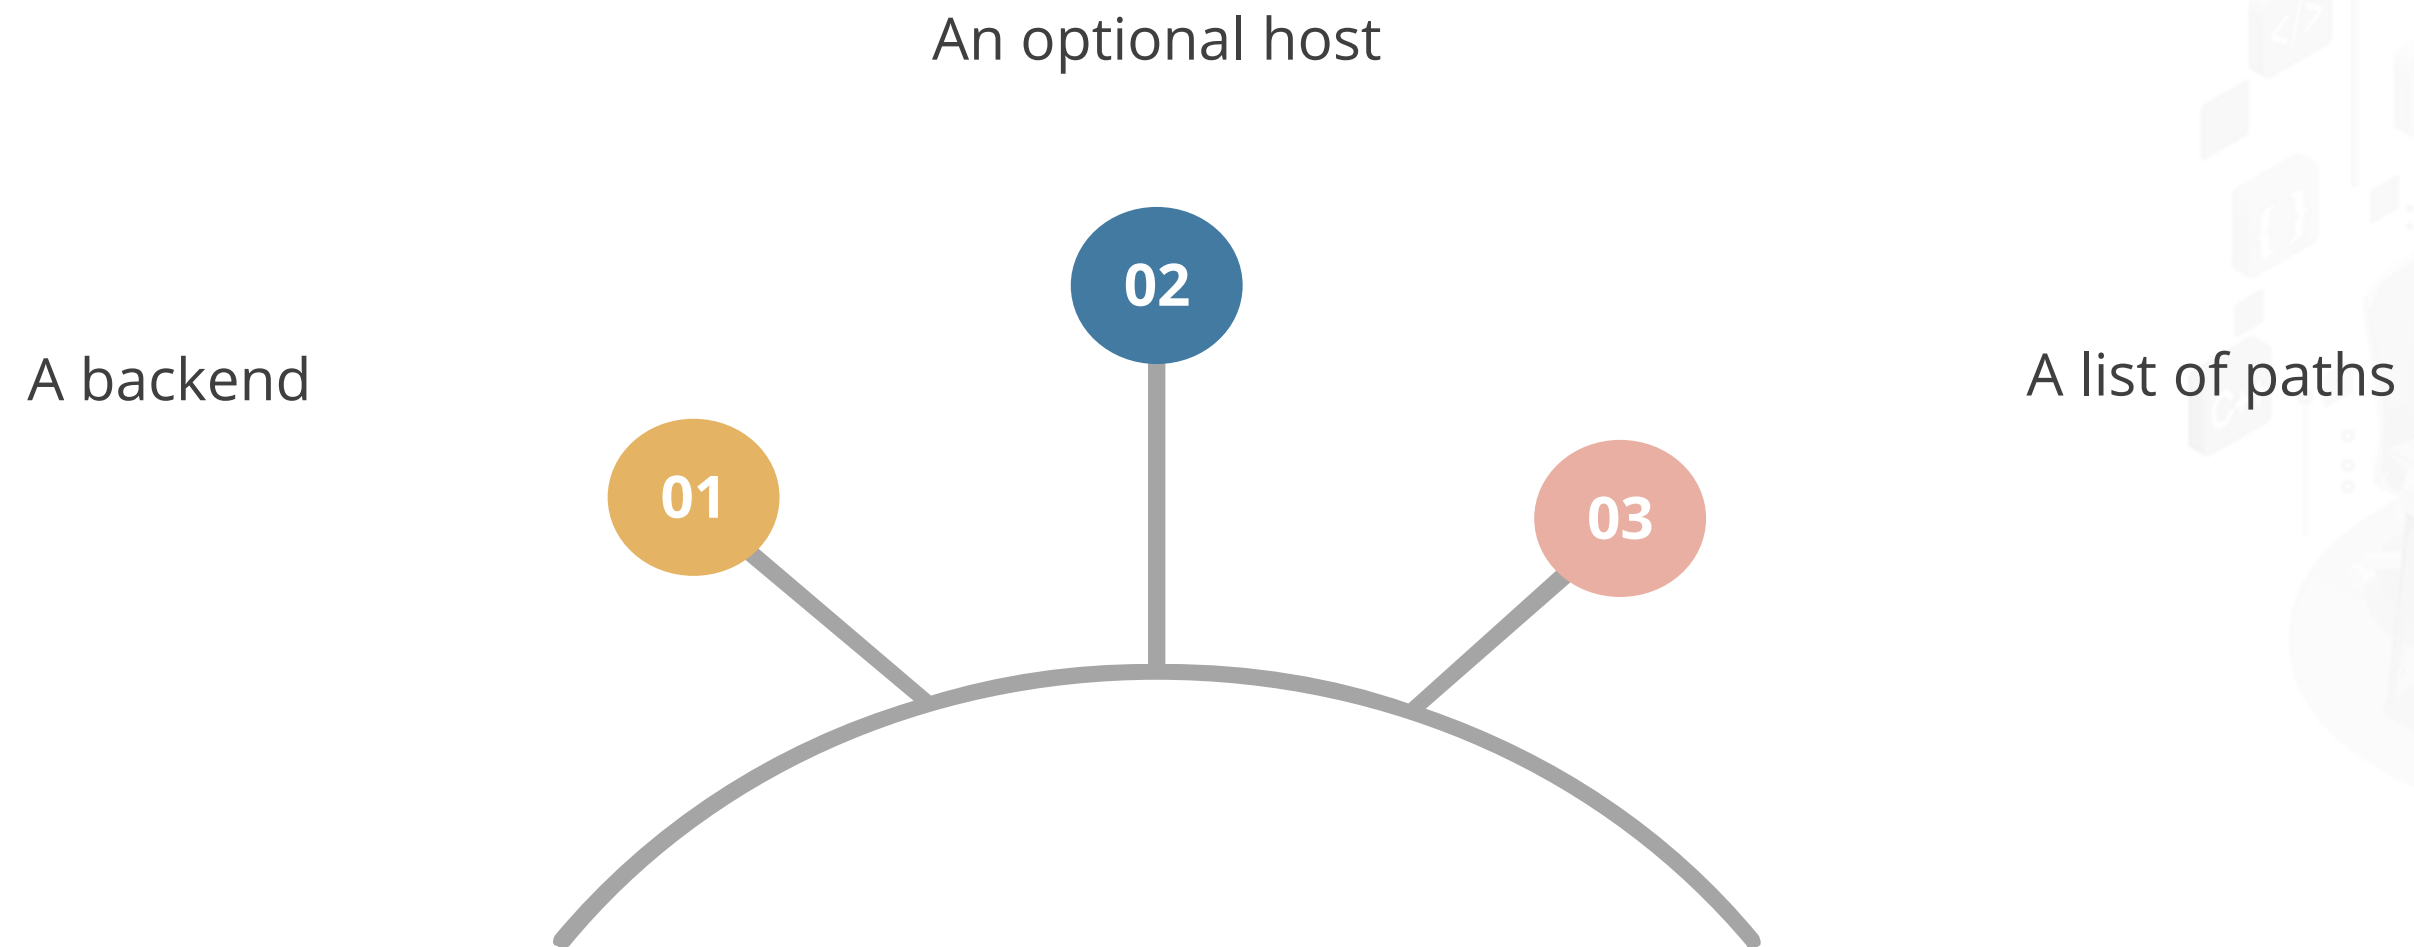
Example:

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: minimal-ingress
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  ingressClassName: nginx-example
  rules:
  - http:
      paths:
      - path: /testpath
        pathType: Prefix
        backend:
          service:
            name: test
            port:
              number: 80
```



Ingress Rules

Each HTTP rule contains the following information:



Default and Resource Backend

Default backend

The **defaultBackend** is conventionally a configuration option of the Ingress controller and is not specified in the Ingress resources.

Resource backend

A resource backend is an ObjectRef to another Kubernetes resource within the same namespace as the Ingress object.

Default Backend

An Ingress with no rules sends all traffic to a single default backend and is the backend that should handle requests.

Below are configuration to be set for backend:

If no **.spec.rules** are specified, **.spec.defaultBackend** must be specified.

If **defaultBackend** is not set, the handling of requests that do not match any of the rules will go to the ingress controller.

If the feature gate **MixedProtocolLBService** is enabled for the **kube-apiserver**, it can use different protocols when multiple ports are specified.

Resource Backend

A common usage for a resource backend is to ingress data to an object storage backend with static assets.

Demo

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: ingress-resource-backend
spec:
  defaultBackend:
    resource:
      apiGroup: k8s.example.com
      kind: StorageBucket
      name: static-assets
  rules:
    - http:
        paths:
          - path: /icons
            pathType: ImplementationSpecific
            backend:
              resource:
                apiGroup: k8s.example.com
                kind: StorageBucket
                name: icon-assets
```



View Ingress

To view the created Ingress, use the command shown below:

Demo

```
kubectl describe ingress ingress-resource-backend
```



Path Types

Each path in an Ingress must have a corresponding path type. There are three supported **path types**:

ImplementationSpecific

Matches with the **IngressClass**

Exact

Matches the URL path exactly along with case sensitivity

Prefix

Matches based on a URL path prefix split by the **/** separator



Hostname Wildcards

Precise matches require that the HTTP **host** header matches the **host** field. Wildcard matches require the HTTP **host** header to be equal to the suffix of the wildcard rule.

Hosts can be precise matches (for example **foo.bar.com**) or a wildcard (for example ***.foo.com**).

Host	Host header	Match
*.foo.com	bar.foo.com	Matches based on shared suffix
*.foo.com	baz.bar.foo.com	No match, wildcard only covers a single DNS label
*.foo.com	foo.com	No match, wildcard only covers a single DNS label

Hostname Wildcards

An example of the configuration file showing the use of wildcards:

Demo

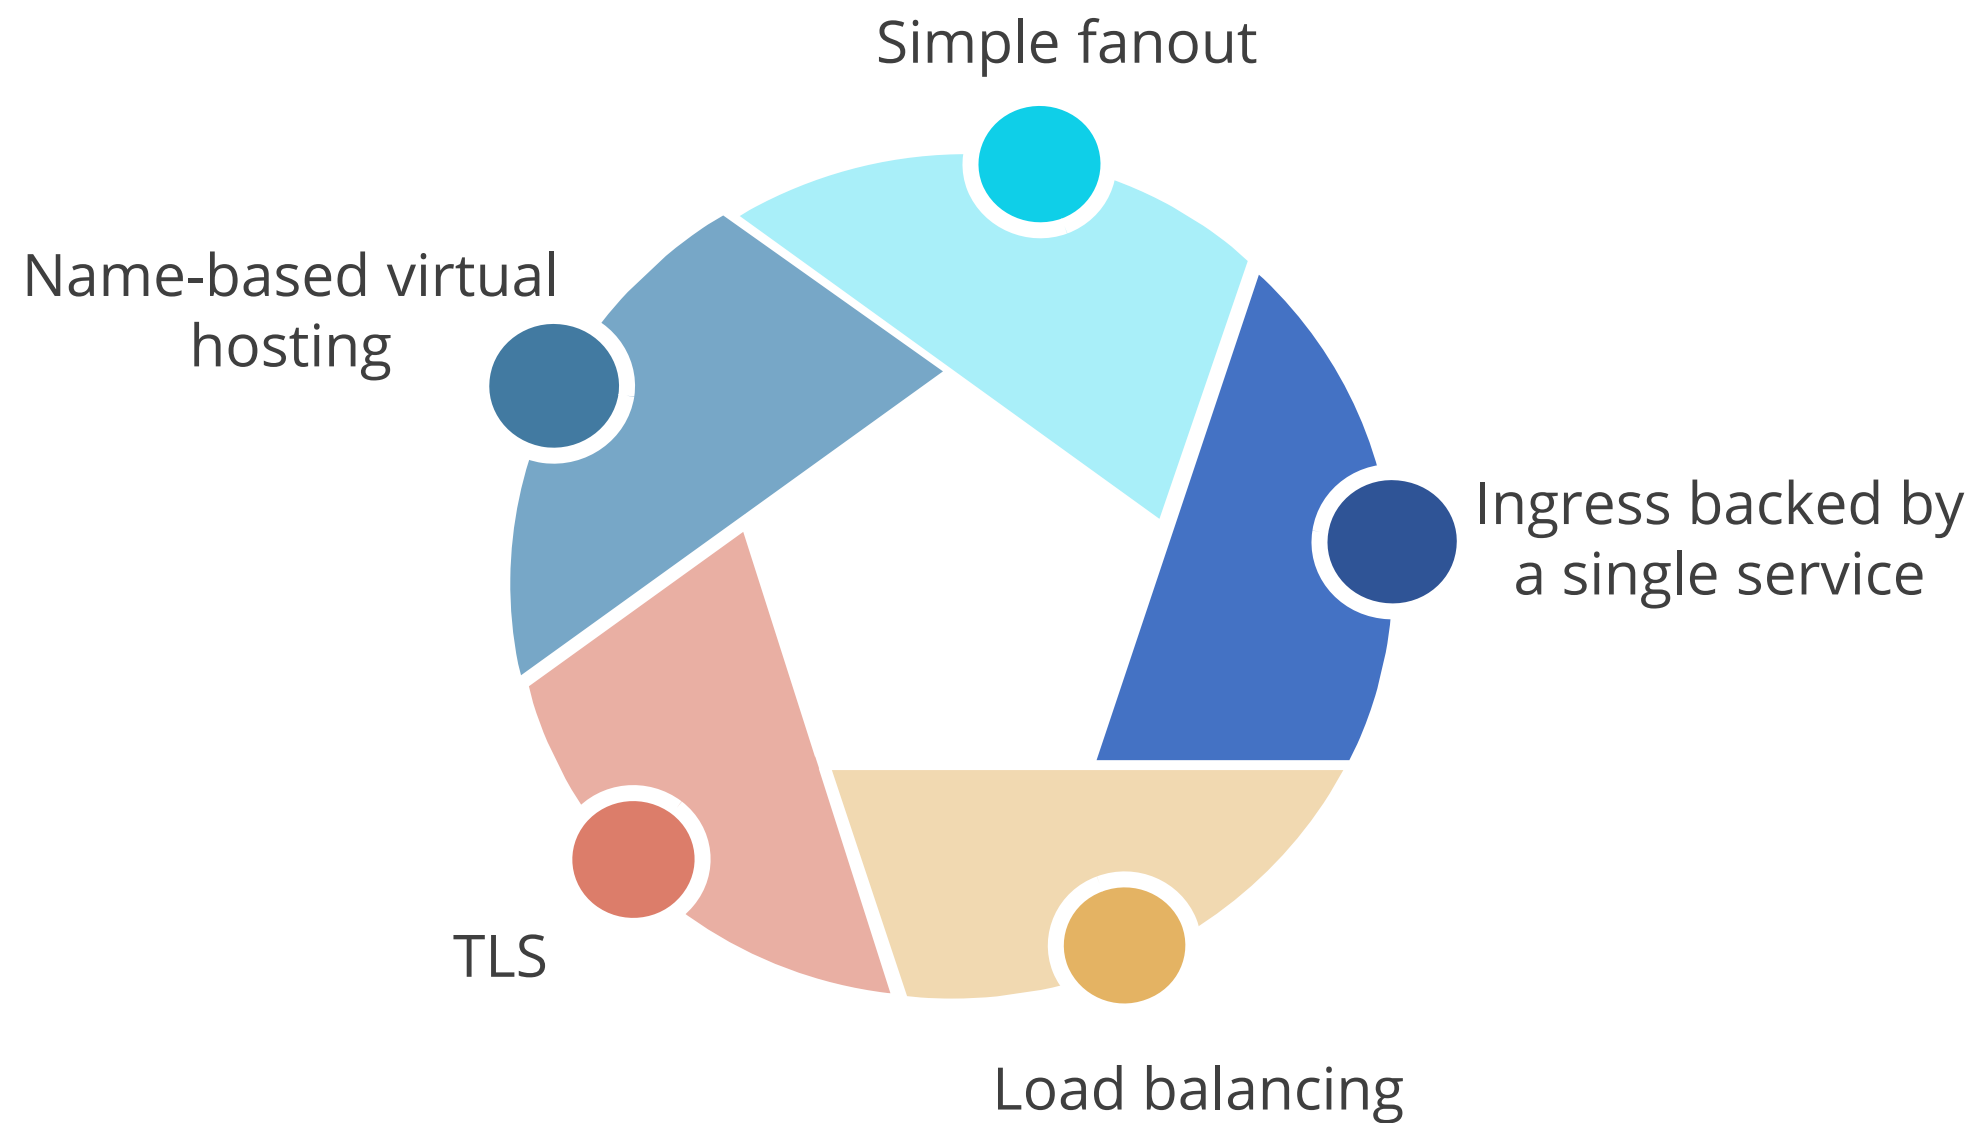
```
apiVersion: networking.k8s.10/v1
kind: Ingress
metadata:
  name: ingress-wildcard-host
spec:
  rules:
  - host: "foo.bar.com"
  http:
  paths:
  - pathtype: Prefix
    path: "/bar"
    backend:
      service:
        Name: service1
        port:
          number 80
  ...
```

Demo

```
...
- host: "*foo.com"
  http:
  paths:
  - pathtype: Prefix
    path: "/foo"
    backend:
      service:
        Name: service2
        port:
          number 80
```

Types of Ingress

In Kubernetes, there are five different types of ingress:



Ingress Backed by a Single Service

A single service can be exposed using an Ingress by specifying a default backend with no rules.

Demo

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: test-ingress
spec:
  defaultBackend:
    service:
      name: test
      port:
        number: 80
```

Demo

```
#Command to view the state of the created Ingress:
kubectl get ingress test-ingress
```

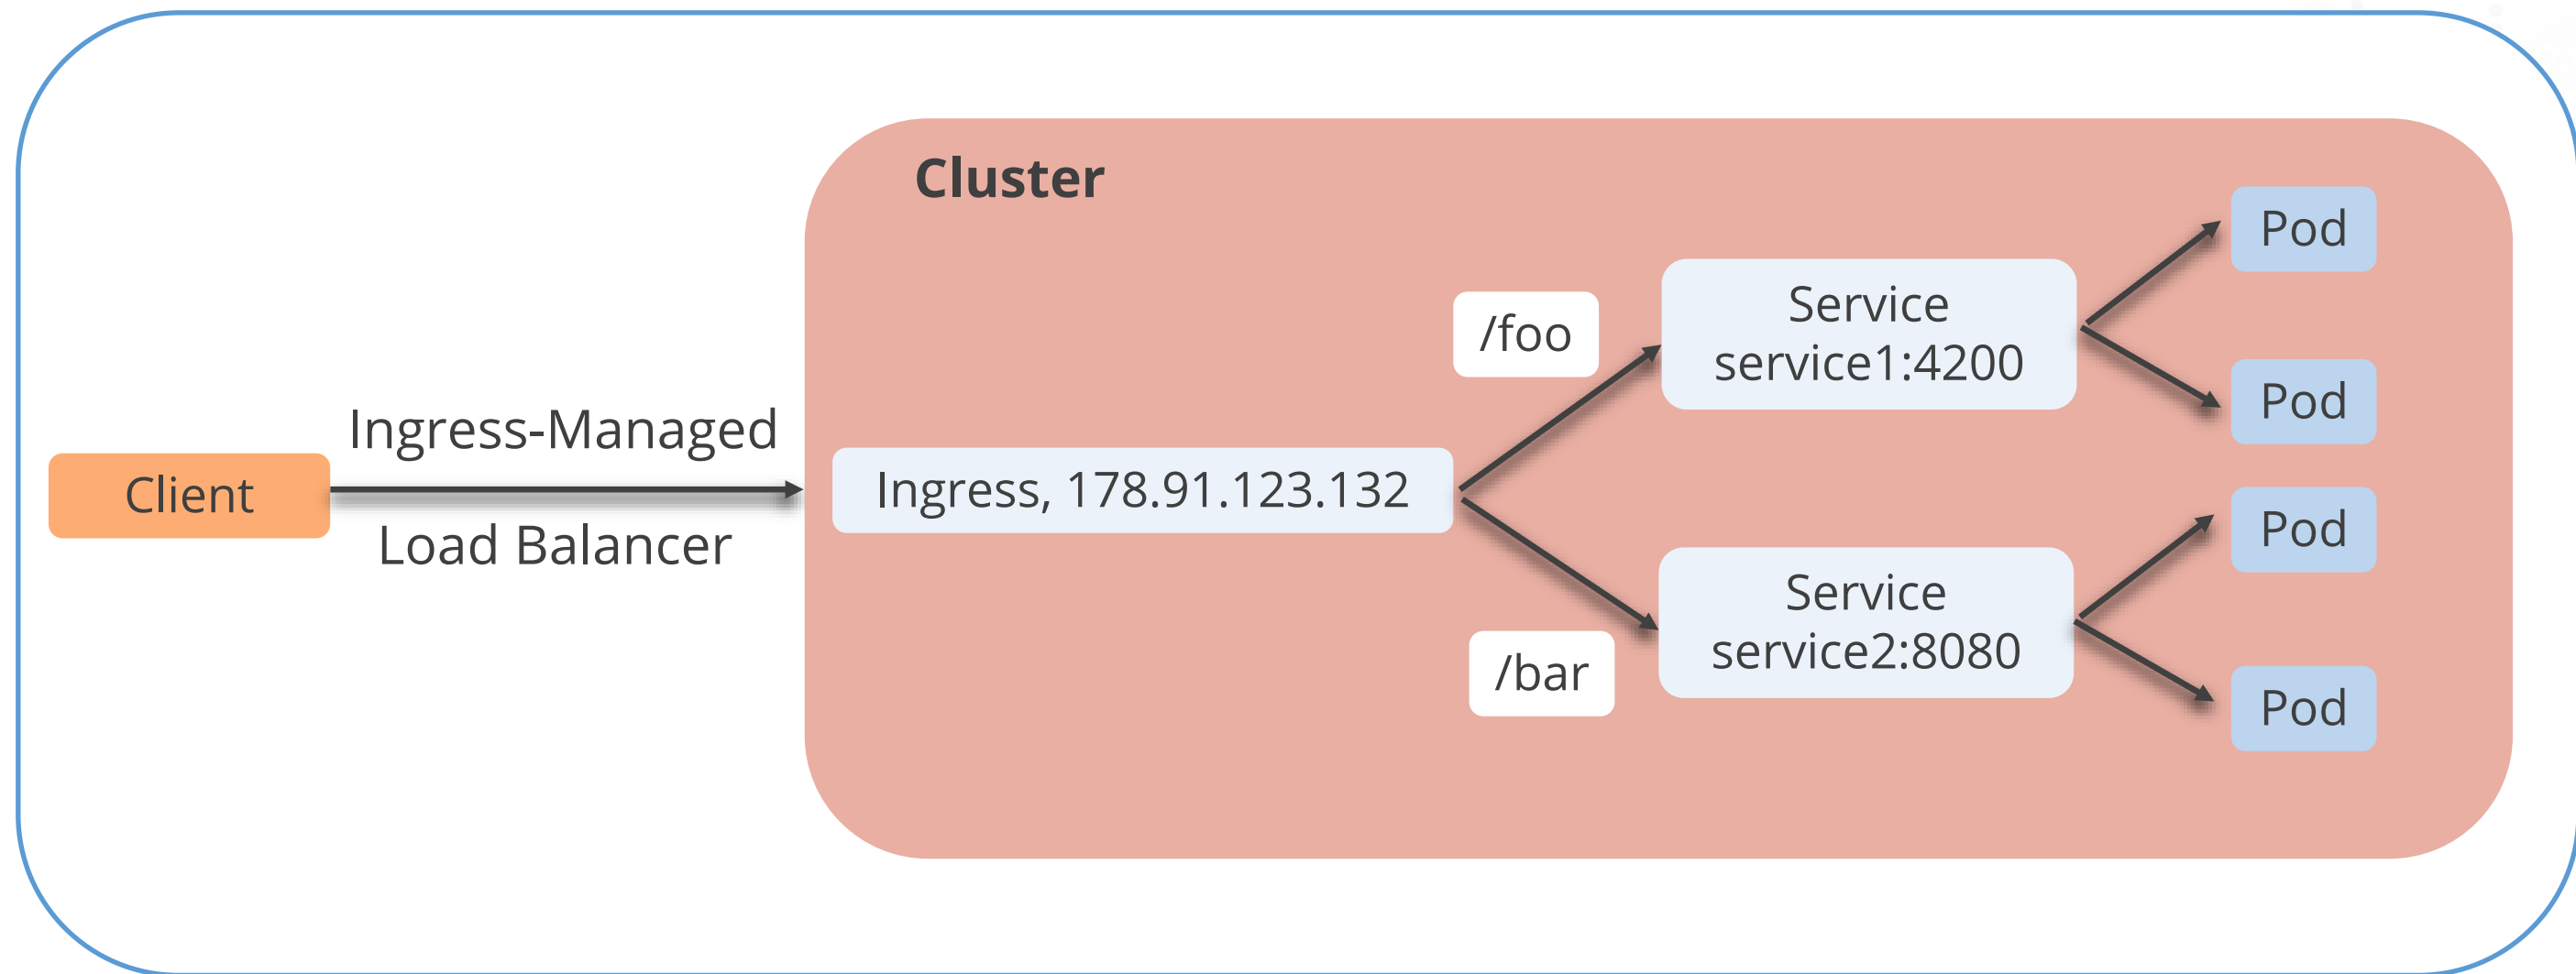
Result:

NAME	CLASS	HOSTS	ADDRESS	PORTS	AGE
test-ingress	external-lb	*	203.0.113.123	80	59s

Where 203.0.113.123 is the IP allocated by the Ingress controller to satisfy this Ingress.

Simple Fanout

A fanout configuration routes traffic from a single IP address to more than one service, based on the HTTP URI being requested.



Simple Fanout

A fanout would require an Ingress such as:

Demo

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: simple-fanout-example
spec:
  rules:
  - host: foo.bar.com
    http:
      paths:
      - path: /foo
        pathType: Prefix
        backend:
          service:
            name: service1
            port:
              number: 4200
      - path: /bar
        pathType: Prefix
        backend:
          service:
            name: service2
            port:
              number: 8080
```



Simple Fanout

The command to describe the simple fanout Ingress is as follows:

Demo

Command:

```
kubectl describe ingress simple-fanout-example
```

Result:

```
Name:                simple-fanout-example
Namespace:           default
Address:             178.91.123.132
Default backend:     default-http-backend:80 (10.8.2.3:8080)
```

Rules:

Host	Path	Backends
----	----	-----
foo.bar.com	/foo	service1:4200 (10.8.0.90:4200)
	/bar	service2:8080 (10.8.0.91:8080)

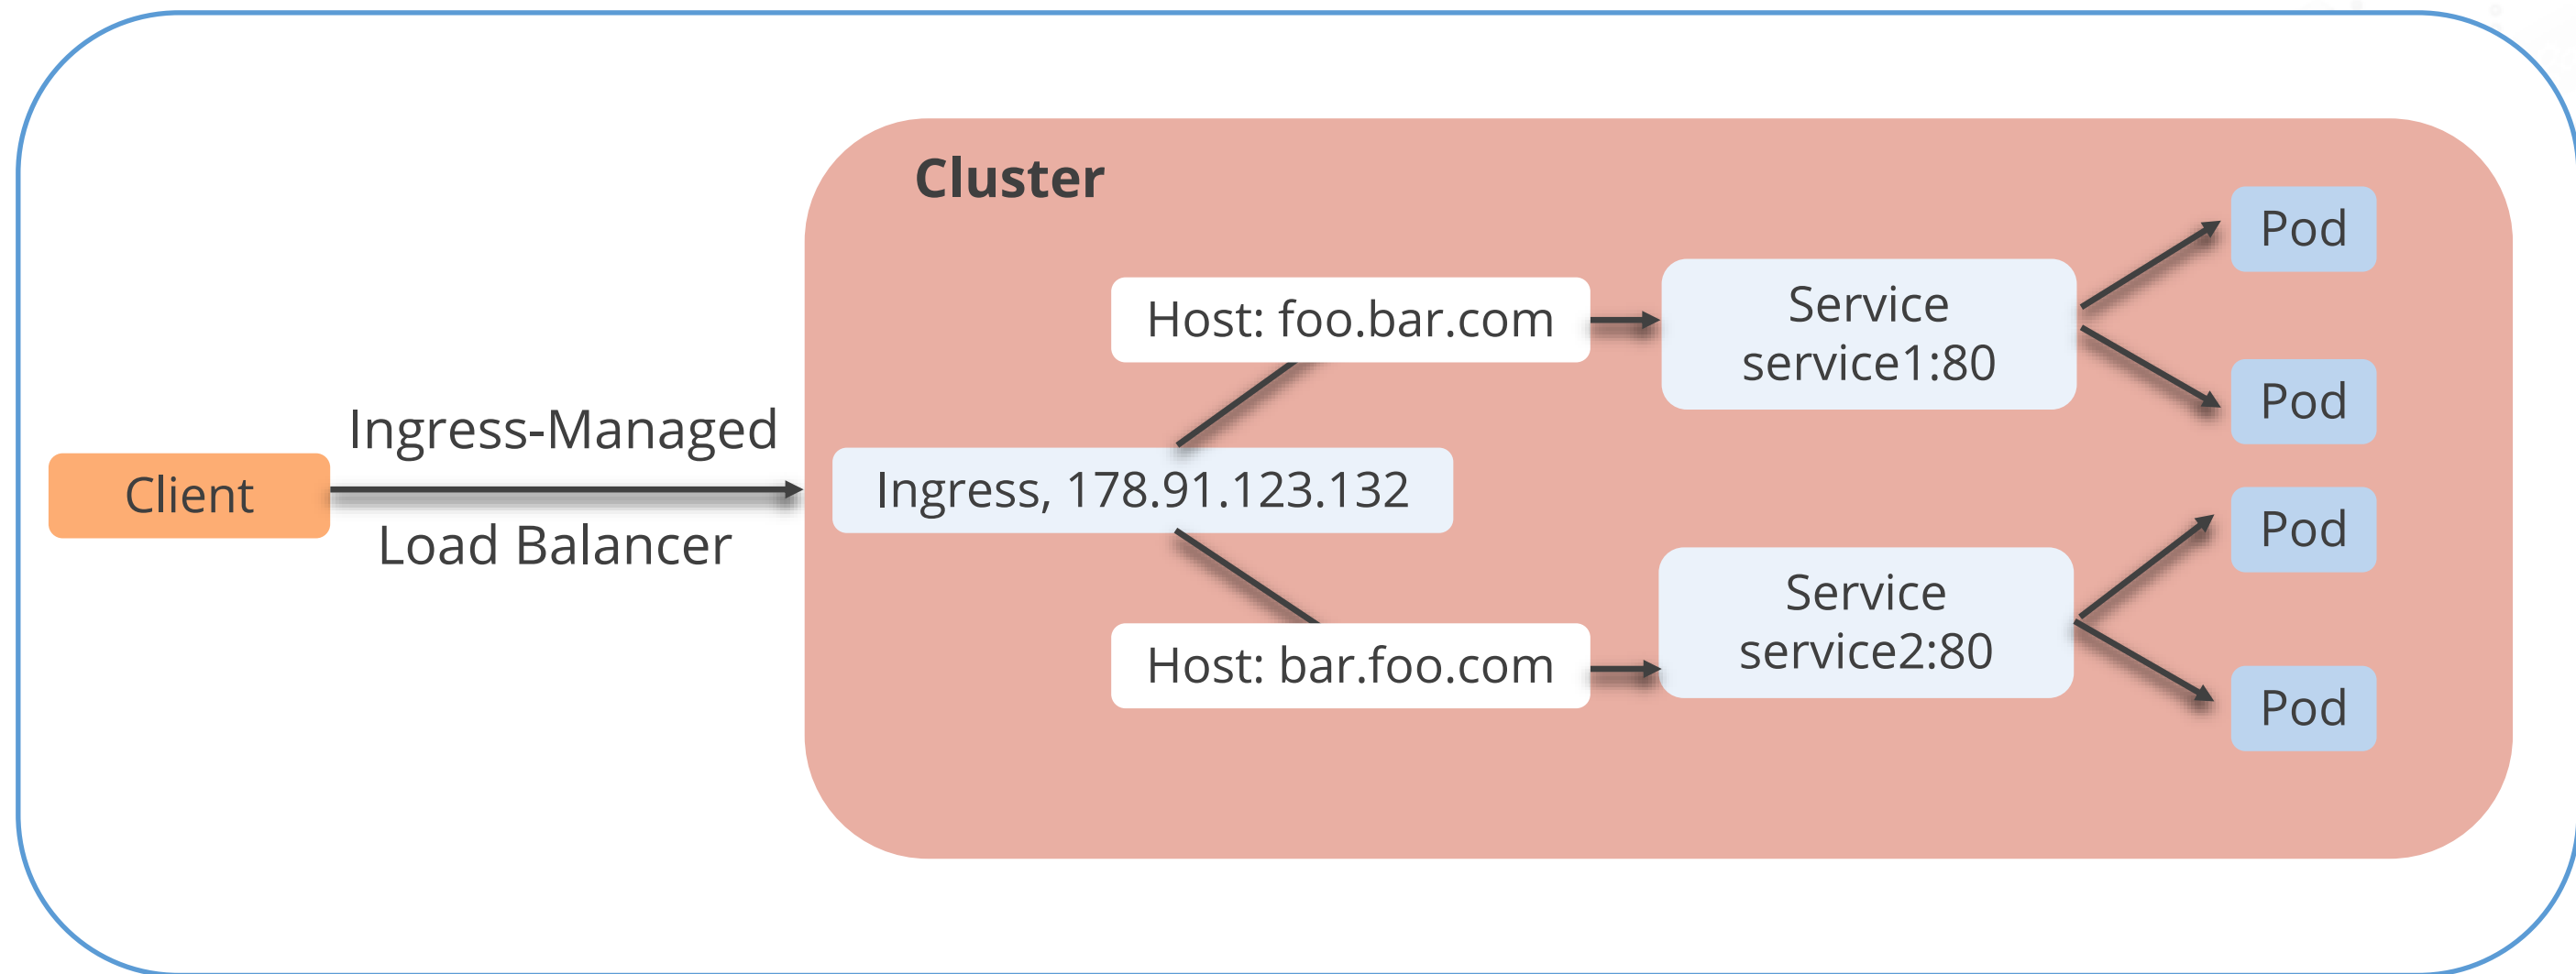
Events:

Type	Reason	Age	From	Message
----	-----	----	----	-----
Normal	ADD	22s	loadbalancer-controller	default/test



Name-Based Virtual Hosting

Name-based virtual hosts support routing HTTP traffic to multiple host names at the same IP address.



Name-Based Virtual Hosting

The Ingress shown below tells the backing load balancer to route requests based on the host header:

Demo

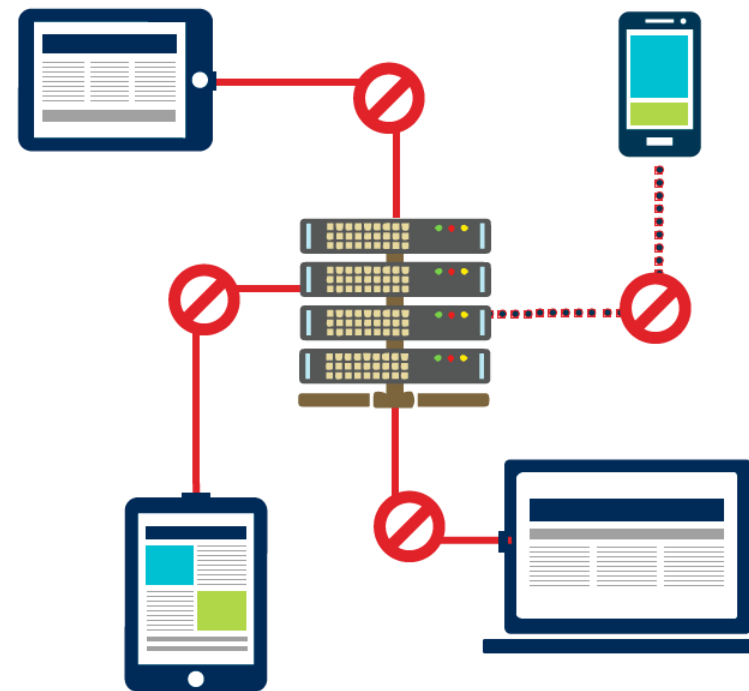
```
apiVersion: networking.k8s.io/v1
Kind: Ingress
Metadata:
  name: name-virtual-host-ingress
spec:
  rules:
  - host: foo.bar.com
    http:
      paths:
      - pathType: Prefix
        path: "/"
        backend:
          service:
            Name: service1
            port:
              number 80
  ...
```

Demo

```
...
- host: bar.foo.com
  http:
    paths:
    - pathType: Prefix
      path: "/"
      backend:
        service:
          Name: service2
          port: number 80
```

TLS

Transport Layer Security (TLS) is a cryptographic protocol designed to provide communications security over a computer network.



TLS

An Ingress can be secured using the configuration shown below:

Demo

```
apiVersion: v1
kind: Secret
metadata:
  name: testsecret-tls
  namespace: default
  data:
    tls.crt: base64 encoded cert
    tls.key: base64 encoded key
type: kubernetes.io/tls
```

The **tls.crt** and **tls.key** contain the certificate and private key to be used for TLS.



TLS

Ensure that the TLS secret created has come from a certificate that contains a Common Name (CN), also known as a Fully Qualified Domain Name (FQDN) for https-example.foo.com.

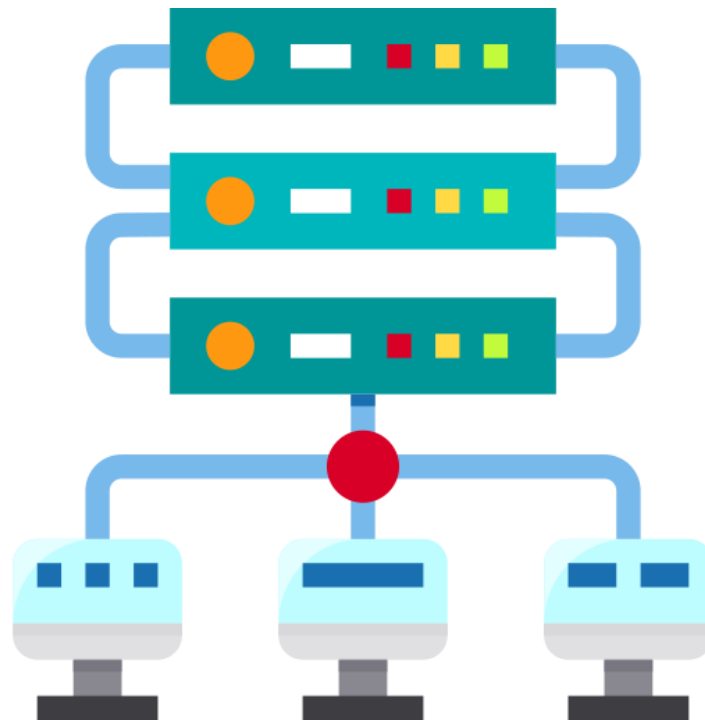
Demo

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: tls-example-ingress
spec:
  tls:
  - host:
    - https-example.foo.com
  rules:
  - host: bar.foo.com
    http:
      paths:
      - path: /
        - pathtype: Prefix
        backend:
          service:
            Name: service2
            port:
              number 80
```



Load Balancing

Load balancing settings are bootstrapped to an Ingress controller. These settings can thus apply to all Ingresses.



Ingress does not expose health checks directly. They can be exposed using readiness probes.



Update an Ingress

To update an existing Ingress to be added to a new host, edit the resource as shown:

Demo

Command:

```
kubectl describe ingress test
```

Result:

```
Name:          test
Namespace:     default
Address:       178.91.123.132
```

Demo

```
spec:
  rules:
  - host: foo.bar.com
    https:
      paths:
      - backend
        service:
          Name: service1
          port:
            number 80
          path: /foo
  - pathtype: Prefix
    host: foo.bar.com
    https:
      paths:
      - backend
        service:
          Name: service2
          port:
            number 80
          path: /foo
          pathtype: Prefix
```

Setting up Ingress Controller with Transport Layer Security



Duration: 20 mins

Problem Statement:

You've been asked to implement the transport layer security by deploying an Ingress rule to generate an SSL certificate

ASSISTED PRACTICE

Assisted Practice: Guidelines

Steps to be followed:

1. Deploy Ingress
2. Deploy HTTPD and OpenShift
3. Generate a self-signed SSL certificate and a TLS certificate
4. Verify the Ingress rule



Ingress Controllers

Overview

For an Ingress resource to function, the cluster must have an active Ingress controller running.

Kubernetes as a project supports and maintains three controllers:



AWS Ingress
Controllers



Nginx Ingress
Controllers



GCE Ingress
Controllers



Additional Controllers

These are the third-party controllers that give Kubernetes the features it requires:

1

AKS Application Gateway Ingress Controller for AZure

2

Ambassador API Gateway

3

Apache APISIX Ingress Controller

4

Avi Kubernetes Operator for Load Balancing

5

Citrix Ingress Controller

6

Contour

7

EnRoute

8

Gloo, an open-source Ingress Controller based on Envoy

Additional Controllers

These are the third-party controllers that give Kubernetes the features it requires:

9

HAProxy Ingress is an ingress

10

The HAProxy Ingress Controller

11

Skipper HTTP Router and Reverse Proxy

12

The Kong Ingress Controller

13

portworxVolume

14

Istio Ingress

15

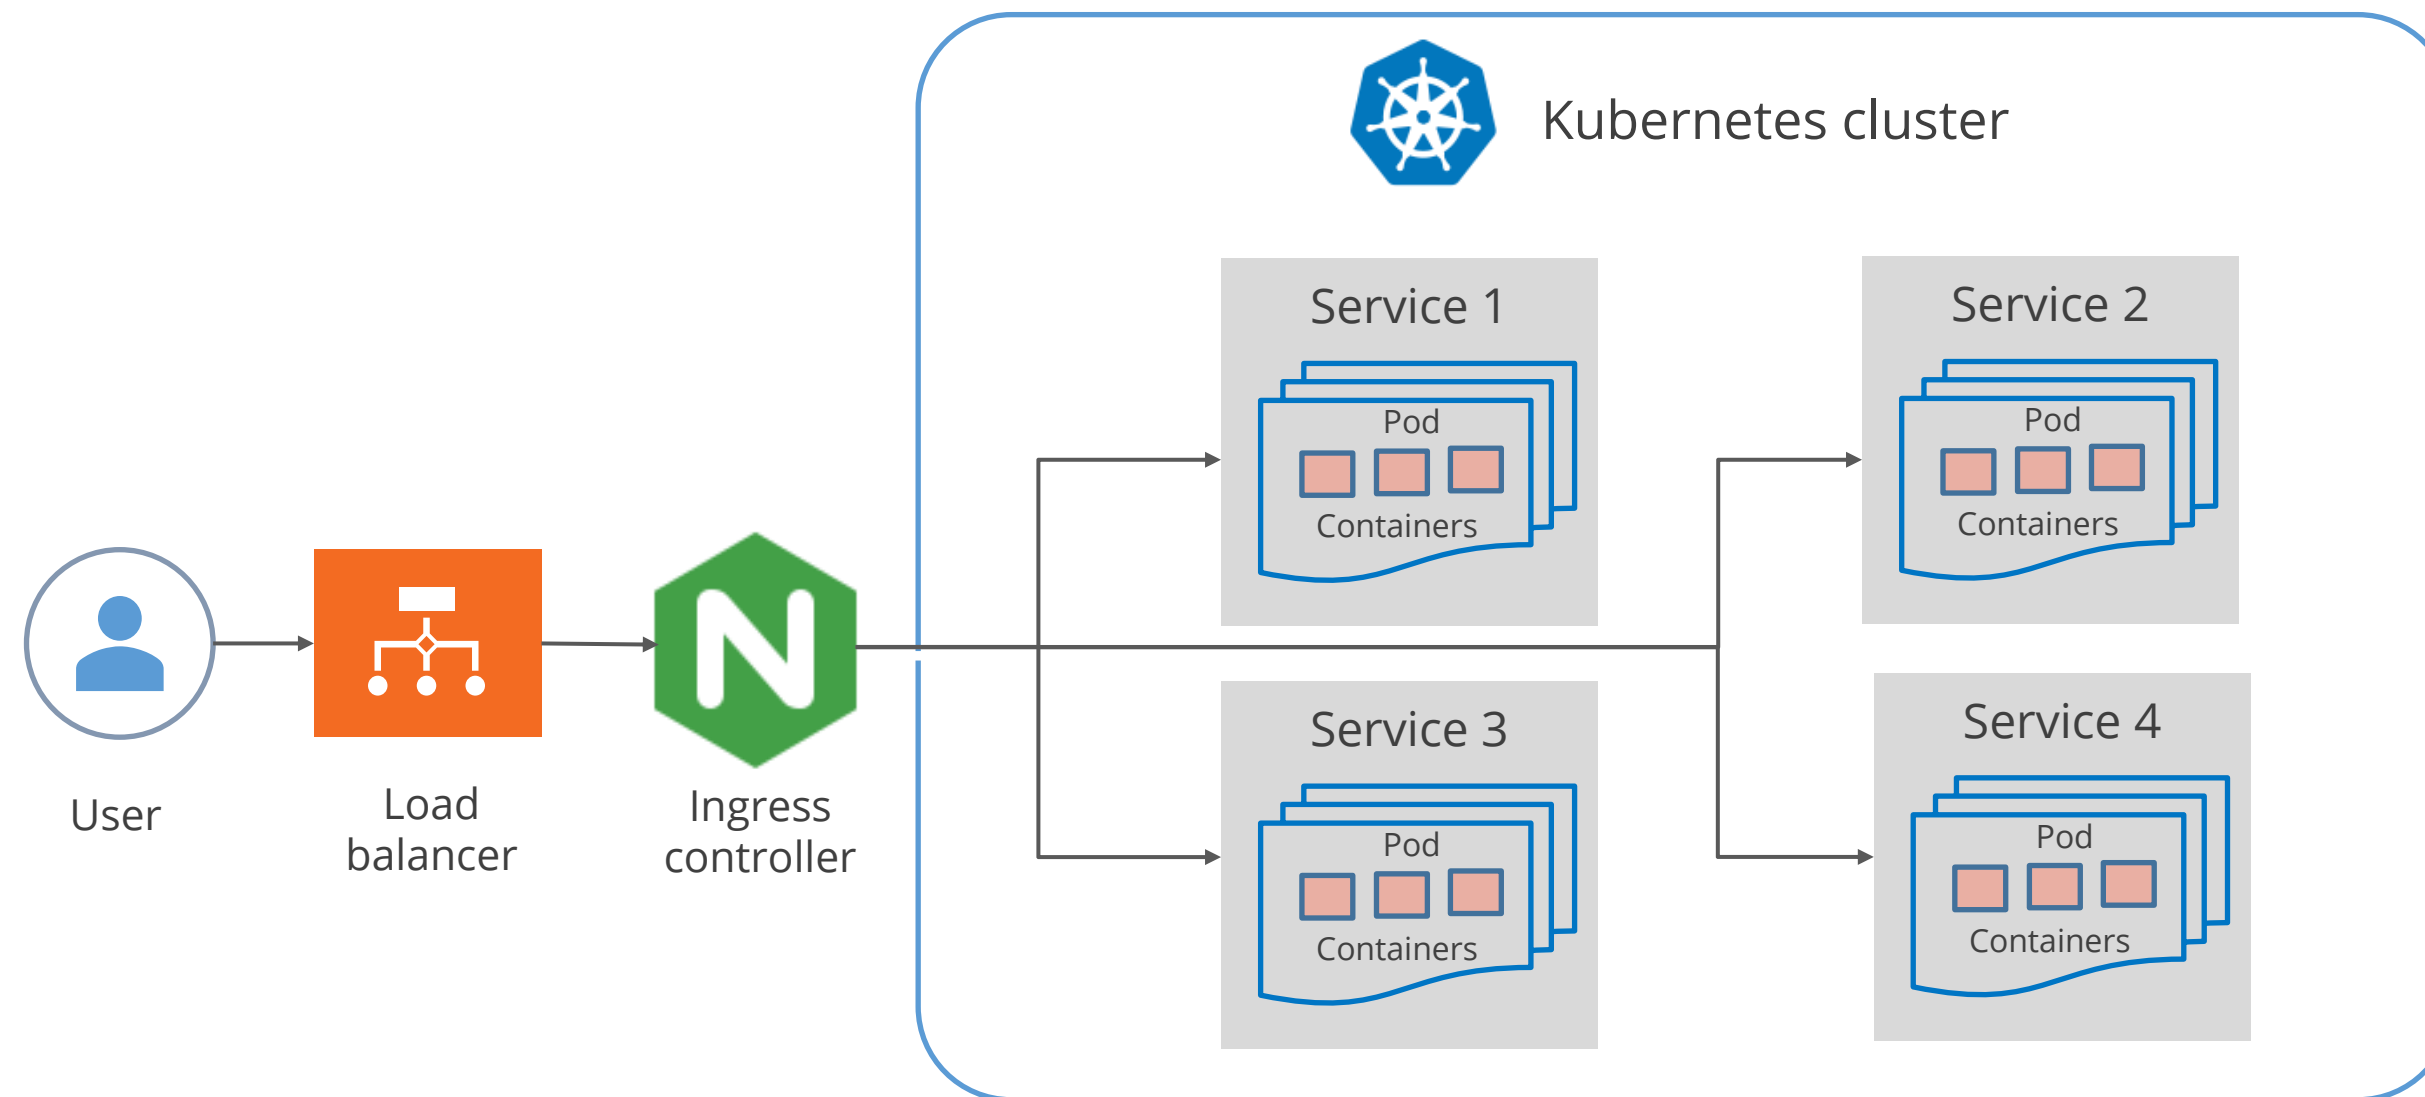
Voyager

16

The Traefik Kubernetes Ingress Provider

Using Ingress Controllers and Rules to Manage Network Traffic

The load balancer, as shown below, handles traffic distribution over several clusters, while the clusters contain Ingress controllers to ensure equal distribution to the services.



Using Ingress Controllers and Rules to Manage Network Traffic

Users can deploy any number of **Ingress controllers** using **IngressClass** within a cluster.

When creating an Ingress, the **ingressClassName** field on the Ingress object must be specified using **.metadata.name**.

If an Ingress does not have an **IngressClass** specified and the cluster has precisely one **IngressClass** marked as default, Kubernetes applies the cluster's default **IngressClass** to the Ingress.

By adding the text value **true** to the **ingressclass.kubernetes.io/is-default-class** annotation on the **IngressClass**, it can be declared as default.

Using Ingress Controllers and Rules to Manage Network Traffic

Ingress controller features:

- ✓ Accept traffic from outside the Kubernetes platform and distribute it to pods within the platform
- ✓ Handle **Egress** traffic within a cluster for services that need to connect with other services outside of the cluster
- ✓ Deploy objects called **Ingress Resources** via the Kubernetes API
- ✓ Monitor the pods running in Kubernetes and automatically adjust the loadbalancing rules when pods are added or withdrawn from a service

Setting up an Ingress



Duration: 15 mins

Problem Statement:

You've been asked to create an Ingress controller on a Kubernetes cluster.

ASSISTED PRACTICE

Assisted Practice: Guidelines

Steps to be followed:

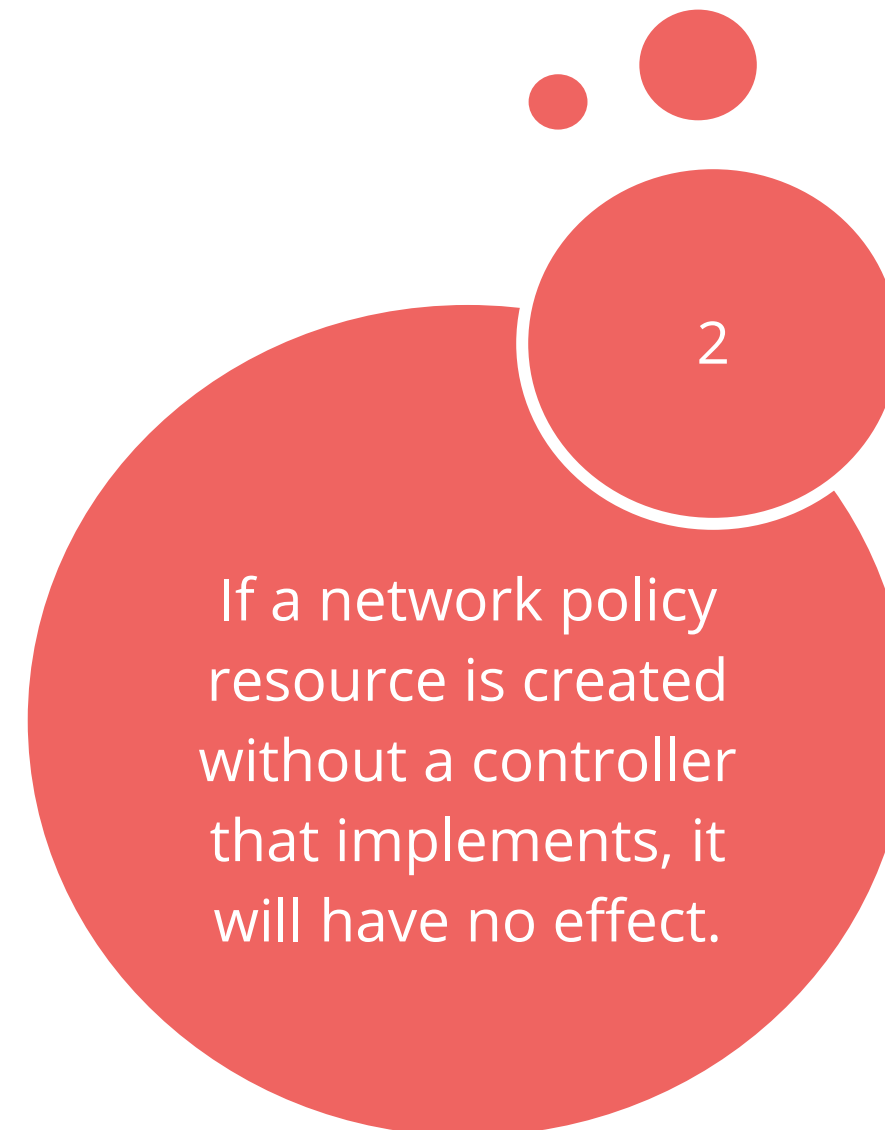
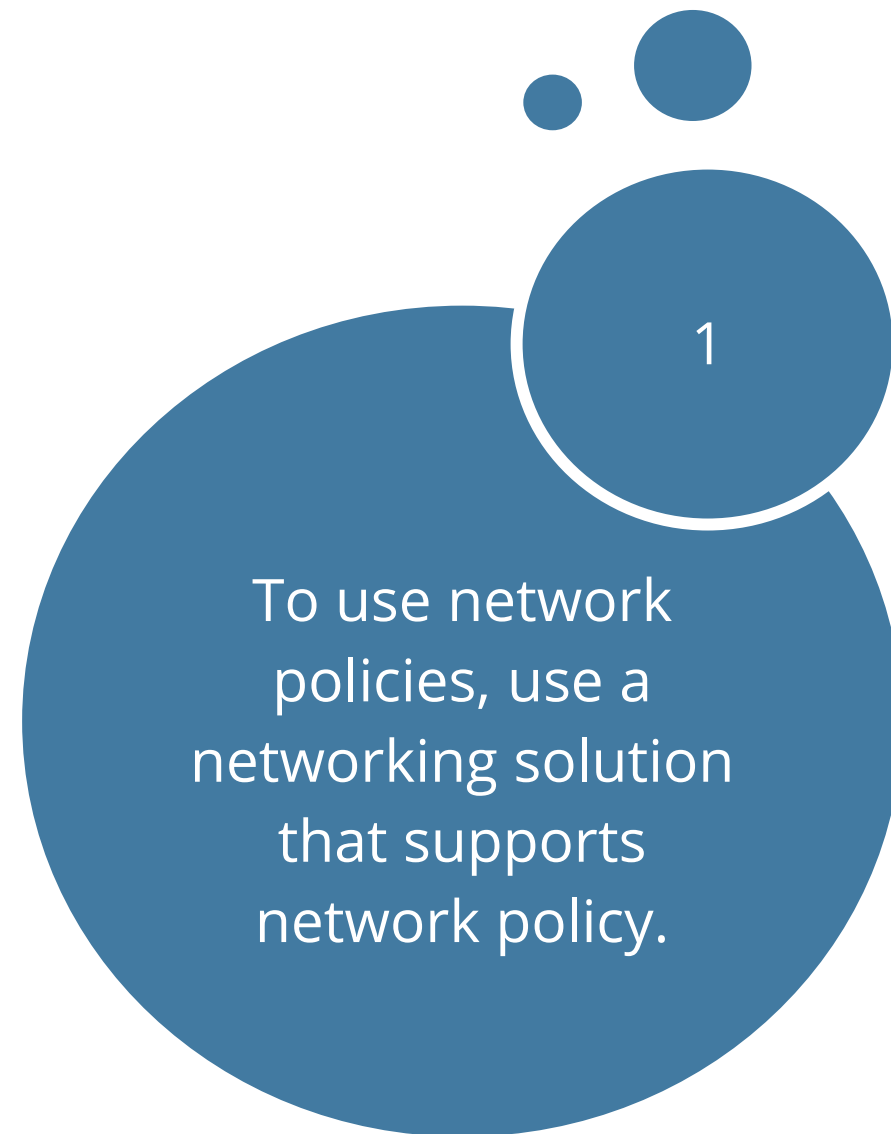
1. Creating an Nginx Ingress controller
2. Adding a rule
3. Verifying the Ingress



Network Policies

Network Policies

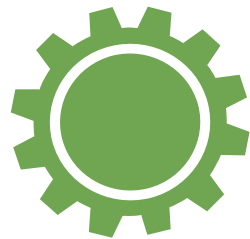
Network policies are implemented by the network plugin.



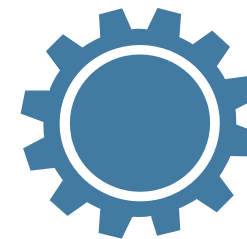
Introduction

Network policies are application-centric constructs that enable specifying how a pod is allowed to communicate with various network entities over the network.

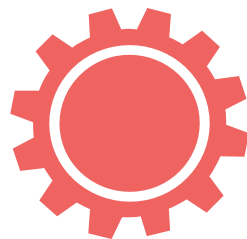
The entities that a pod can communicate with are identified through a combination of three identifiers:



Other pods that are
allowed



Namespaces that are
allowed



IP blocks

Isolated and Non-isolated Pods

By default, pods are non-isolated; they accept traffic from any source.

When there is a NetworkPolicy in a namespace selecting a particular pod, the pod will reject all connections not allowed by the NetworkPolicy.

If a policy (or policies) selects a pod, the pod is restricted to what is allowed by the union of the Ingress/Egress rules of the policy (or policies).

For a network flow to happen between two pods, both the Egress policy on the source pod and the Ingress policy on the destination pod should allow the traffic.

Network Policy Resource: Example

An example of **NetworkPolicy** might look like this:

Demo

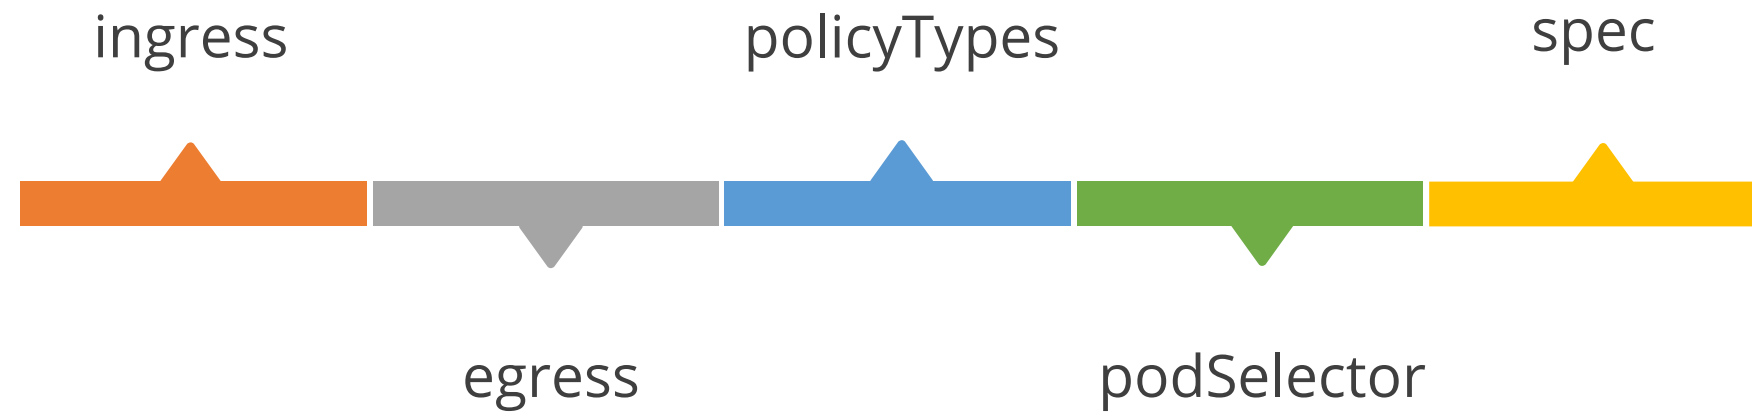
```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: test-network-policy
  namespace: default
spec:
  podselector:
    matchLabels:
      role: db
  policyTypes:
    - Ingress
    - Egress
  ingress:
    - from:
      - ipBlock:
          cidr: 172.17.0.0/16
        except:
          - 172.17.1.0/24
      ...
```

Demo

```
- namespaceSelector:
    matchLabels:
      project: myproject
- podSelector:
    matchLabels:
      role: frontend
ports:
  - protocol: TCP
    port: 6379
egress:
  - to:
    - ipBlock:
        cidr: 10.0.0.0/24
    ports:
      - protocol: TCP
        port: 5978
```

Network Policy Resource

The configuration file has six attributes:



Behavior of To and From Selectors

There are four kinds of selectors that can be specified in an **ingress from** section or **egress to** section:

namespaceSelector and podSelector

A single **to** or **from** entry, which specifies both **namespaceSelector** and **PodSelector**, selects specific **pods** within particular **namespaces**.

podSelector

This selects specific **pods** in the same namespace as the **NetworkPolicy** which should be allowed as ingress sources or egress destinations.

Behavior of To and From Selectors

There are four kinds of selectors that can be specified in an **ingress from** section or **egress to** section:

namespaceSelector

This selects particular **namespaces** for which all **Pods** should be allowed as ingress sources or egress destinations.

ipBlock

This selects particular IP CIDR ranges to allow as ingress sources or egress destinations.

Behavior of To and From Selectors

A single **from** element in **namespaceSelector** and **podSelector** allows connections from pods with the label **role=client** in namespaces with the label **user=alice**.

Demo

```
...
  ingress:
  - from:
    namespaceSelector:
      matchLabels:
        user: alice
    podSelector:
      matchLabels:
        role: client
...
```



Default Policies

By default, if no policies exist in a namespace, all ingress and egress traffic is allowed to and from pods in that namespace.

```
Demo
...
  apiVersion: networking.k8s.10/v1
  kind: NetworkPolicy
  metadata:
    name: default-deny-ingress
  spec:
    podselector: {}
    policyTypes:
      - Ingress
```

Default **deny all** Ingress traffic

```
Demo
...
  apiVersion: networking.k8s.10/v1
  kind: NetworkPolicy
  metadata:
    name: allow-all-ingress
  spec:
    podselector: {}
    ingress:
      - {}
    policyTypes:
      - Ingress
```

Default **allow all** Ingress traffic



Default Policies

By default, if no policies exist in a namespace, all ingress and egress traffic is allowed to and from pods in that namespace.

Demo

```
...
apiVersion: networking.k8s.10/v1
kind: NetworkPolicy
metadata:
  name: default-deny-egress
spec:
  Podselector: {}
  Ingress:
    - {}
  policyTypes:
    - Ingress
```

Default **deny all** Egress traffic

Demo

```
...
apiVersion: networking.k8s.10/v1
kind: NetworkPolicy
metadata:
  name: allow-all-egress
spec:
  Podselector: {}
  egress:
    - {}
  policyTypes:
    - egress
```

Default **allow all** Egress traffic

Demo

```
...
apiVersion: networking.k8s.10/v1
kind: NetworkPolicy
metadata:
  name: default-deny-all
spec:
  Podselector: {}
  policyTypes:
    - egress
    - Ingress
```

Default **deny all** Ingress and all Egress traffic

Adding Entries to Pod /etc/hosts With HostAliases

Default Hosts File Content

Adding entries to a pod's **/etc/hosts** file provides pod-level override of hostname resolution when DNS and other options are not applicable. These custom entries can be added with the HostAliases field in PodSpec:

Demo

Command to start a Nginx Pod that is assigned to a Pod IP is as follows:

```
kubectl run nginx --image nginx
```

Result:

```
Pod/nginx created
```

Command Examine a Pod IP:

```
kubectl get Pods --output=wide
```

Result:

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
nginx	1/1	Running	0	13s	10.200.0.4	worker0



Default Hosts File Content

The following is an example of the hosts file content:

Demo

Command:

```
kubectl exec nginx -- cat /etc/hosts
```

Result:

```
# Kubernetes-managed hosts file.
127.0.0.1    localhost
::1         localhost ip6-localhost ip6-loopback
fe00::0     ip6-localnet
fe00::0     ip6-mcastprefix
fe00::1     ip6-allNodes
fe00::2     ip6-allrouters
10.200.0.4   nginx
```



By default, the **hosts** file only includes IPv4 and IPv6 boilerplates like **localhost** and its own hostname.

Additional Entries with HostAliases

In addition to the default boilerplate, more entries can be made to the **hosts** file.

Demo

```
apiVersion: v1
Kind: Pod
Metadata:
  name: hostaliases-Pod
spec:
  restartPolicy: Never
  hostAliases:
    - ip: "127.0.0.1"
      hostnames:
        - "foo.local"
        - "bar.local"
    - ip: "10.1.2.3"
      ...
```

Demo

```
...
hostnames:
  - "foo.remote"
  - "bar.remote"
containers:
  - name: cat-hosts
    image: busybox
    command:
      - cat
    args:
      - "/etc/hosts"
```

For example, to resolve **foo.local**, **bar.local** to **127.0.0.1** and **foo.remote**, **bar.remote** to **10.1.2.3**, configure HostAliases for a pod under **.spec.hostAliases**.

Additional Entries with HostAliases

A pod can be configured by running the following command:

Demo

Command:

```
kubectl apply -f https://k8s.io/examples/service/networking/hostaliases-Pod.yaml
```

Result:

```
Pod/hostaliases-Pod created
```

Examine a Pod's details to see its IPv4 address and its status, using the following command:

```
kubectl get Pod --output=wide
```

Result:

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
hostaliases-Pod	0/1	Completed	0	6s	10.200.0.5	worker0

Additional Entries with HostAliases

The **hosts** file content looks like this:

Demo

Command:

```
kubectl logs hostaliases-Pod
```

Result:

```
# Kubernetes-managed hosts file.
```

```
127.0.0.1    localhost
```

```
::1    localhost ip6-localhost ip6-loopback
```

```
fe00::0     ip6-localnet
```

```
fe00::0     ip6-mcastprefix
```

```
fe00::1     ip6-allNodes
```

```
fe00::2     ip6-allrouters
```

```
10.200.0.5   hostaliases-Pod
```

```
# Entries added by HostAliases.
```

```
127.0.0.1    foo.local      bar.local
```

```
10.1.2.3     foo.remote     bar.remote
```

with the additional entries specified at the bottom.

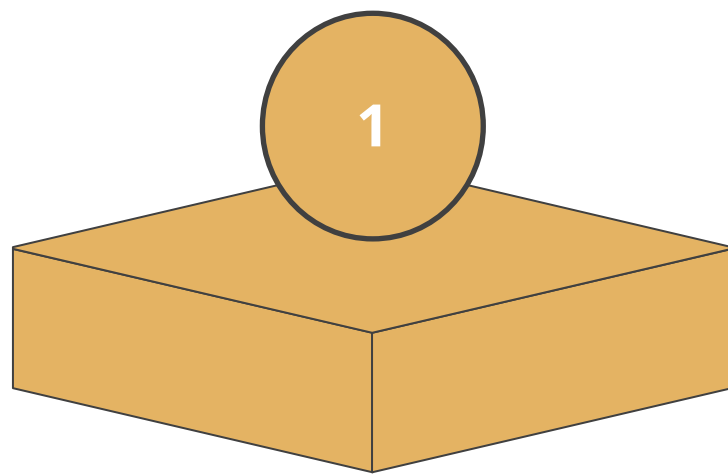


IPv4/IPv6 Dual-Stack

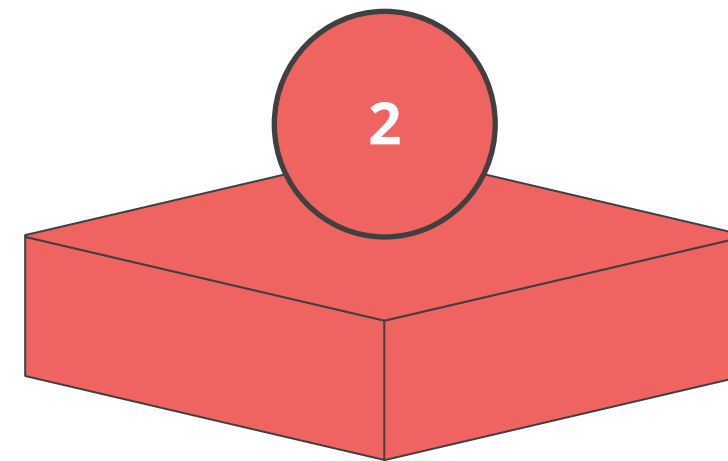
IPv4/IPv6 Dual-Stack

Dual-stacking refers to the ability of the devices to run IPv4 and IPv6 simultaneously.

IPv4/IPv6 dual-stack networking enables the allocation of both IPv4 and IPv6 addresses to pods and Services.



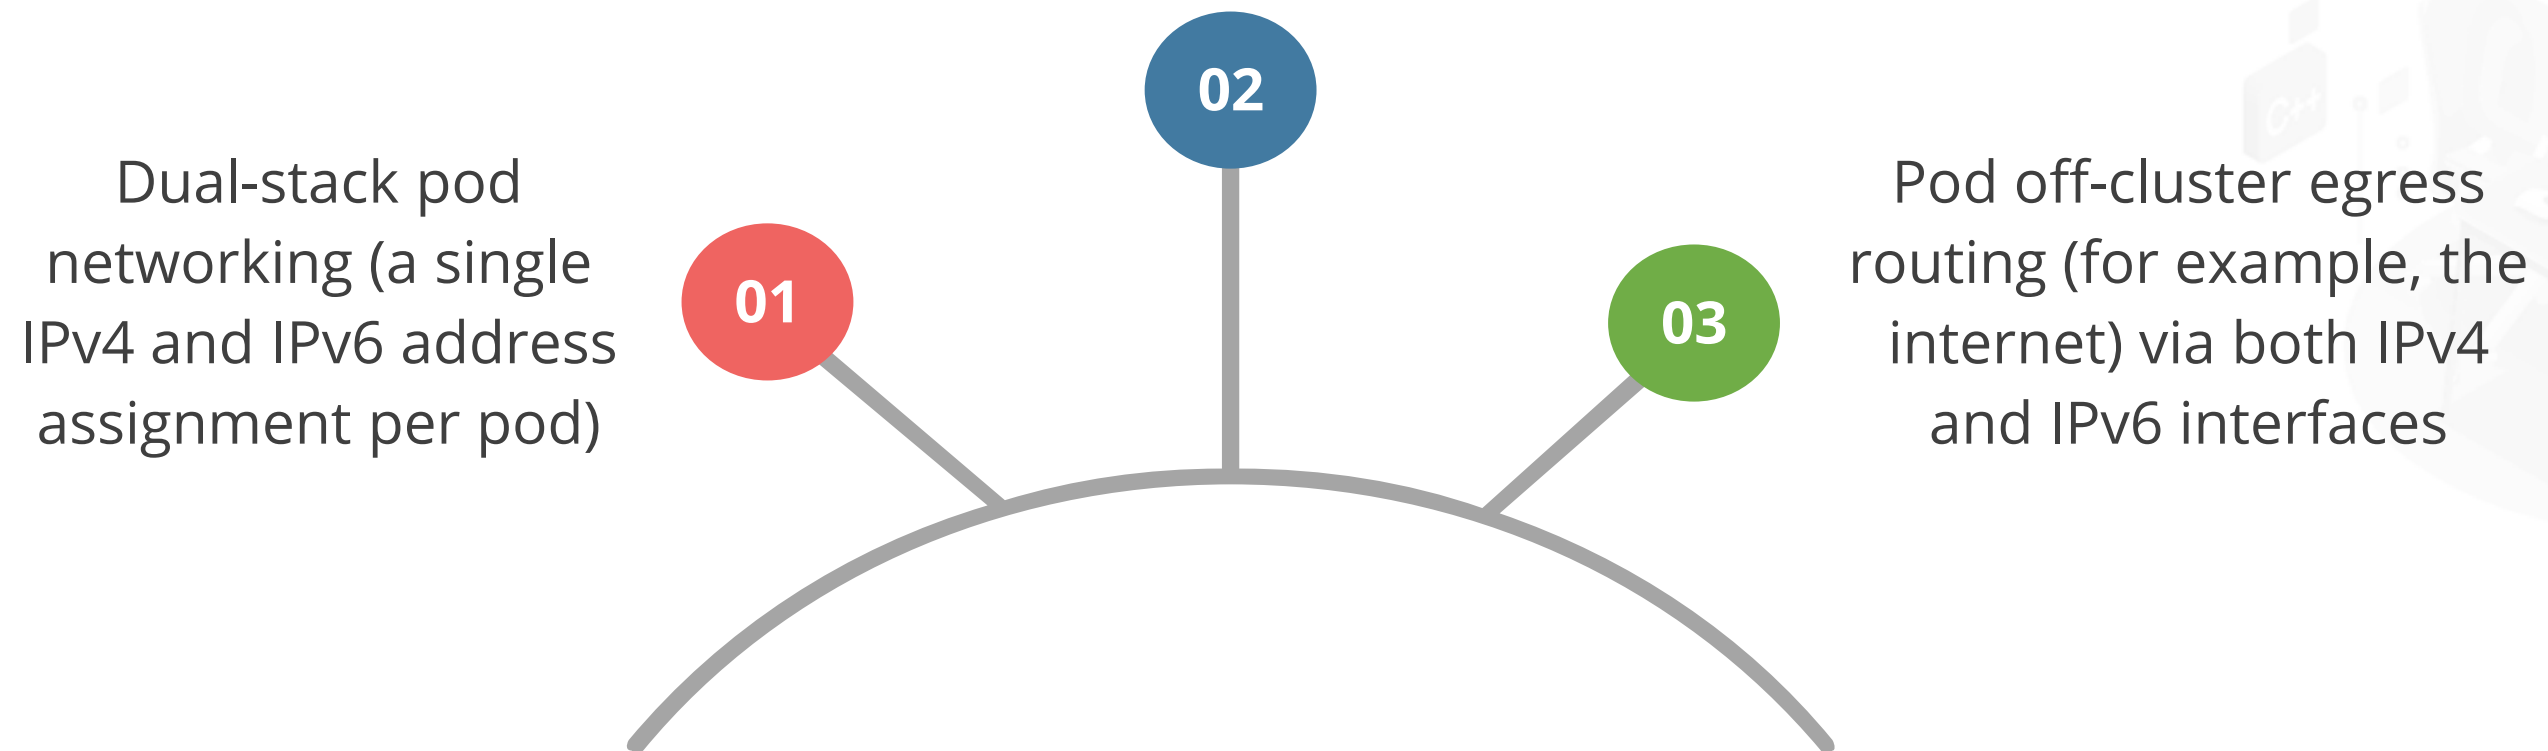
IPv4/IPv6 dual-stack networking is enabled by default for the Kubernetes, allowing simultaneous assignment of both IPv4 and IPv6 addresses.



IPv4/IPv6 Dual-Stack: Features

The IPv4/IPv6 dual-stack on the Kubernetes cluster provides the following services:

IPv4 and IPv6 enabled services



IPv4/IPv6 Dual-Stack: Prerequisites

There are three prerequisites for utilizing IPv4/IPv6 dual-stack Kubernetes clusters:

01

Kubernetes v1.20 or later

02

Provider support for dual-stack networking

03

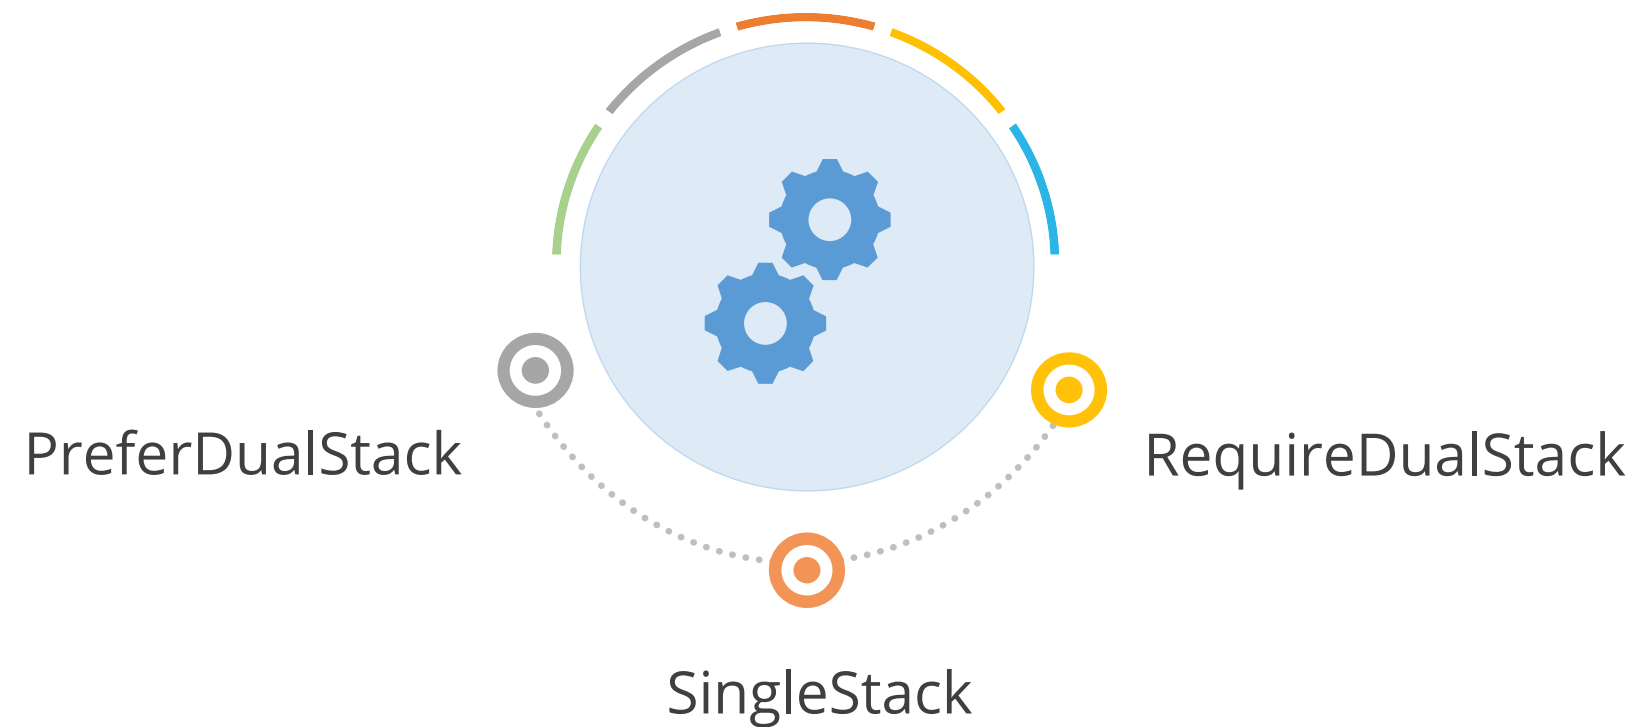
Network plugin that supports dual-stack



IPv4/IPv6 Dual-Stack: Services

When defining a service, it can optionally be configured as dual-stack by using the **.spec.ipFamilyPolicy** field.

It takes one of the following values:



Dual-Stack Options on New Services

This service specification does not define **.spec.ipFamilyPolicy** explicitly.

Demo

```
apiVersion: v1
kind: service
metadata:
  name: my-service
  labels:
    app: MyApp
spec: IPv4
  selector:
    app: MyApp
  ports:
    - protocol: TCP
      port: 80
```

When the service is created, Kubernetes assigns a cluster IP from the first configured **service-cluster-ip-range** and sets the **.spec.ipFamilyPolicy** to **SingleStack**.



Dual-Stack Options on New Services

The service specification shown below explicitly defines **PreferDualStack** in **.spec.ipFamilyPolicy**:

Demo

```
apiVersion: v1
kind: service
metadata:
  name: my-service
  labels:
    app: MyApp
spec:
  ipFamilyPolicy: PreferDualStack
  selector:
    app: MyApp
  ports:
    - protocol: TCP
      port: 80
```



Dual-Stack Options on New Services

The service specification shown below explicitly defines **IPv6** and **IPv4** in **.spec.ipFamilies** and **PreferDualStack** in **.spec.ipFamilyPolicy**:

Demo

```
apiVersion: v1
kind: service
metadata:
  name: my-service
  labels:
    app: MyApp
spec:
  ipFamilyPolicy: PreferDualStack
  ipFamilies:
  - IPv6
  - IPv4
  selector:
    app: MyApp
  ports:
  - protocol: TCP
    port: 80
```



Dual-Stack Defaults on Existing Services

The examples shown below demonstrate the default behavior of dual-stack when it is newly enabled on a cluster with services:

Demo

```
apiVersion: v1
kind: service
metadata:
  name: my-service
  labels:
    app: MyApp
spec:
  selector:
    app: MyApp
  ports:
    - protocol: TCP
      port: 80

#Validate this behavior by using the following
command:
  kubectl get svc my-service -o yaml
```

Demo

```
apiVersion: v1
kind: Service
metadata:
  labels:
    app: MyApp
  name: my-service
spec:
  clusterIP: None
  clusterIPs:
    - None
  ipFamilies:
    - IPv4
  ipFamilyPolicy: SingleStack
  ports:
    - port: 80
      protocol: TCP
      targetPort: 80
  selector:
    app: MyApp
```

Egress Traffic

To enable Egress traffic from a pod that utilizes non-publicly routable IPv6 addresses to reach off-cluster destinations (such as the public Internet), users must configure the pod to use a publicly routed IPv6 address.



Switch Services Between Single-Stack and Dual-Stack

Services can be switched between dual-stack and single-stack services and vice versa.

To change a service from single-stack to dual-stack, change **.spec.ipFamilyPolicy** from SingleStack to **PreferDualStack** or **RequireDualStack** as desired.

Before:

```
spec:  
  ipFamilyPolicy SingleStack
```

After:

```
spec:  
  ipFamilyPolicy PreferDualStack
```

Blocking All Traffic to an Application



Duration: 10 mins

Problem Statement:

You've been asked to Deny all traffic to an application.

ASSISTED PRACTICE

Assisted Practice: Guidelines

Steps to be followed:

1. Set up the application pod and policy
2. Verify the network policy



Limiting the Traffic to an Application



Duration: 15 mins

Problem Statement:

You've been asked to limit all network traffic to a specific application within a Kubernetes cluster for controlled access and resource utilization

ASSISTED PRACTICE

Assisted Practice: Guidelines

Steps to be followed:

1. Launch the API server
2. Configure the YAML file for the network policy
3. Verify the network policy
4. Clear the created resources



Blocking Traffic from Other Namespaces



Duration: 15 mins

Problem Statement:

You've been asked to implement a Kubernetes network policy for blocking all inter-namespace traffic

ASSISTED PRACTICE

Assisted Practice: Guidelines

Steps to be followed:

1. Launch the web service
2. Configure the YAML file for the network policy
3. Create a new namespace
4. Verify the network policy
5. Clear the created resources



Key Takeaways

- The Kubernetes service is a logical abstraction for a deployed group of pods in a cluster.
- Kubernetes supports two primary modes of finding a service, namely, environment variables and DNS.
- Network policies are application-centric constructs that allow to specify how a pod is allowed to communicate with various network entities over the network.
- IPv4/IPv6 dual-stack networking supports IPv4 and IPv6 addresses to be allocated to pods and services.



Implement Ingress for Multiple Containers with AKS

Duration: 25 Min



Project agenda: To implement Ingress for multiple containers with AKS

Description:

Your organization has an AKS cluster and wants to access multiple microservices using the Ingress controller. Configure AKS with Ingress to get an external URL for accessing applications based on the context root with the same external URL.

Steps to perform:

1. Creating an AKS cluster
2. Deploying the Ingress controller on an AKS cluster using Helm
3. Deploying the application deployment and service
4. Deploying the Ingress YAML for redirecting the traffic to multiple containers in a pod